

# Computer Vision 1

THEO GEVERS

MASTER AI

UNIVERSITY OF AMSTERDAM

# Lectures/Theory

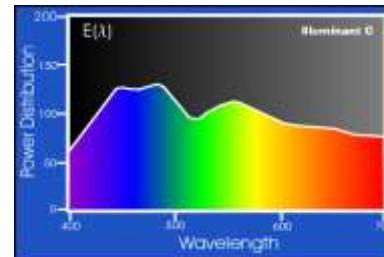
- 06-02-2018, 17:00-19:00, C0.05, **Introduction** (*Szeliski 1*)
- 13-02-2018, 17:00-19:00, C0.05, **Image Formation** (*Szeliski: 2.1.1 + 2.1.2 + 2.2 + 2.3.2 + 2.3.3*)
- 20-02-2018, 17:00-19:00, C0.05, **Color and Image Processing** (*Szeliski: 3.1 + 3.2 + 3.3*)
- 27-02-2018, 17:00-19:00, C0.05, **Feature Detection, Motion and Classification** (*Szeliski: 4, 8.1.1 + 8.1.3 + 8.2.1 + 8.4; Bengio: 4 + 5.1 + 5.2 + 5.3 + 5.7 + 5.8 + 5.9*)
- 06-03-2018, 17:00-19:00, C0.05, **Object Recognition: BoW and ConvNets** (*Szeliski: 5.1.1 + 5.1.4 + 5.1.5 + 5.2 + 5.3 + 5.4, 6.1 + 6.3, 14.1 + 14.2.1 + 14.3 + 14.4.1; Bengio: 7.2 + 7.4 + 9.1 + 9.2 + 9.3*)
- 13-03-2018, 17:00-19:00, C0.05, **Deep Learning, Stereo and 3D Reconstruction** (*Szeliski: 11.1 + 11.2 + 11.3 + 11.4, 12.1 + 12.2; Bengio: 12.1 + 12.2*)
- 20-03-2018, 17:00-19:00, C0.05, **Applications** (*Szeliski: 12.6.2 + 12.6.3 + 12.2.4*)
- 26-03-2018, Monday, 9:00-12:00, **Written Exam**

# Today's Class

- 1. Reflection Models**
- 2. Color Invariance**
- 3. Image Processing**

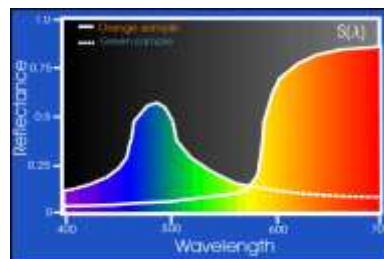
# What makes an image

Light source



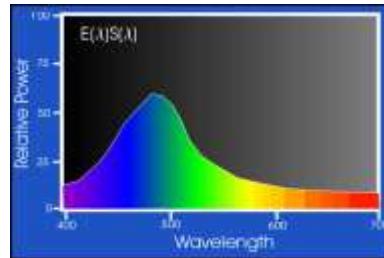
$$e(\lambda)$$

Object



$$\rho(\lambda)$$

Sensor



$$f_c(\lambda)$$

$$R = \int_{\lambda} e(\lambda) \rho(\lambda) f_R(\lambda) d\lambda, \quad G = \int_{\lambda} e(\lambda) \rho(\lambda) f_G(\lambda) d\lambda, \quad B = \int_{\lambda} e(\lambda) \rho(\lambda) f_B(\lambda) d\lambda$$

# Recap: What makes an image?

the triplet light-objects-observer

$$R = \int_{\lambda} e(\lambda) \rho(\lambda) f_R(\lambda) d\lambda, \quad G = \int_{\lambda} e(\lambda) \rho(\lambda) f_G(\lambda) d\lambda, \quad B = \int_{\lambda} e(\lambda) \rho(\lambda) f_B(\lambda) d\lambda$$

# Basic Principles of Surface Reflectance

Slides from Shree Nayar, Ravi Ramamoorthi, Pat Hanrahan

# Surface Appearance

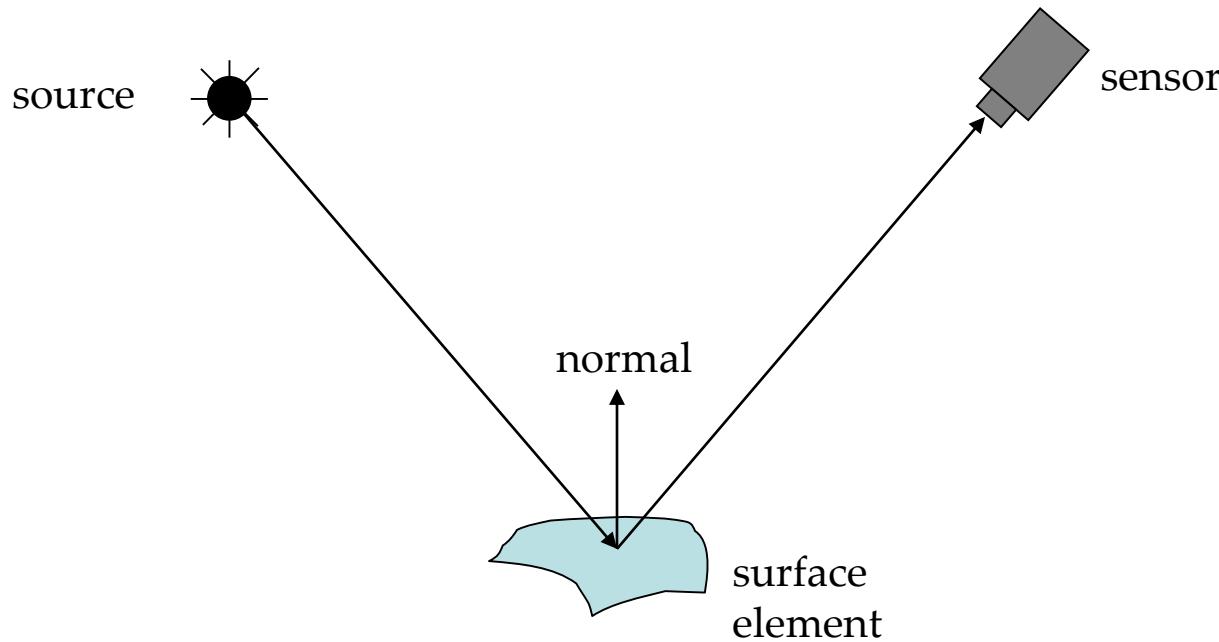
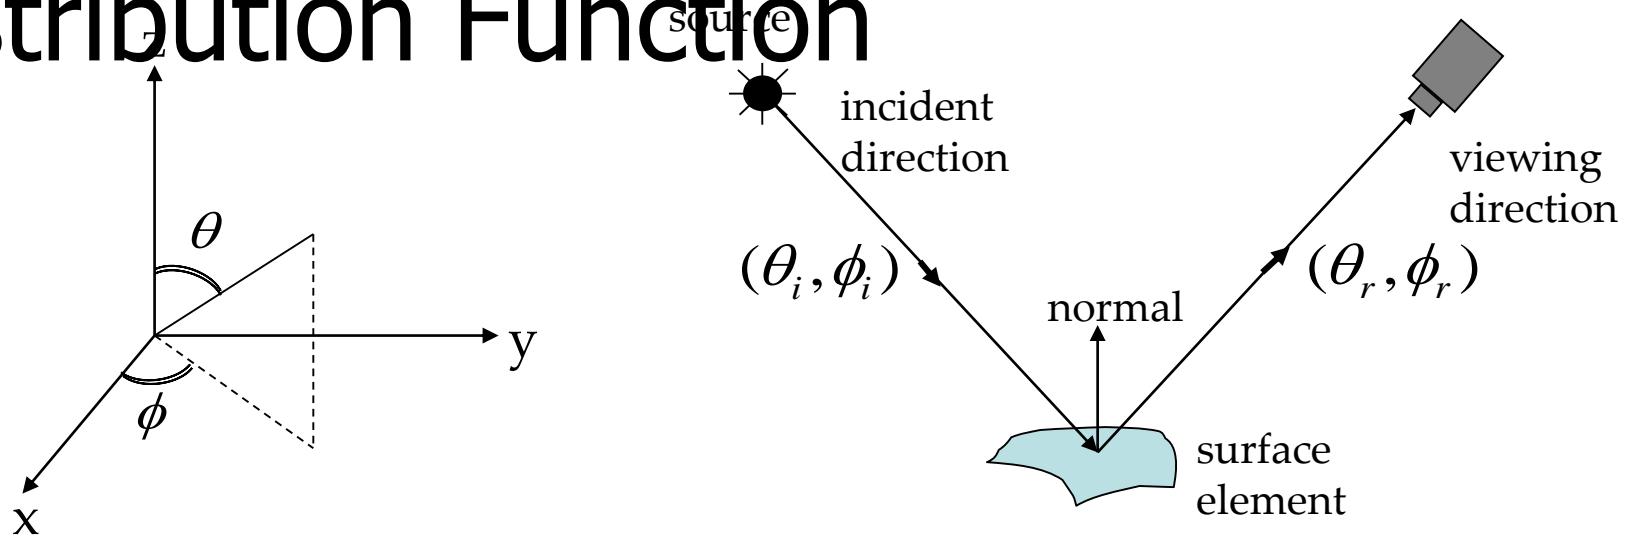


Image intensities =  $f(\text{normal}, \text{surface reflectance}, \text{illumination})$

Surface Reflection depends on both the viewing and illumination direction.

# BRDF: Bidirectional Reflectance Distribution Function

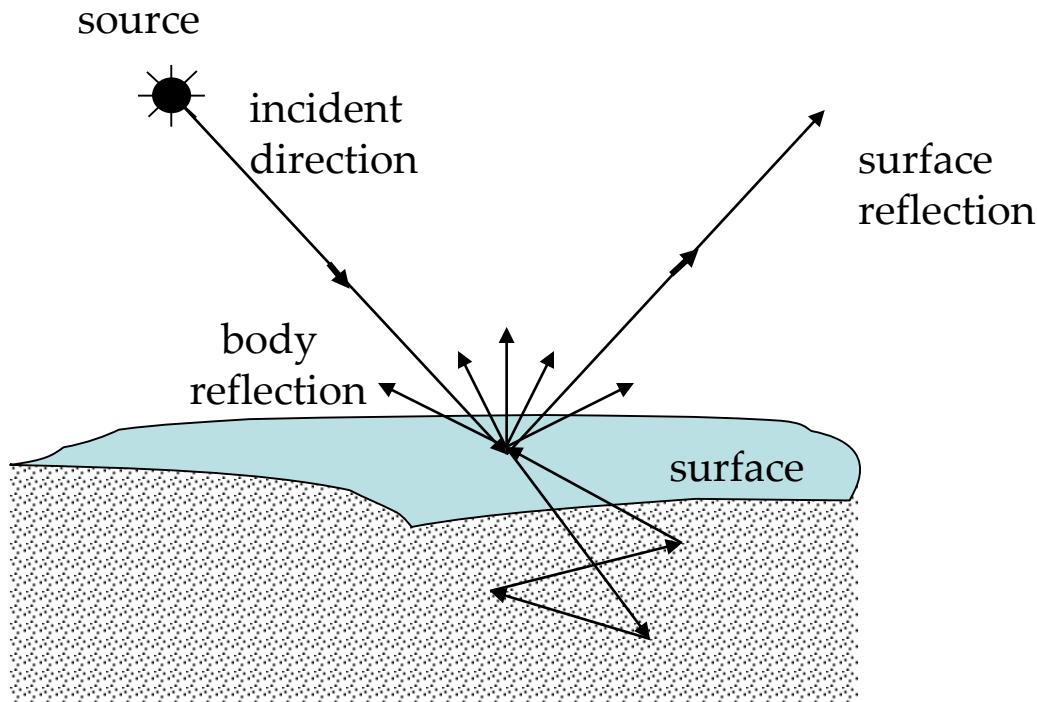


$E^{surface}(\theta_i, \phi_i)$  Irradiance at Surface in direction  $(\theta_i, \phi_i)$

$L^{surface}(\theta_r, \phi_r)$  Radiance of Surface in direction  $(\theta_r, \phi_r)$

$$\text{BRDF : } f(\theta_i, \phi_i; \theta_r, \phi_r) = \frac{L^{surface}(\theta_r, \phi_r)}{E^{surface}(\theta_i, \phi_i)}$$

# Mechanisms of Surface Reflection



Body Reflection:

- Diffuse Reflection
- Matte Appearance
- Non-Homogeneous Medium
- Clay, paper, etc

Surface Reflection:

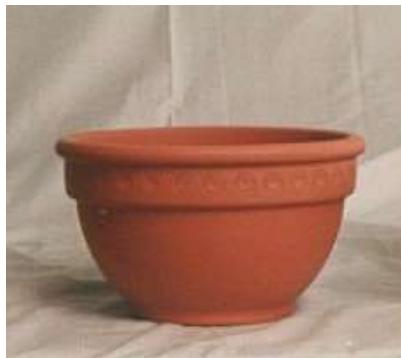
- Specular Reflection
- Glossy Appearance
- Highlights
- Dominant for Metals

Image Intensity = Body Reflection + Surface Reflection

# Mechanisms of Surface Reflection

Body Reflection:

Diffuse Reflection  
Matte Appearance  
Non-Homogeneous Medium  
Clay, paper, etc



Surface Reflection:

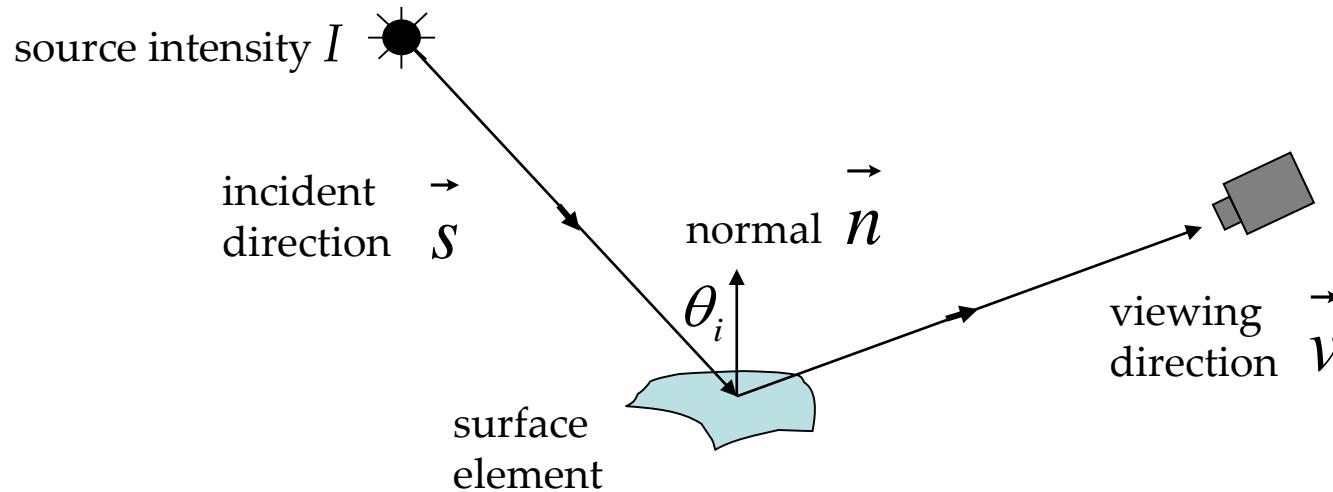
Specular Reflection  
Glossy Appearance  
Highlights  
Dominant for Metals



Many materials exhibit both Reflections:

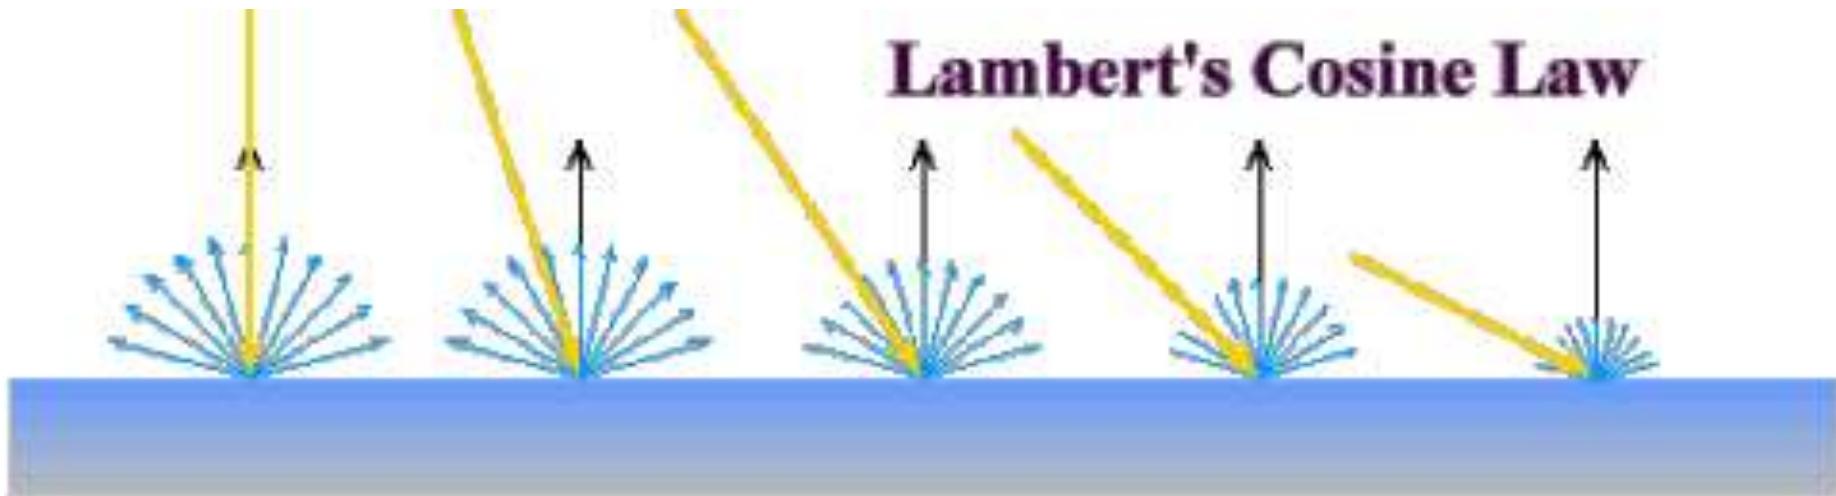


# Diffuse Reflection and Lambertian BRDF



- Surface appears equally bright from ALL directions! (independent of  $\vec{v}$ )
- Lambertian BRDF is simply a constant:  $f(\theta_i, \phi_i; \theta_r, \phi_r) = \frac{\rho_d}{\pi}$  albedo
- Surface Radiance:  $L = \frac{\rho_d}{\pi} I \cos \theta_i = \frac{\rho_d}{\pi} I \vec{n} \cdot \vec{s}$  source intensity
- Commonly used in Vision and Graphics!

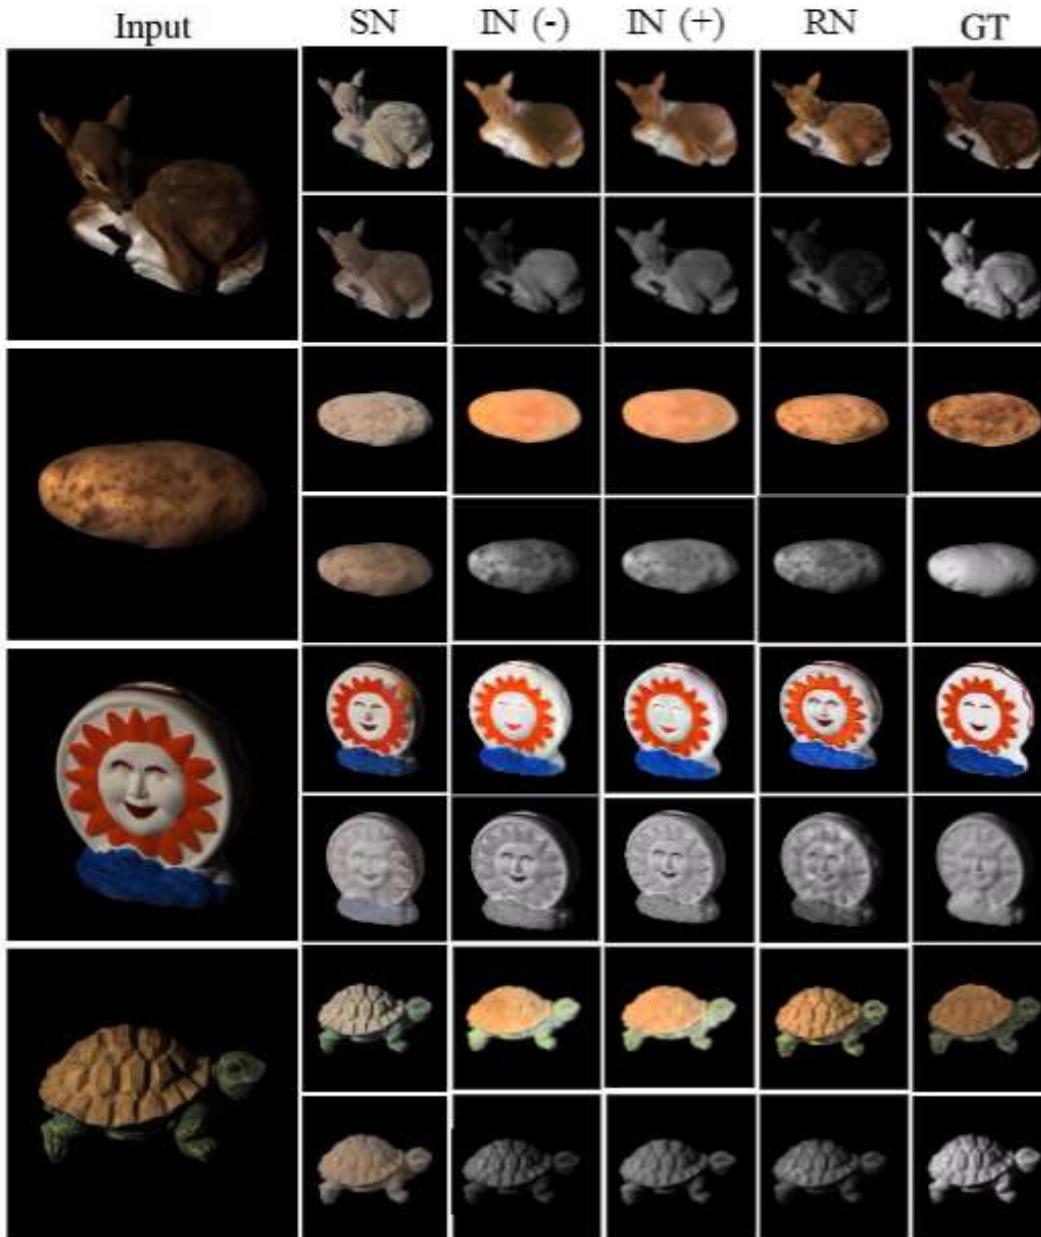
# Diffuse Reflection and Lambertian BRDF



# Diffuse Reflection and Lambertian BRDF



# MIT Dataset (Small)



Surface Radiance:

$$I = \rho I_i \cos \theta = \rho I_i \vec{n} \cdot \vec{s}$$

$$I = \rho S, \text{ where } S = I_i \vec{n} \cdot \vec{s}$$

$$I = R S, \text{ where } R = \rho$$

$$I = R S$$

Intensity = reflectance x shading

# ShapeNet Rendered Dataset (1M)

Input



Albedo GT



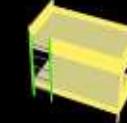
Shading GT



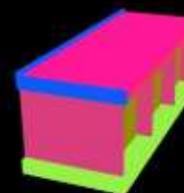
Input

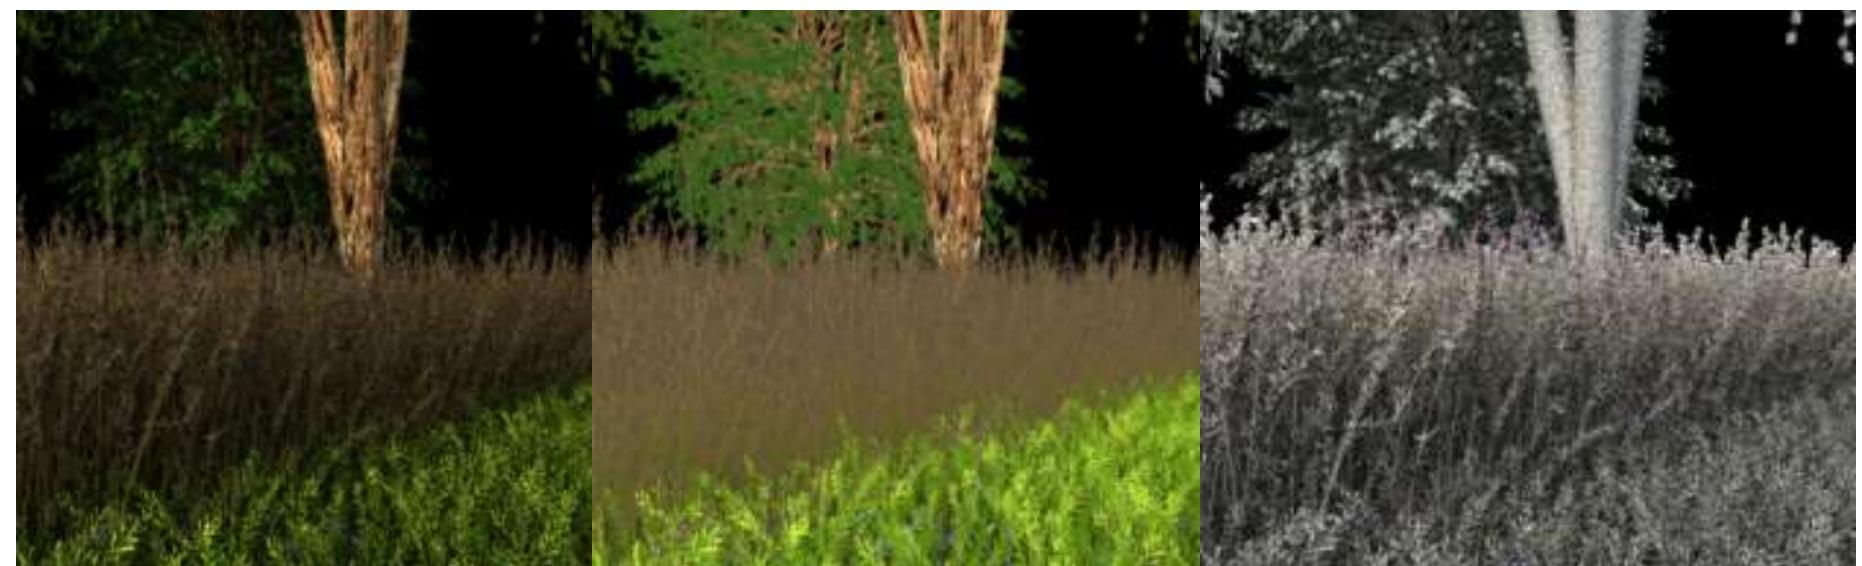
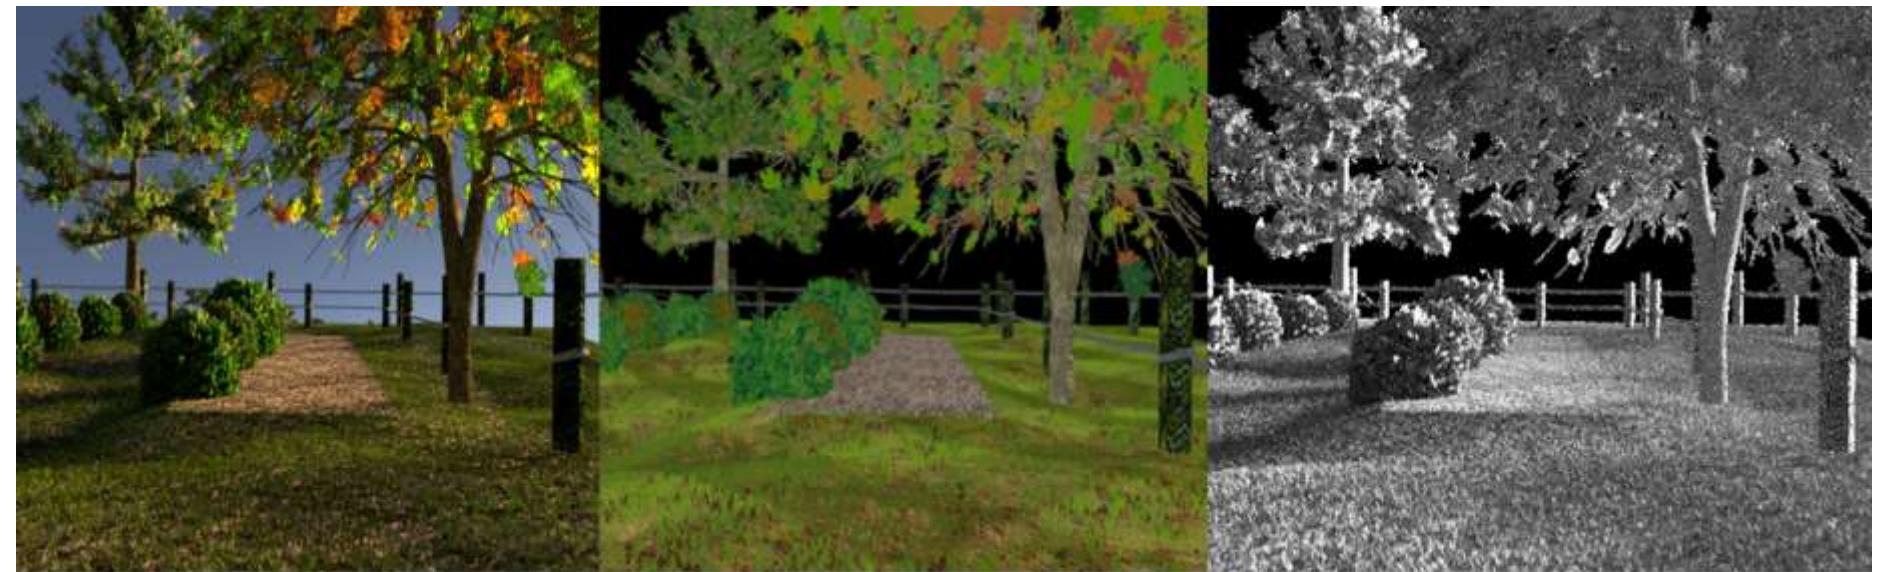


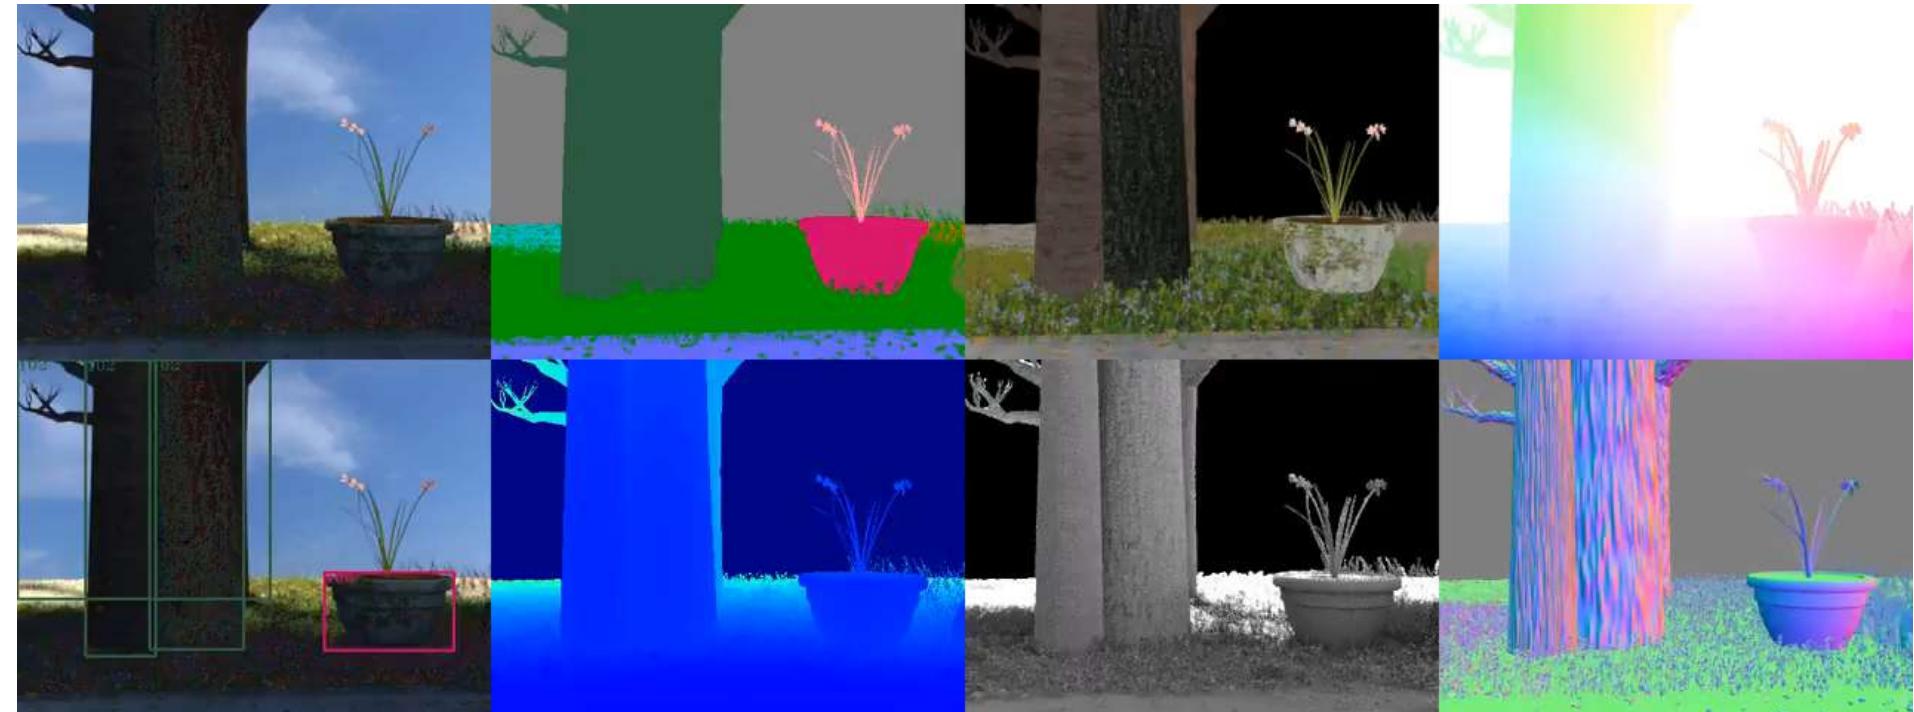
Albedo GT



Shading GT







Video link: <https://youtu.be/I5U85sGdanA>

# Face Analysis

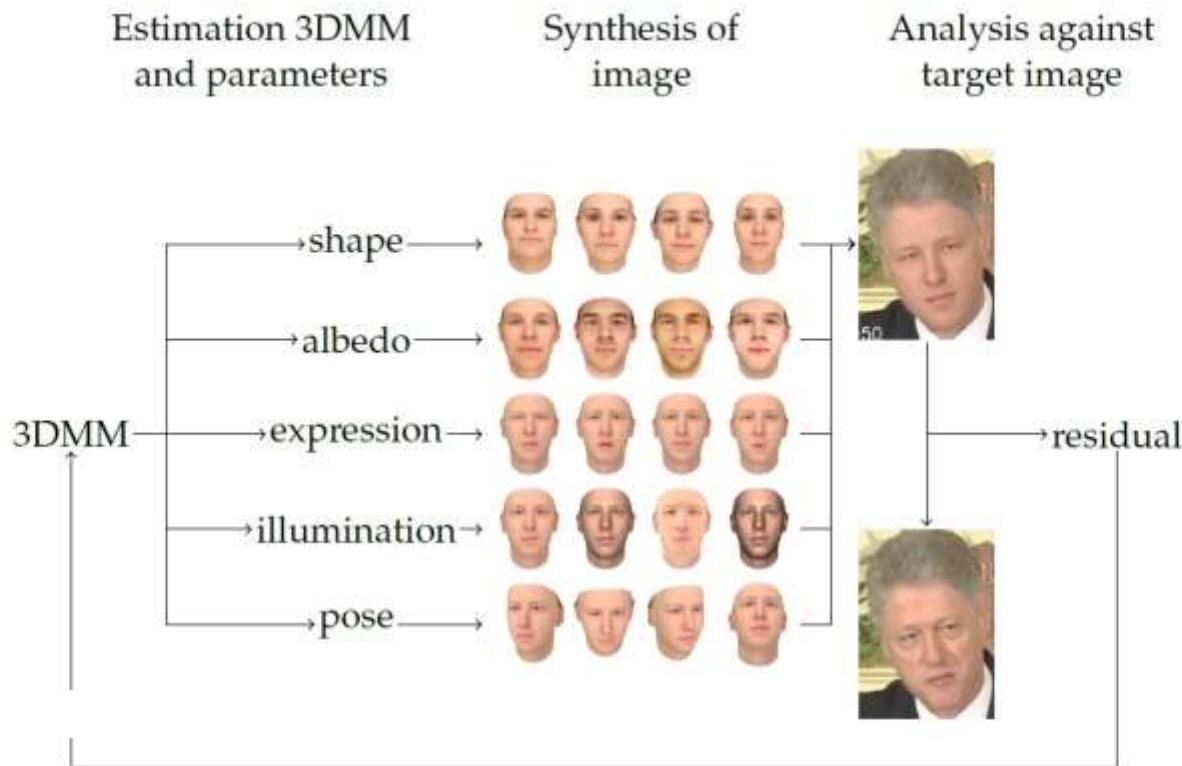
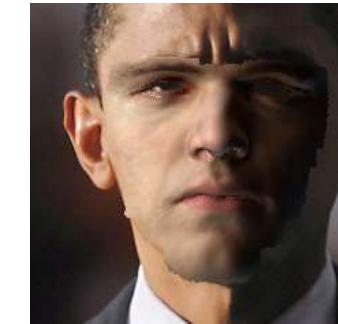
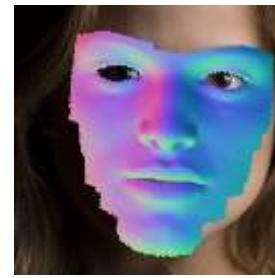
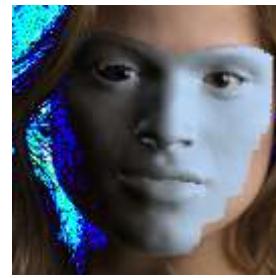
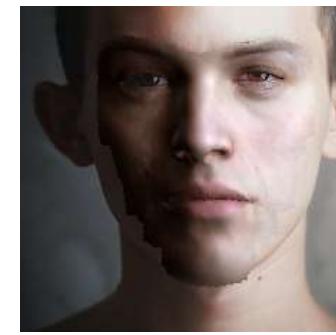
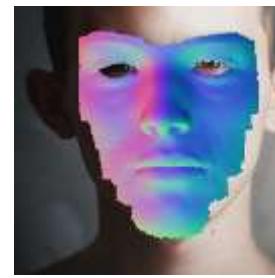
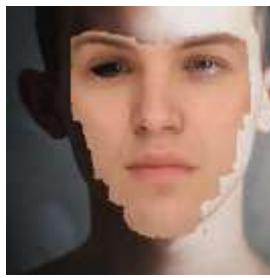


FIGURE 1.1: Example of 3DMM fitting following an analysis-by-synthesis approach. From left to right: The 3D Morphable Model, prediction of parameters for the different components of the optimization problem and rendered results, analyses against target image, and feedback loop.



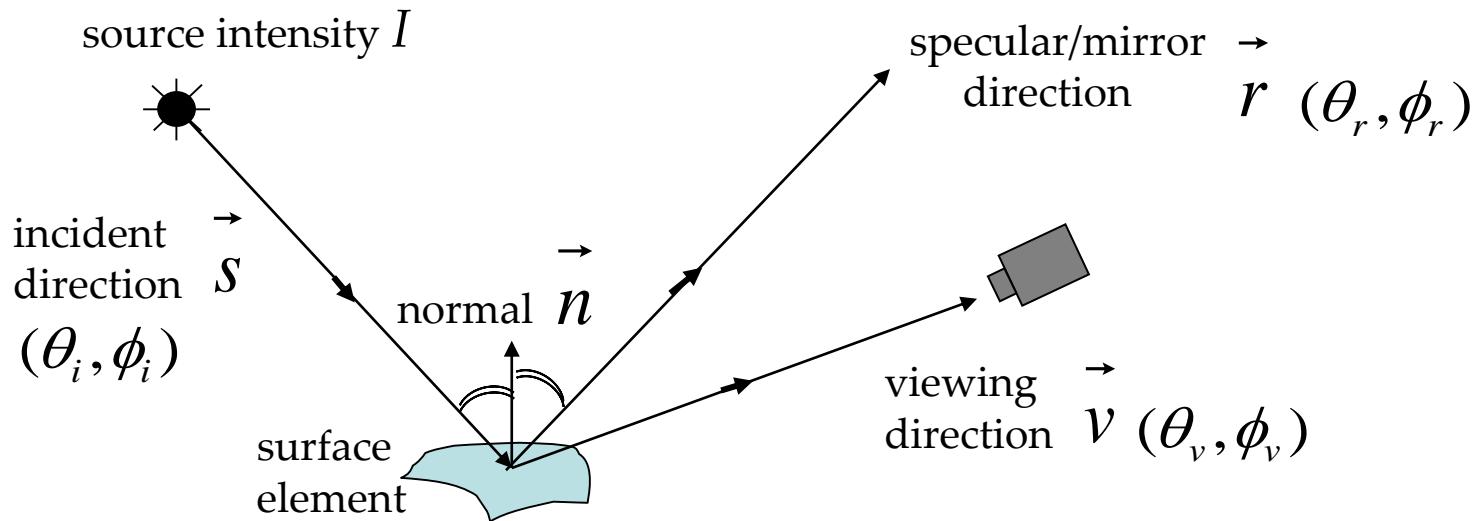








# Specular Reflection and Mirror BRDF



- Very smooth surface.
- All incident light energy reflected in a SINGLE direction. (only when  $\vec{v} = \vec{r}$ )
- Mirror BRDF is simply a double-delta function :

$$f(\theta_i, \phi_i; \theta_v, \phi_v) = \rho_s \delta(\theta_i - \theta_v) \delta(\phi_i + \pi - \phi_v)$$

specular albedo

- Surface Radiance :  $L = I \rho_s \delta(\theta_i - \theta_v) \delta(\phi_i + \pi - \phi_v)$

# Glossy Surfaces

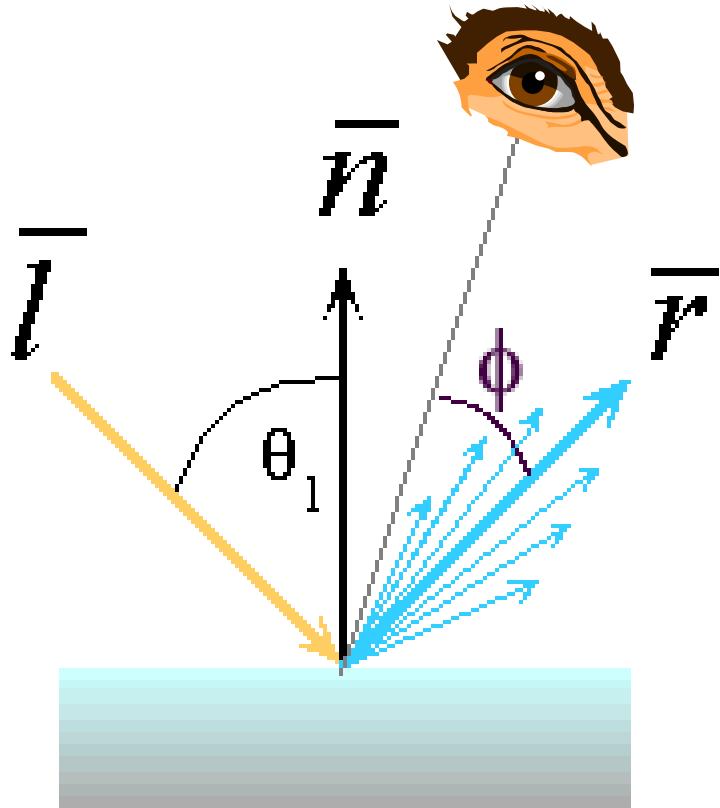
- Many glossy surfaces show broader highlights in addition to specular reflection.



- Example Models : Phong Model (no physical basis, but sort of works (empirical)

# Phong Model: An Empirical Approximation

- An illustration of the angular falloff of highlights:

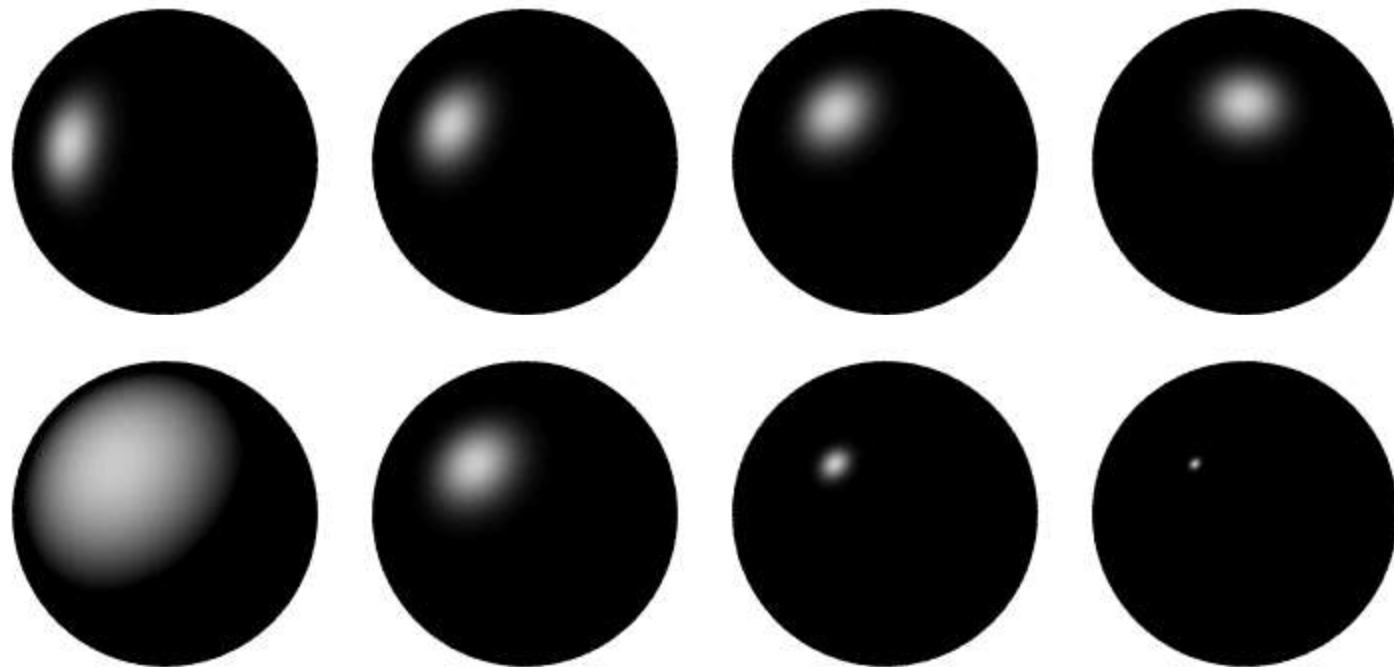


$$L = I \rho_s (\cos \phi)^{n_{shiny}}$$

- Very commonly used in Computer Graphics

# Phong Examples

- These spheres illustrate the Phong model as *lighting direction* and  $n_{shiny}$  are varied:



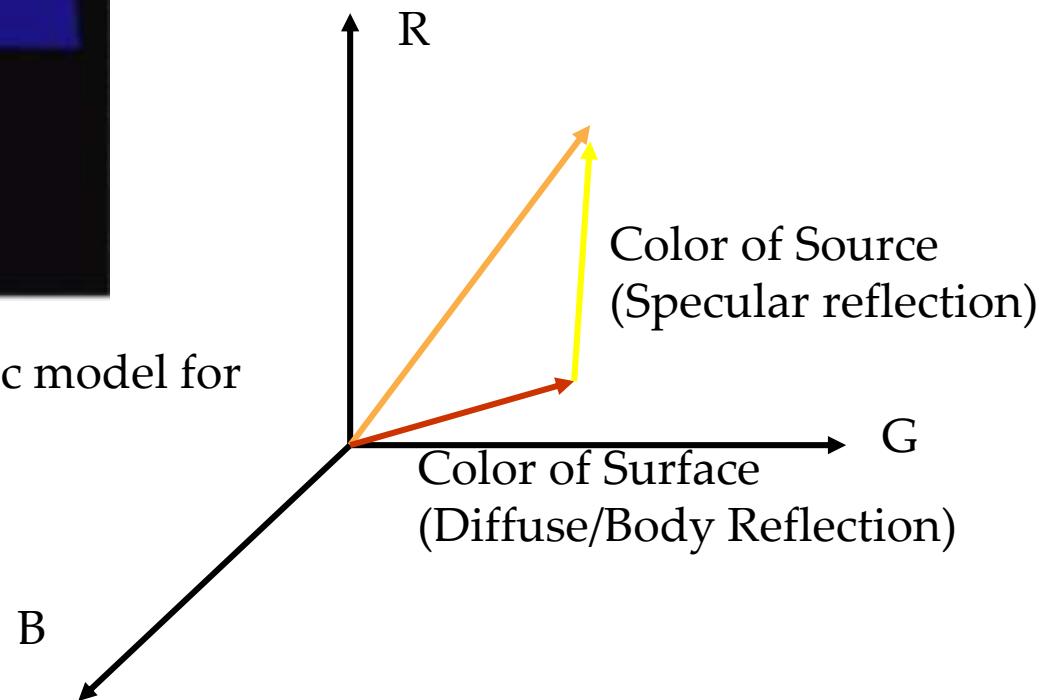
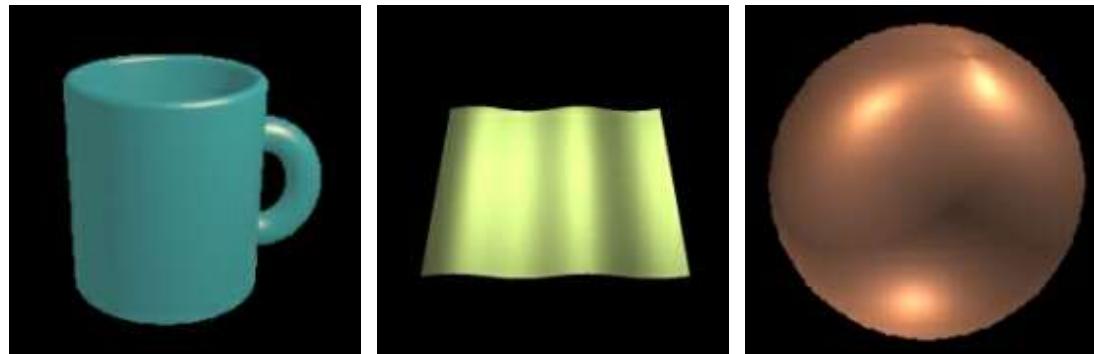
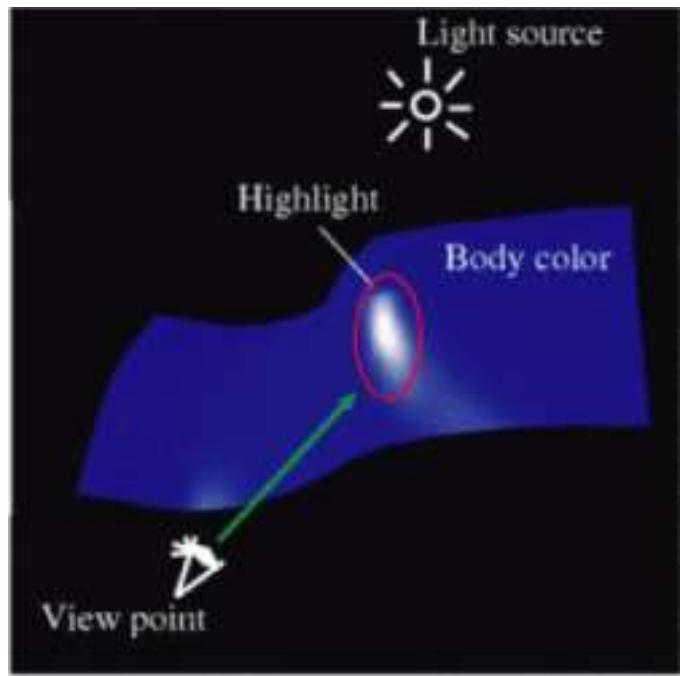
# Surface Reflection



# A Simple Reflection Model - Dichromatic Reflection

Observed Image Color =  $a \times$  Body Color +  $b \times$  Specular Reflection Color

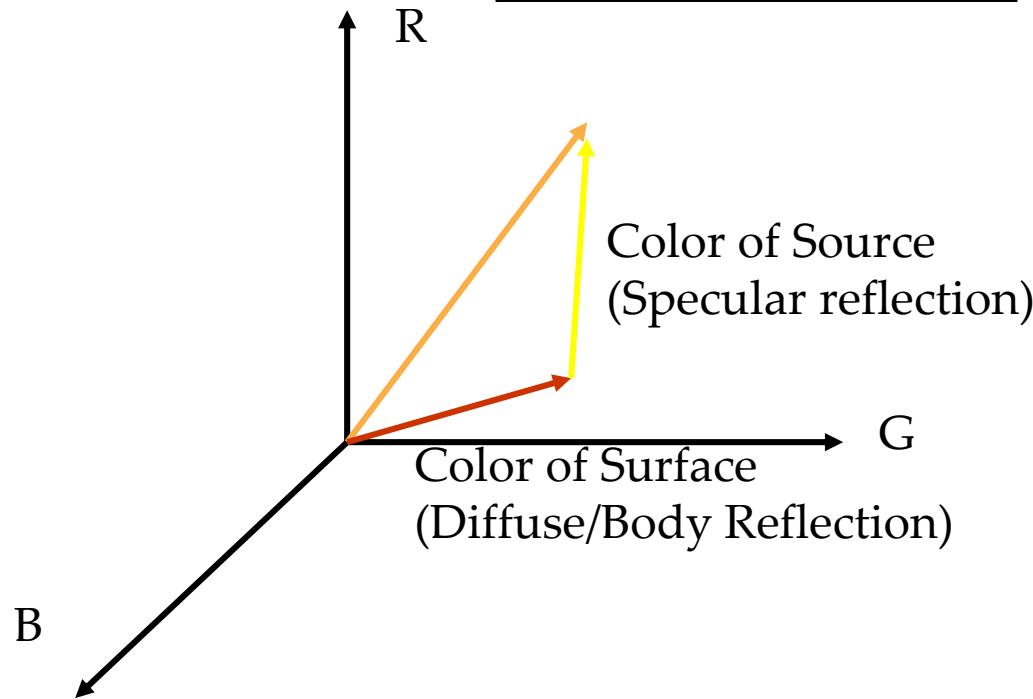
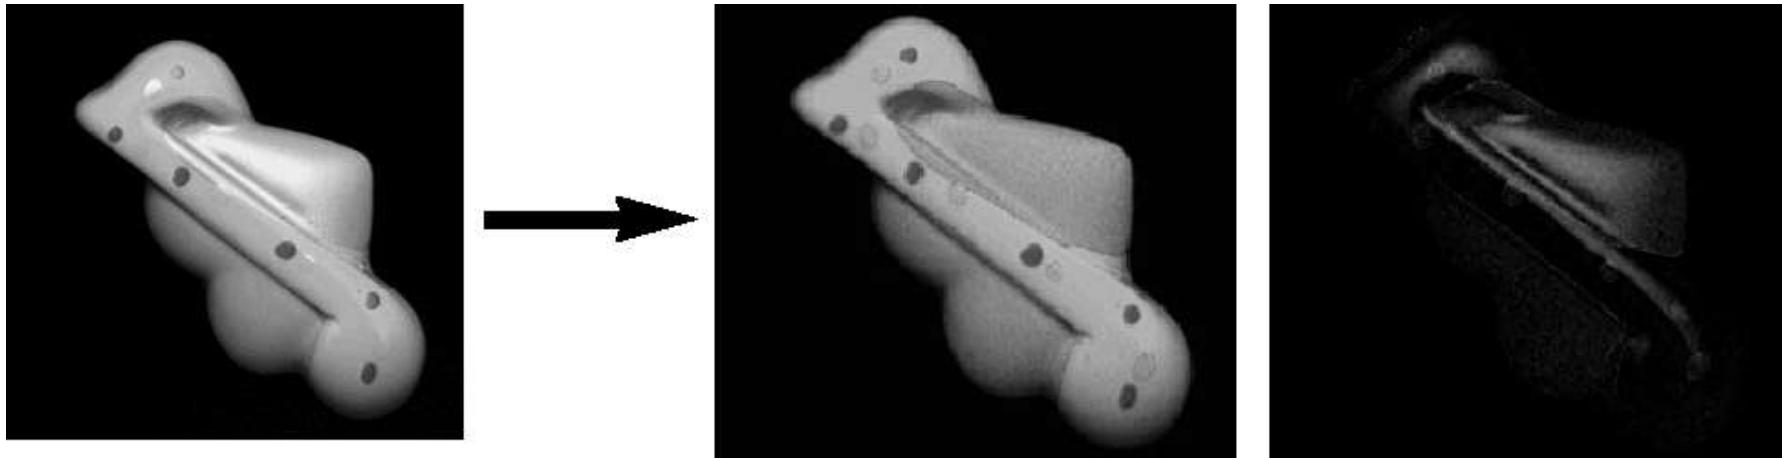
Klinker-Shafer-Kanade 1988



Does not specify any specific model for  
Diffuse/specular reflection

# Separating Diffuse and Specular Reflections

Observed Image Color =  $a \times$  Body Color +  $b \times$  Specular Reflection Color



# Today's Class

- 1. Reflection Models**
- 2. Color Invariance**
- 3. Image Processing**

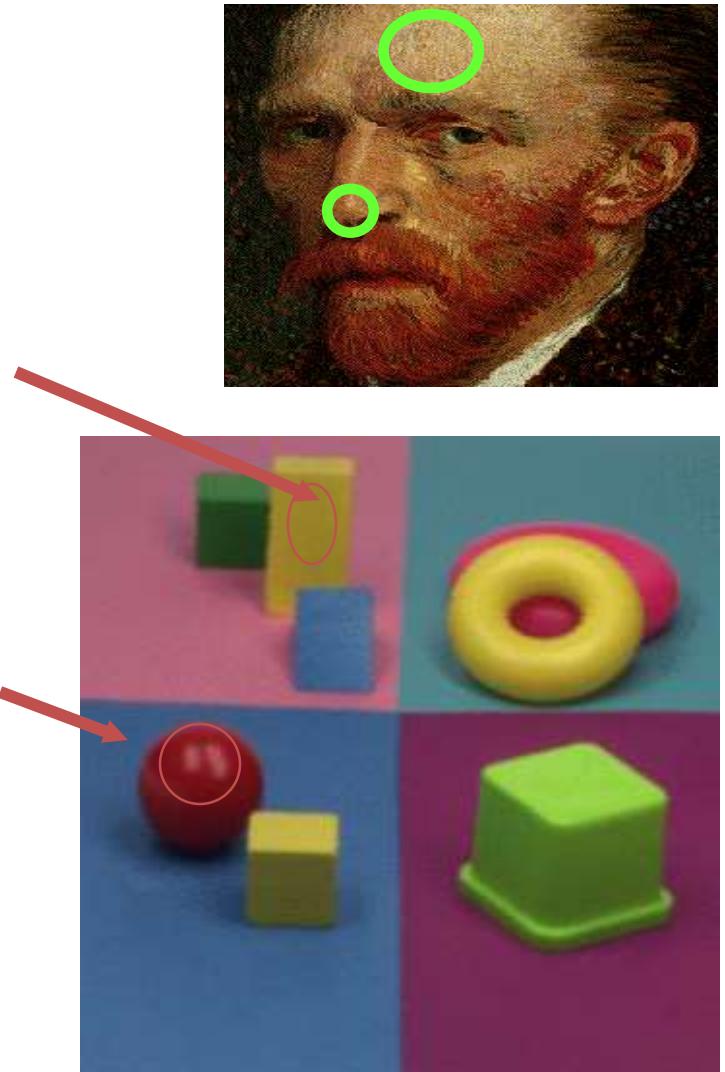
# Reflection Model

Dichromatic reflection model [Shafer]

$$\text{body} = m_b(\mathbf{n}, \mathbf{s}) \int_{\lambda} e(\lambda) \rho(\lambda) f_C(\lambda) d\lambda$$

$$\text{surface} = m_s(\mathbf{n}, \mathbf{s}, \mathbf{v}) \int_{\lambda} e(\lambda) c(\lambda) f_C(\lambda) d\lambda$$

for {R,G,B} giving an R-, B-, G-sensor response



# Reflection Model



$$C = m_b(\mathbf{n}, \mathbf{s}) \int_{\lambda} e(\lambda) \rho(\lambda) f_c(\lambda) d\lambda + m_s(\mathbf{n}, \mathbf{s}, \mathbf{v}) \int_{\lambda} e(\lambda) c(\lambda) f_c(\lambda) d\lambda$$

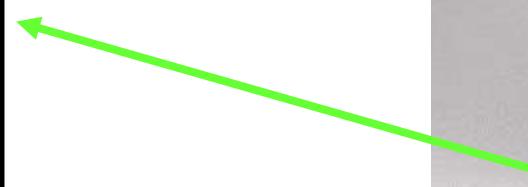
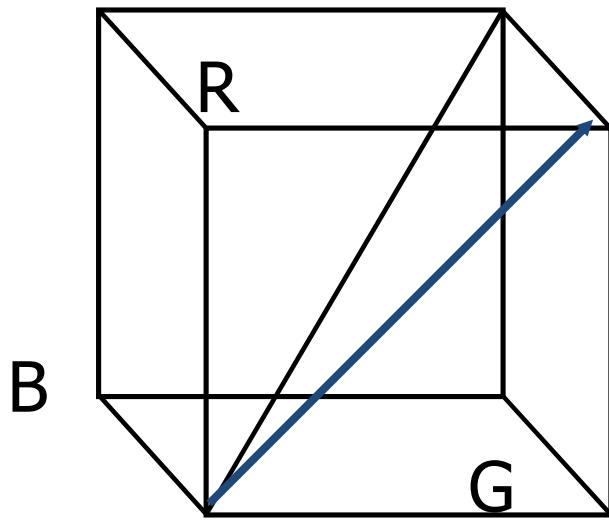
$\rho(\lambda)$	surface albedo	scene & viewpoint invariant
$e(\lambda)$	illumination	scene dependent
$\mathbf{n}$	object surface normal	object shape variant
$\mathbf{s}$	illumination direction	scene dependent
$\mathbf{v}$	viewer's direction	viewpoint variant
$f_c(\lambda)$	sensor sensitivity	camera dependent
$c(\lambda)$	specular	surface dependent

# Body Reflectance in RGB - Space

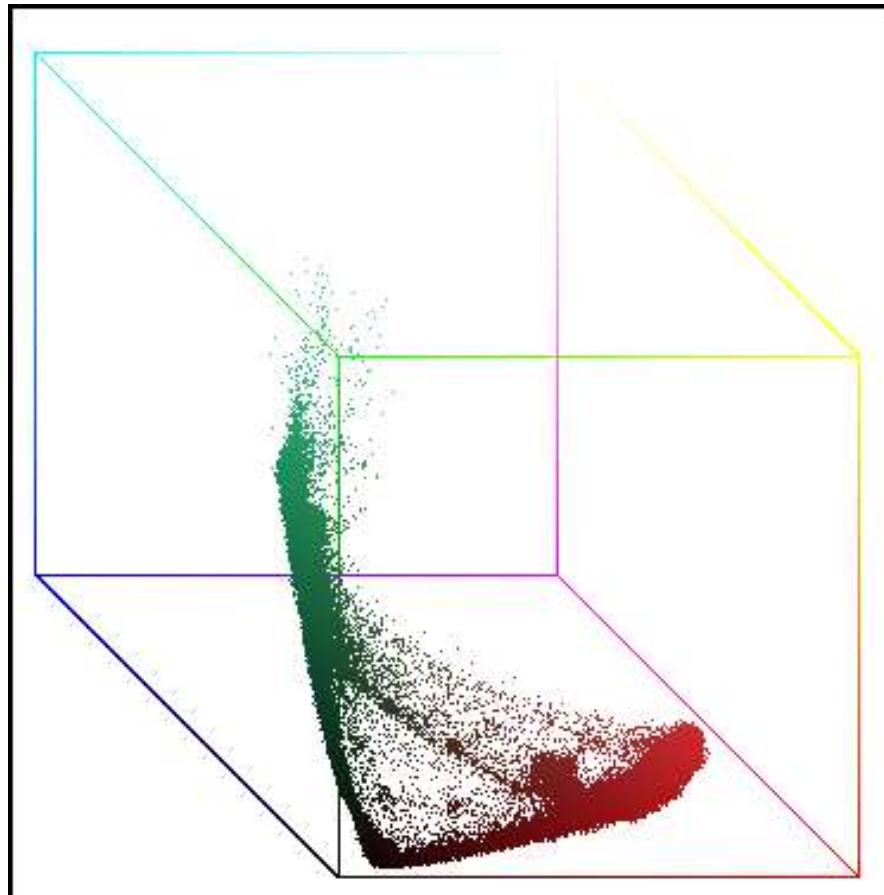
Consider the body reflection term:

$$\begin{bmatrix} R \\ G \\ B \end{bmatrix} = \begin{bmatrix} m_b(\mathbf{n}, \mathbf{s}) \int_{\lambda} e(\lambda) \rho(\lambda) f_R(\lambda) d\lambda \\ m_b(\mathbf{n}, \mathbf{s}) \int_{\lambda} e(\lambda) \rho(\lambda) f_G(\lambda) d\lambda \\ m_b(\mathbf{n}, \mathbf{s}) \int_{\lambda} e(\lambda) \rho(\lambda) f_B(\lambda) d\lambda \end{bmatrix}$$

for example the geometric term is Lambertian i.e.  $m_b(\mathbf{n}, \mathbf{s}) = \mathbf{n} \cdot \mathbf{s} = \cos \theta$



# Body Reflectance in RGB - Matte Surfaces



An object with its representation in RGB-colour space.

# Reflectance under White Light

Dichromatic reflection:

$$C = m_b(\mathbf{n}, \mathbf{s}) \int e(\lambda) \rho(\lambda) f_C(\lambda) d\lambda + m_s(\mathbf{n}, \mathbf{s}, \mathbf{v}) \int e(\lambda) c(\lambda) f_C(\lambda) d\lambda$$

Considering dichromatic reflectance and white illumination,  
then  $e = e(\lambda)$  and  $c = c(\lambda)$ . Then,

for  $C_w = \{R, G, B\}$  giving the red, green and blue sensor  
response under white light. This gives:

$$\begin{bmatrix} R_w \\ G_w \\ B_w \end{bmatrix} = \begin{bmatrix} em_b(\mathbf{n}, \mathbf{s}) \int_{\lambda} \rho(\lambda) f_R(\lambda) d\lambda + em_s(\mathbf{n}, \mathbf{s}, \mathbf{v}) cf_R \\ em_b(\mathbf{n}, \mathbf{s}) \int_{\lambda} \rho(\lambda) f_G(\lambda) d\lambda + em_s(\mathbf{n}, \mathbf{s}, \mathbf{v}) cf_G \\ em_b(\mathbf{n}, \mathbf{s}) \int_{\lambda} \rho(\lambda) f_B(\lambda) d\lambda + em_s(\mathbf{n}, \mathbf{s}, \mathbf{v}) cf_B \end{bmatrix}$$

# rgb – Photometric Invariance: Proof

Consider the body reflection term:

$$C_{WB} = em_b(\mathbf{n}, \mathbf{s}) \int \rho(\lambda) f_c(\lambda) d\lambda$$

For  $C_{WB} = \{R_{WB}, G_{WB}, B_{WB}\}$  giving the red, green and blue sensor response under white light. Consider normalized color:

$$r = \frac{R}{R+G+B}, g = \frac{G}{R+G+B}, b = \frac{B}{R+G+B}$$

Then:

$$r = \frac{\cancel{em_b(\mathbf{n}, \mathbf{s})} \int \rho(\lambda) f_R(\lambda) d\lambda}{\cancel{em_b(\mathbf{n}, \mathbf{s})} \left( \int \rho(\lambda) f_R(\lambda) d\lambda + \int \rho(\lambda) f_G(\lambda) d\lambda + \int \rho(\lambda) f_B(\lambda) d\lambda \right)}$$

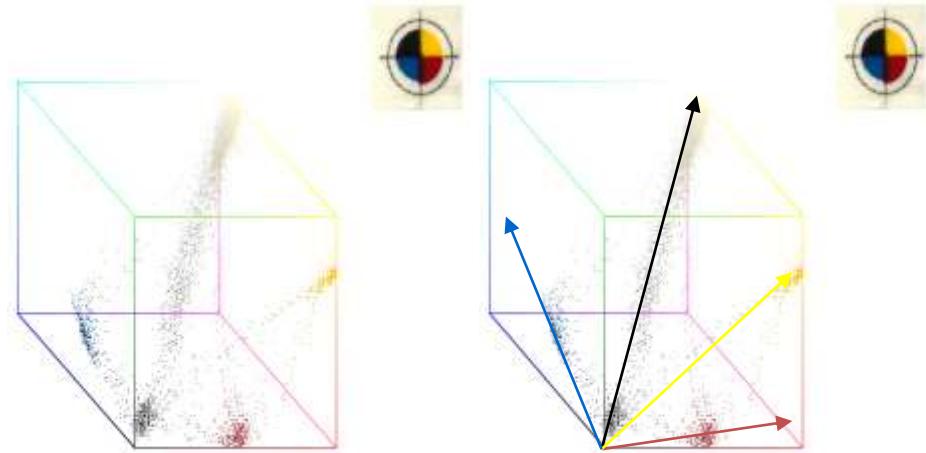
Also holds for g and b. In the sequel:  $k_c = \int \rho(\lambda) f_c(\lambda) d\lambda$

# Colour Invariance / c1c2c3 space

$$c_1(R, G, B) = \arctan \frac{R}{\max\{G, B\}}$$

$$c_2(R, G, B) = \arctan \frac{G}{\max\{R, B\}}$$

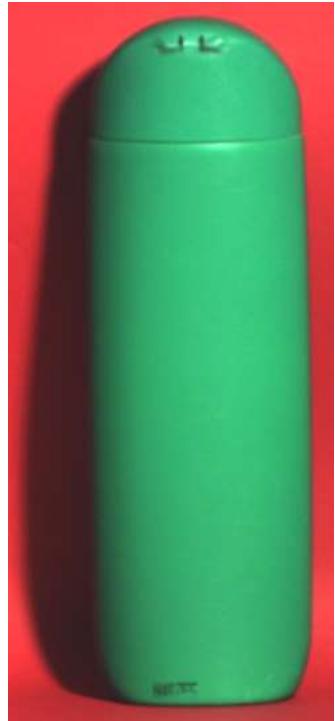
$$c_3(R, G, B) = \arctan \frac{B}{\max\{R, G\}}$$



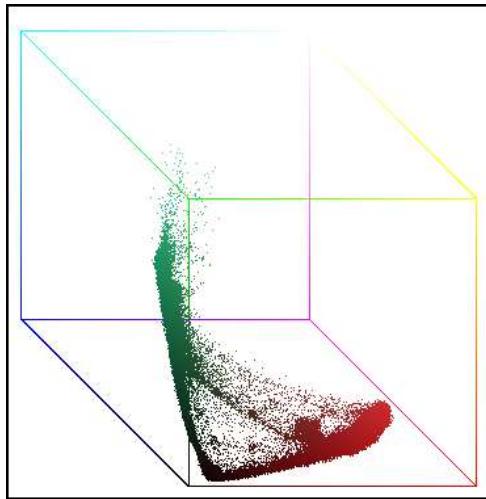
$$c_1(R_b, G_b, B_b) = \arctan \left( \frac{em_b(\mathbf{n}, \mathbf{s}) k_R}{\max\{em_b(\mathbf{n}, \mathbf{s}) k_G, em_b(\mathbf{n}, \mathbf{s}) k_B\}} \right) = \arctan \left( \frac{k_R}{\max\{k_G, k_B\}} \right)$$

$$\text{Where : } k_C = \int_{\lambda} \rho(\lambda) f_C(\lambda) d\lambda$$

# Colour Invariance / rgb space



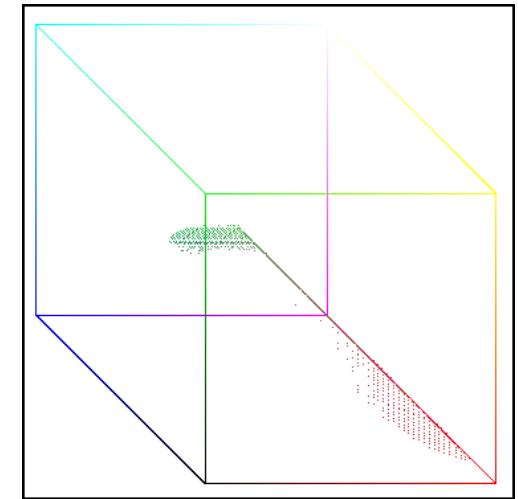
Original RGB  
image



3D plot of  
RGB image

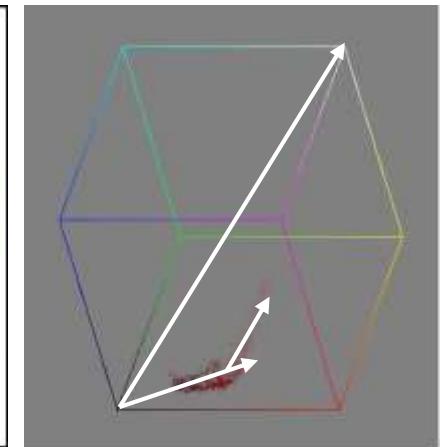
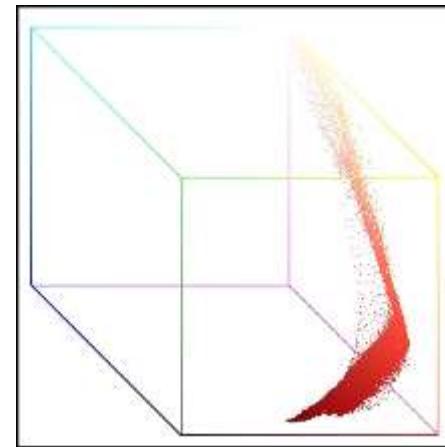
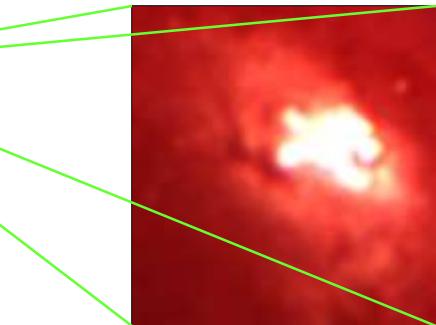
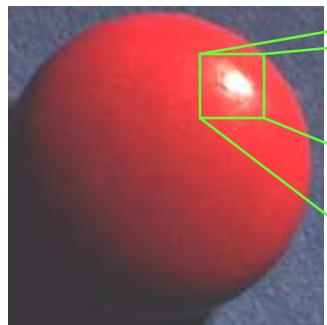


rgb image



3D plot of rgb  
image

# Body Reflectance in RGB - space



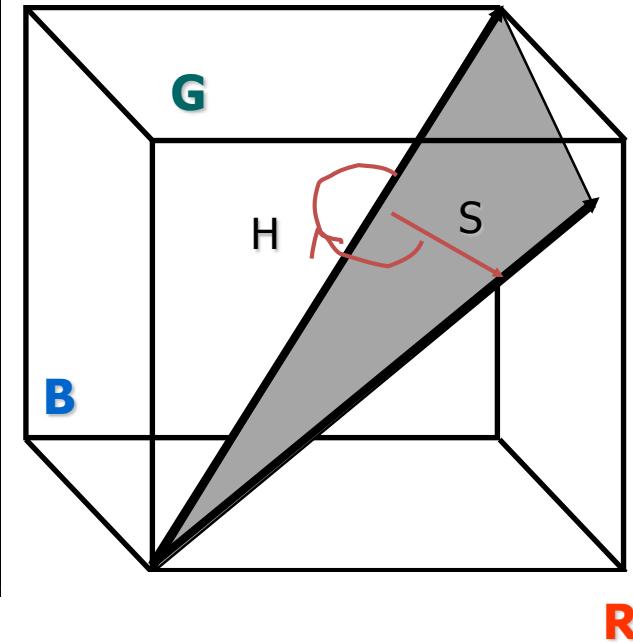
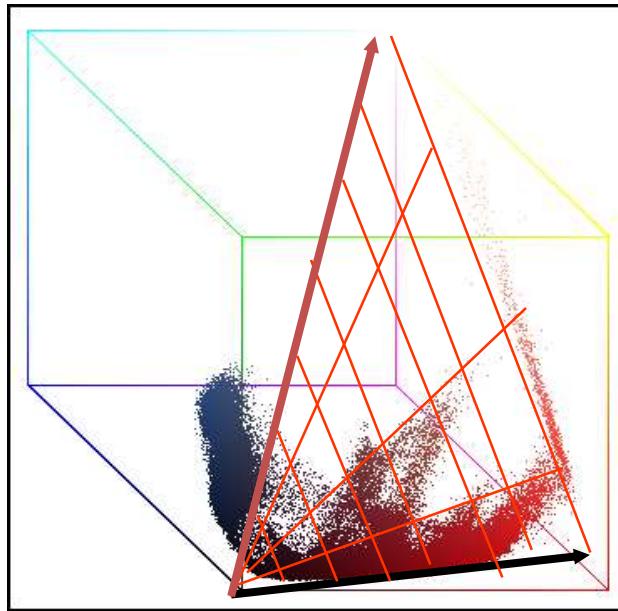
Consider the surface reflection term:

$$m_s(\mathbf{n}, \mathbf{s}, \mathbf{v}) \int_{\lambda} e(\lambda) c(\lambda) f_C(\lambda) d\lambda$$

Where the geometric term is for example the phong model:

$$\cos^n \alpha$$

# Colour Invariance: Hue



$$\text{Hue: } H(R, G, B) = \arctan\left(\frac{\sqrt{3}(G - B)}{(R - G) + (R - B)}\right)$$

# Colour invariance

$$l1(R, G, B) = \frac{(R - G)^2}{(R - G)^2 + (R - B)^2 + (G - B)^2}$$

$$l2(R, G, B) = \frac{(R - B)^2}{(R - G)^2 + (R - B)^2 + (G - B)^2}$$

$$l3(R, G, B) = \frac{(G - B)^2}{(R - G)^2 + (R - B)^2 + (G - B)^2}$$

# Colour Invariance: Hue

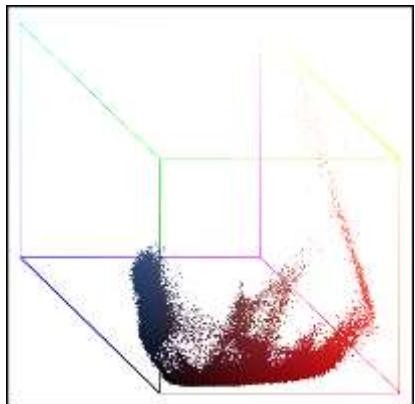


Original colour image

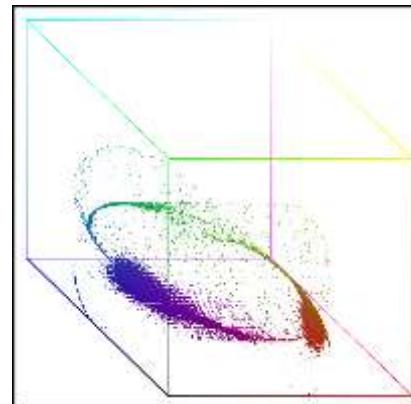
→  
L1I2I3 transformation



L1I2I3 image



3D plot of colour image



3D plot of L1I2I3 image

# Colour Invariance - Summary

	shadows	shading	highlights	ill. intensity	ill. Colour
I	-	-	-	-	-
R,G,B	-	-	-	-	-
r,g,b	+	+	-	+	-
c1,c2,c3	+	+	-	+	-
Hue	+	+	+	+	-
I1,I2,I3	+	+	+	+	-
m1m2m3	+	+	-	+	+

- no invariance

+ invariance

# Summary

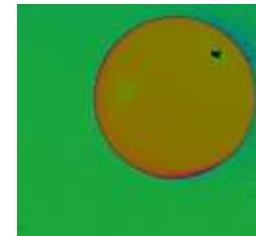
- Reflection models help to understand the image formation process.

$$C = m_b(\mathbf{n}, \mathbf{s}) \int_{\lambda} e(\lambda) \rho(\lambda) f_C(\lambda) d\lambda + m_s(\mathbf{n}, \mathbf{s}, \mathbf{v}) \int_{\lambda} e(\lambda) c(\lambda) f_C(\lambda) d\lambda$$

- Color invariance at the pixel.  
White light source!



- Be aware of instabilities



# Today's Class

- 1. Reflection Models**
- 2. Color Invariance : Constancy**
- 3. Image Processing**

# Color constancy

## Gray-world

### Training

Unknown



### Test

Unknown



### Algorithm

Canonical



Mean color = constant

## Scale-by-max

### Training



### Test

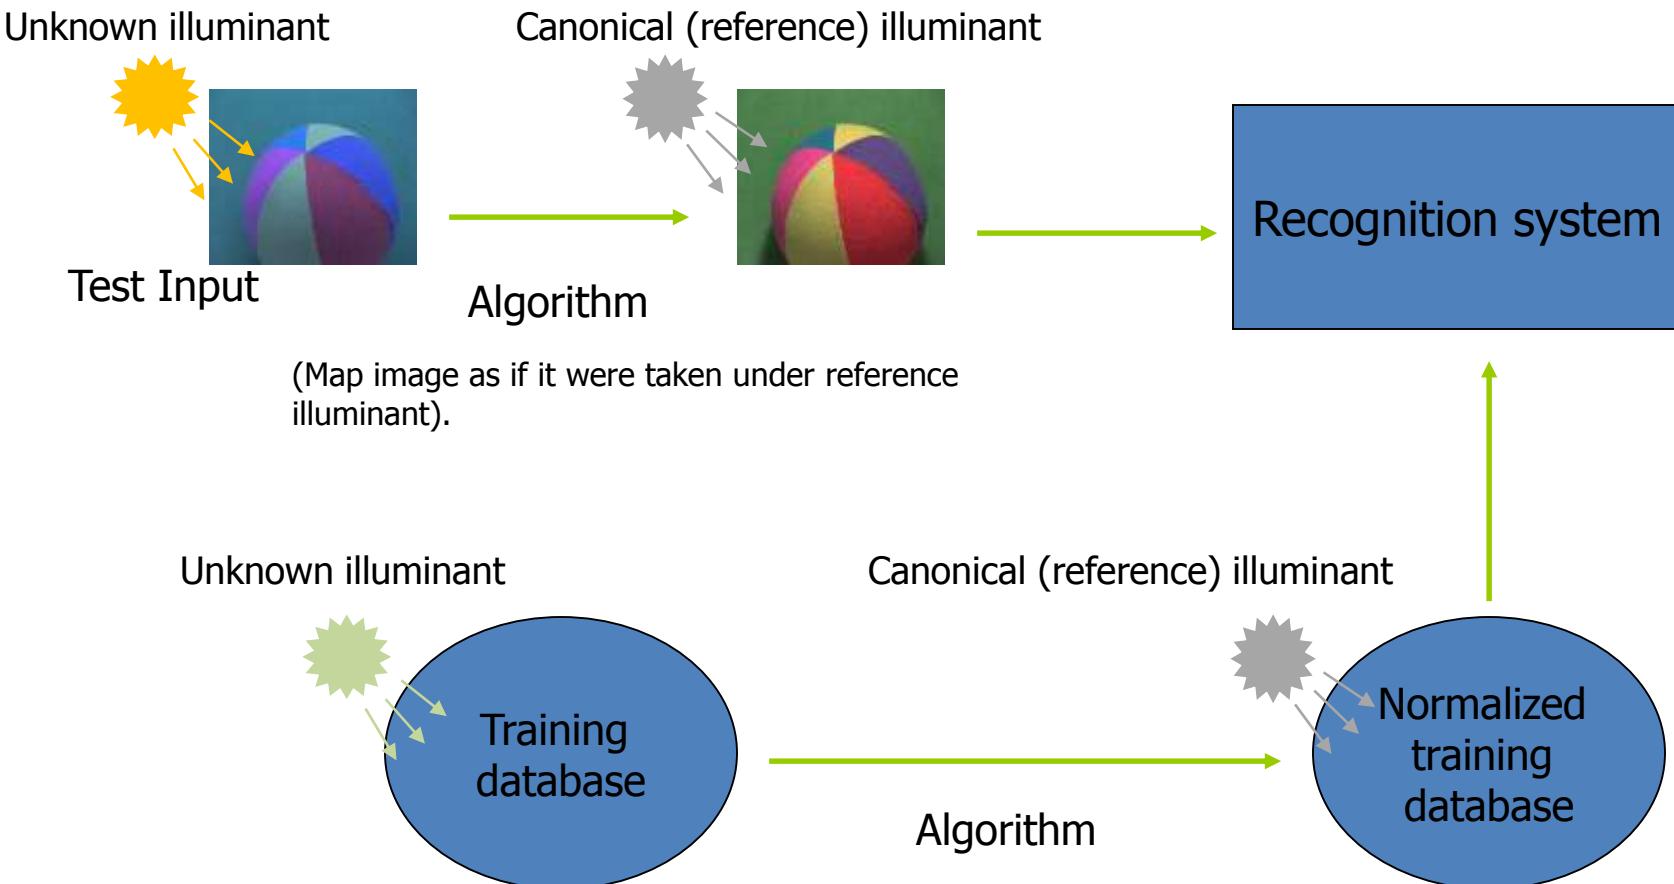


### Algorithm

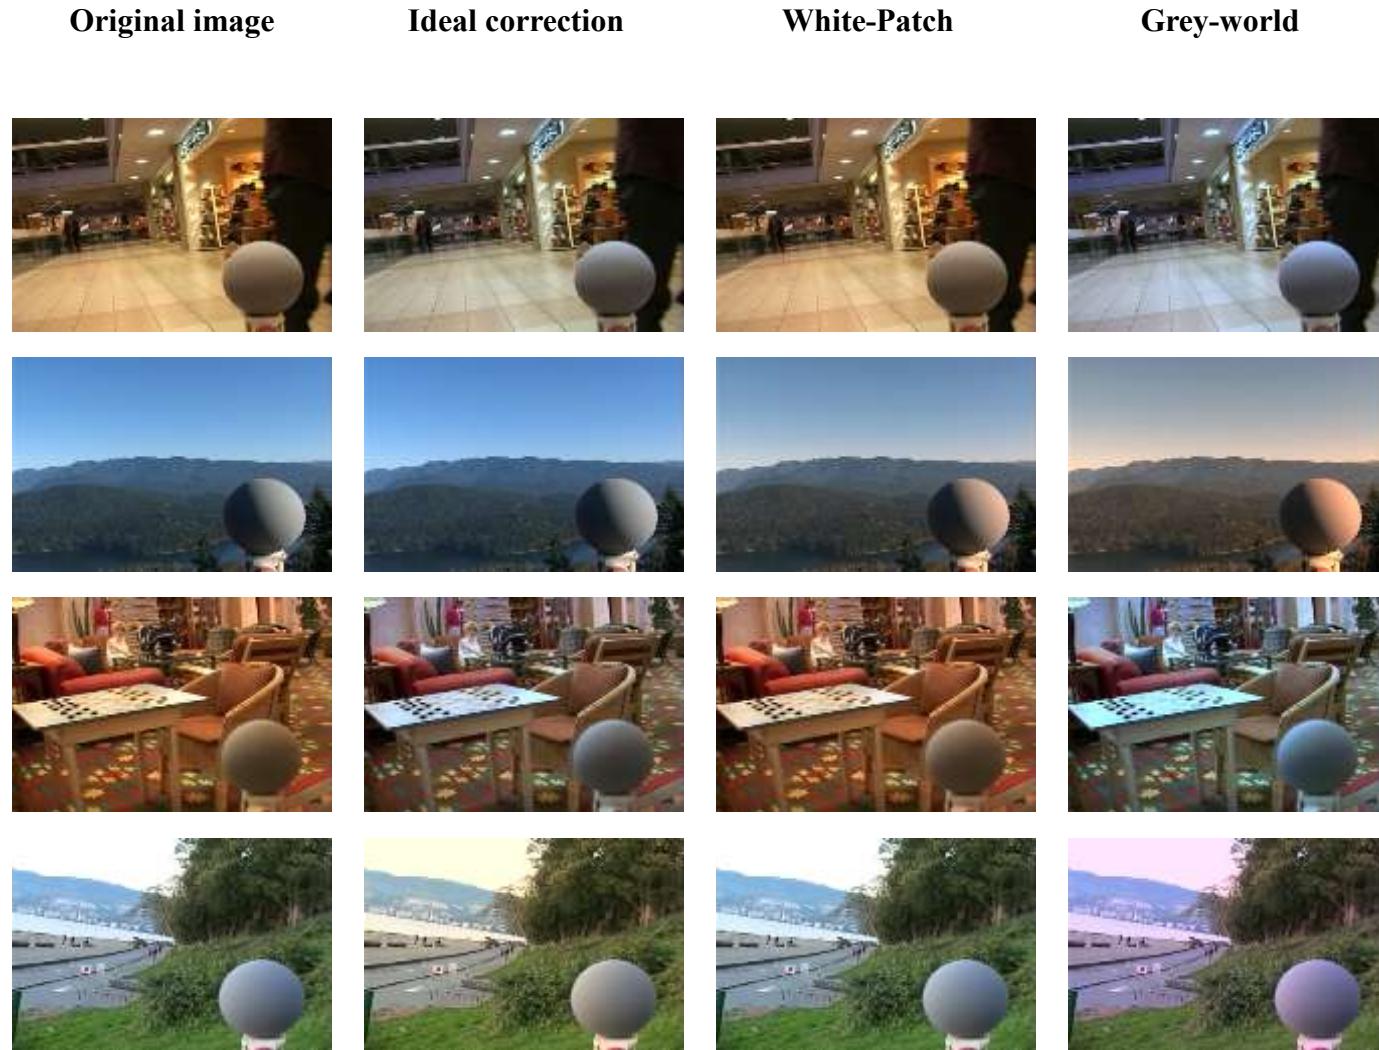


Max color = constant

# Color normalization



# White Patch and Grey-world Examples



Images from “A large image data set for color constancy research”, by F. Ciurea and B. Funt in CIC 2003.

# Color Constancy

- Model images assuming Lambertian reflectance:

$$C = m_b(\mathbf{n}, \mathbf{s}) \int_{\lambda} e(\lambda) \rho(\lambda) f_C(\lambda) d\lambda$$

- Image transformation using von Kries model<sup>(1)</sup>:

$$\begin{pmatrix} R^c \\ G^c \\ B^c \end{pmatrix} = \begin{pmatrix} \alpha & 0 & 0 \\ 0 & \beta & 0 \\ 0 & 0 & \gamma \end{pmatrix} \begin{pmatrix} R^u \\ G^u \\ B^u \end{pmatrix}$$

(1) - J. von Kries. "Influence of adaptation on the effects produced by luminous stimuli." In D. MacAdam, editor, *Sources of Color Vision*, pages 109–119. MIT Press, 1970.

# Color : Von Kries Model

Assuming body reflection

$$C = m_b(\mathbf{n}, \mathbf{s}) \int e(\lambda) \rho(\lambda) f_C(\lambda) d\lambda \quad \text{for } C = \{R, G, B\}$$

And narrow-band filters:  $f_C(\lambda) = \delta(\lambda - \lambda_C)$

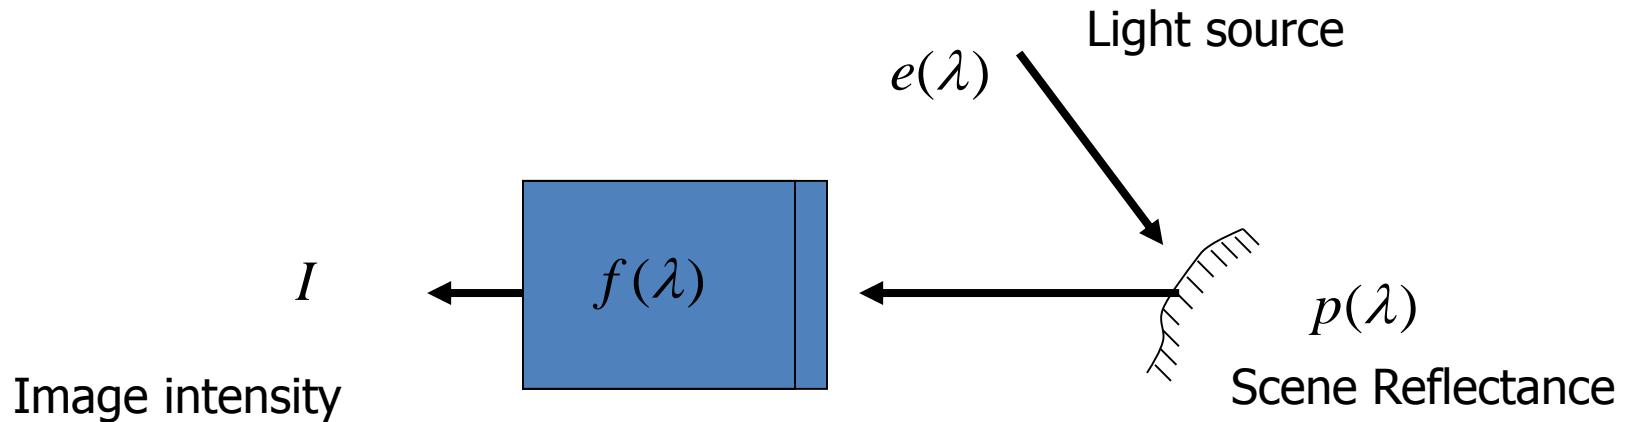
Then:  $C = m_b(\mathbf{n}, \mathbf{s}) e(\lambda_C) \rho(\lambda_C)$  for  $C = \{R, G, B\}$

Therefore, we have:  $R = m_b(\mathbf{n}, \mathbf{s}) e(\lambda_R) \rho(\lambda_R)$

$$G = m_b(\mathbf{n}, \mathbf{s}) e(\lambda_G) \rho(\lambda_G)$$

$$B = m_b(\mathbf{n}, \mathbf{s}) e(\lambda_B) \rho(\lambda_B)$$

# Image Brightness (Intensity)



- Monochromatic Light: ( $\lambda = \lambda_i$ )  $f(\lambda_i) = 1$

$$I = m_b(\mathbf{n}, \mathbf{s}) e \rho$$

$$I = \mathbf{n} \cdot \mathbf{s} e \rho \quad \text{for Lambertian reflection}$$

# Recovering Lightness: Retinex

- Image Intensity:

$$I = e\rho$$

- Take Logarithm:

$$\log I = \log e + \log \rho$$

OR

$$I = e + \rho$$

- Use Laplacian:

$$d = \nabla^2 I = \nabla^2 e + \nabla^2 \rho \quad \nabla^2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$$

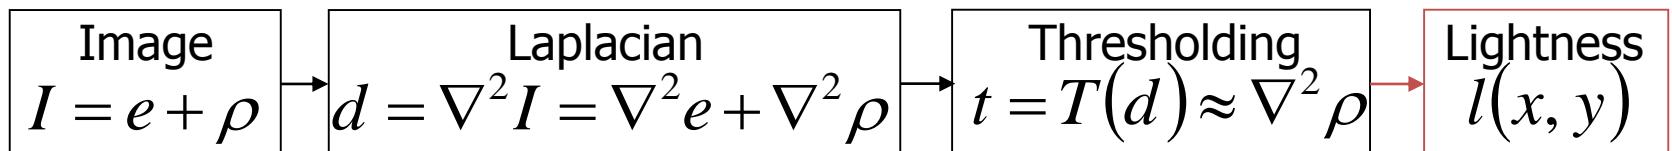
- Sharp changes in reflectance  $\rho$

$\nabla^2 \rho$  has 2 infinite spikes near edges and  $\nabla^2 \rho = 0$  elsewhere

- Smooth changes in illumination  $e$

$\nabla^2 e \approx 0$  everywhere

# Solving the Inverse Problem



Find lightness  $l(x, y)$  from  $t(x, y)$ :

Poisson's Equation

$$\nabla^2 l = t$$

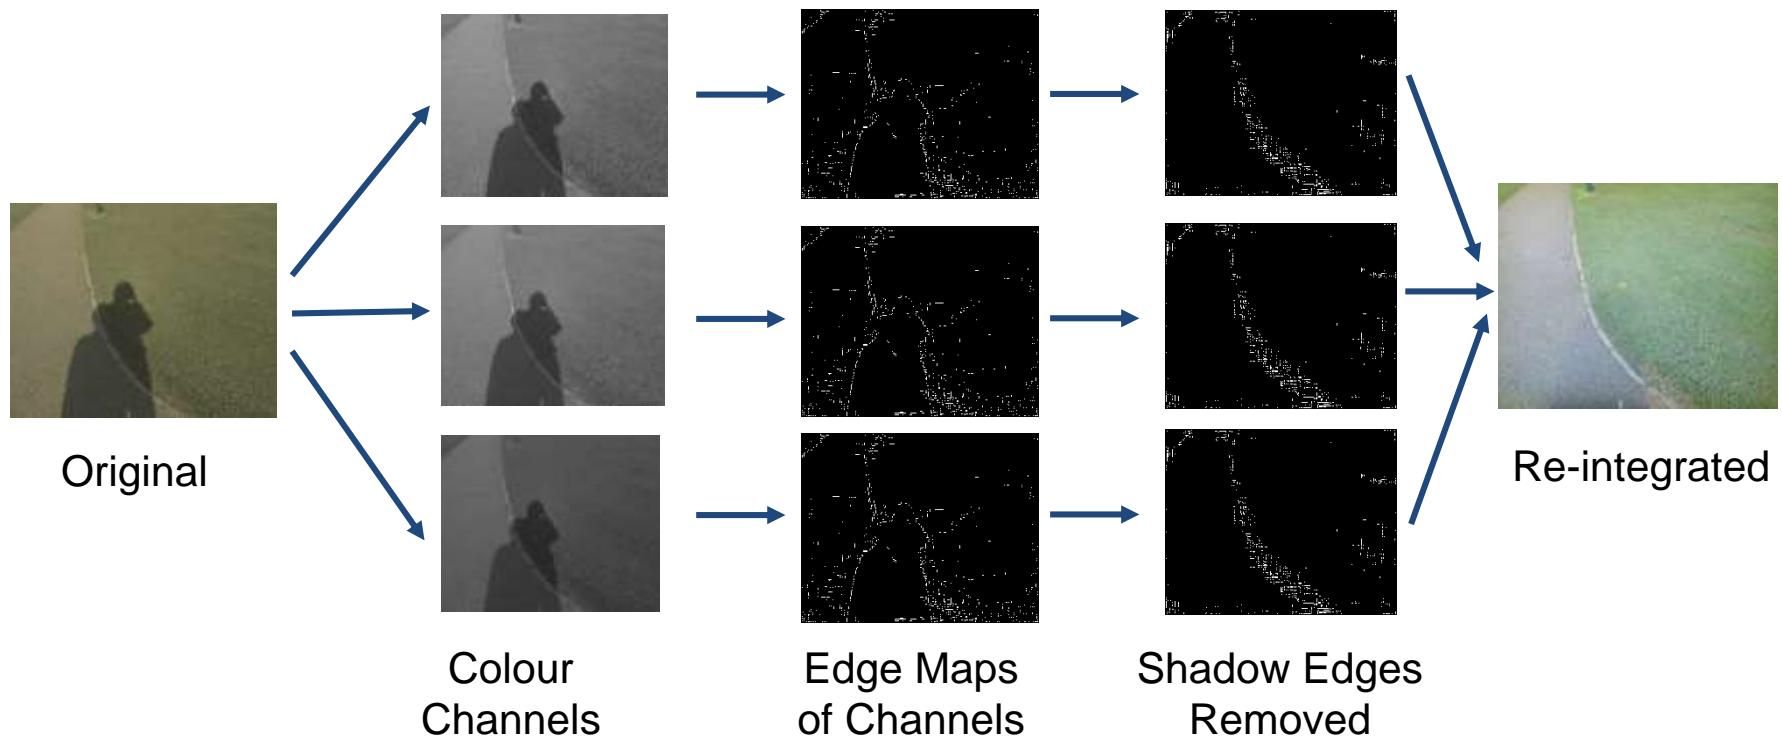
$$\left( \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right) l(x, y) = t(x, y)$$

We have to find  $g(x, y)$  which satisfies

$$l(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} t(u, v) g(x-u, y-v) du dv$$

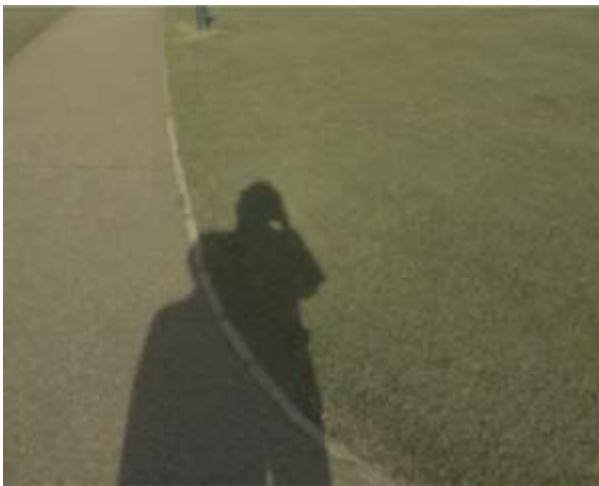
$$l(x, y) = t(x, y) * g(x, y)$$

# Shadow Removal



# An Example

Original  
Image



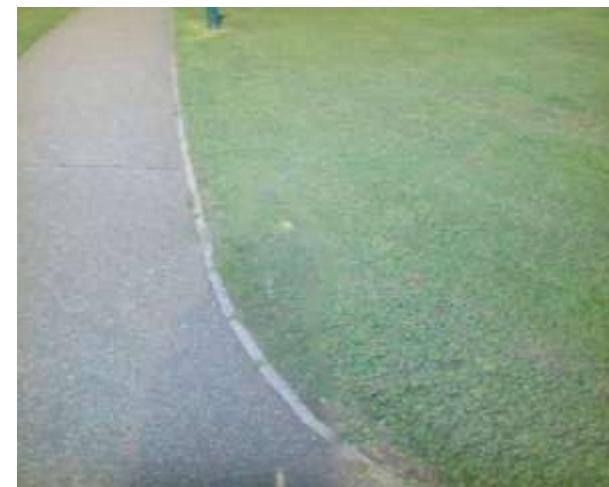
Invariant  
Image



Detected  
Shadow  
Edges



Shadow  
Removed



# A Second Example

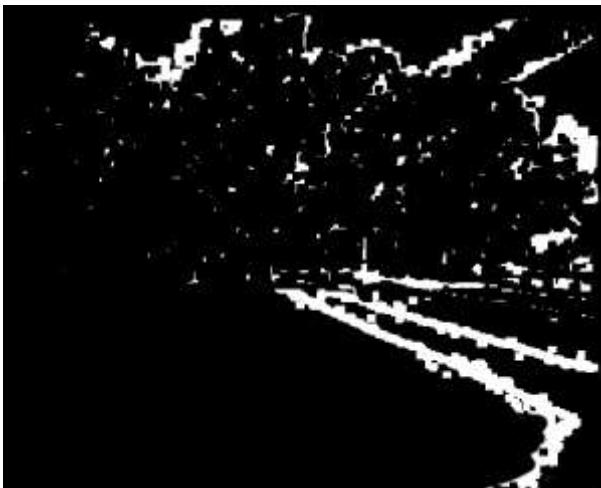
Original  
Image



Invariant  
Image



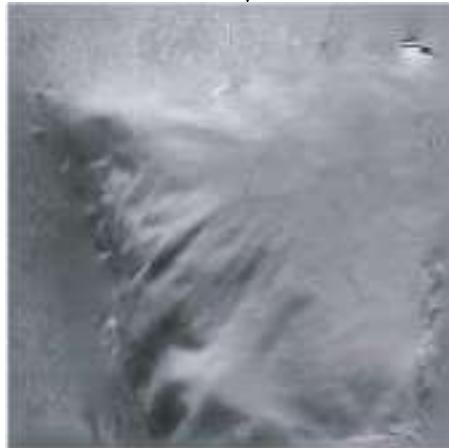
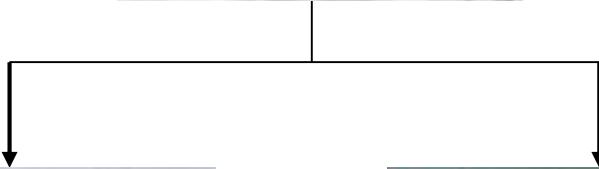
Detected  
Shadow  
Edges



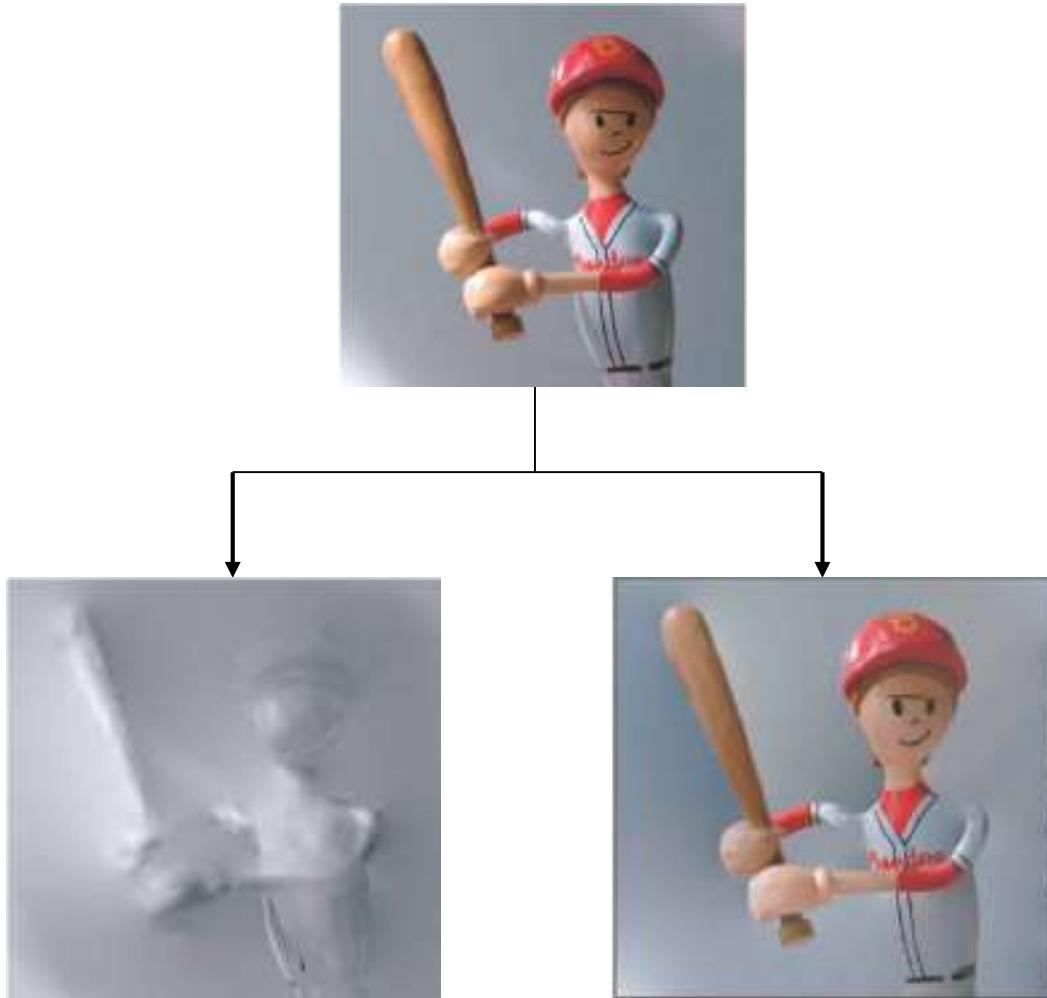
Shadow  
Removed



# Intrinsic Images



# Intrinsic Images



# Summary

- Estimation of the light source

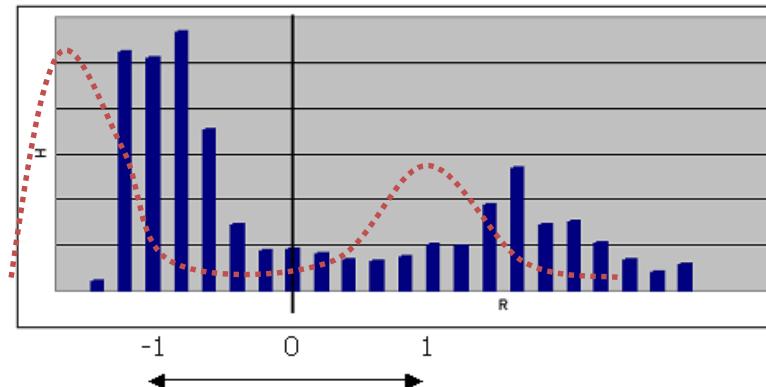


- Invariant color features

- Be aware of instabilities

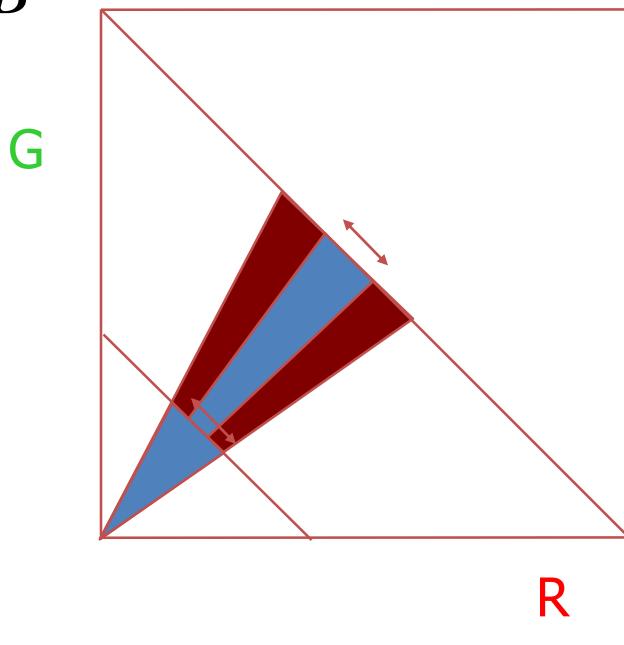


# Kernel Density Estimation for Object Recognition



# Color Invariants: Instabilities

$$r = \frac{R}{R + G + B}$$



# Color Invariants: Instabilities

Measuring two quantities  $x$  and  $y$  to calculate their sum  $x + y$

Then the highest probable value of  $x + y = x_{\text{best}} + y_{\text{best}} + (\delta x + \delta y)$

Then the lowest probable value of  $x + y = x_{\text{best}} + y_{\text{best}} - (\delta x + \delta y)$

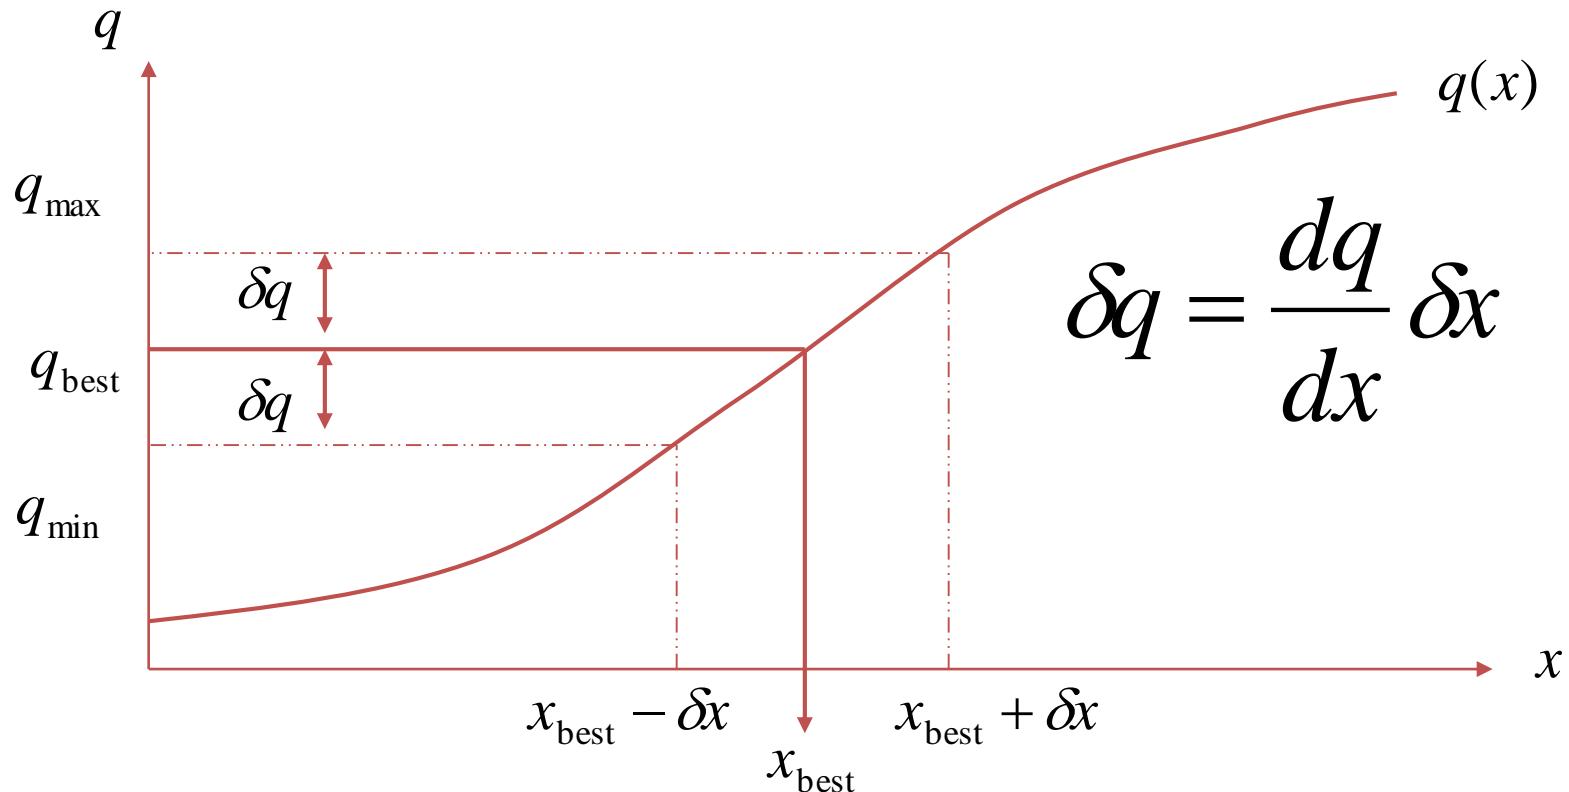
Where  $x_{\text{best}}$  is the best estimate of  $x$  and  $\delta x$  the uncertainty of  $x$

$$q = x + \dots + z - (u + \dots + w)$$

$$\delta q \approx \delta x + \dots + \delta z + \delta u + \dots + \delta w$$

# Color Invariants: Noise Propagation

$$\delta q = q(x_{\text{best}} + \delta x) - q(x_{\text{best}})$$



# Color Invariants: Noise Propagation

Suppose that  $u, \dots, w$  are measured with corresponding uncertainties  $\sigma_u, \dots, \sigma_w$  to compute function  $q(u, \dots, w)$ .

The predicted uncertainty is defined by :

$$\sigma_q = \sqrt{\left(\frac{\partial q}{\partial u} \sigma_u\right)^2 + \dots + \left(\frac{\partial q}{\partial w} \sigma_w\right)^2}$$

The uncertainty in  $q$  is never larger than the ordinary sum

$$\sigma_q \leq \left| \frac{\partial q}{\partial u} \right| \sigma_u + \dots + \left| \frac{\partial q}{\partial w} \right| \sigma_w$$

if and only if the uncertainties  $\sigma_u, \dots, \sigma_w$  are relatively small.

# Noise Propagation: rgb

Function  $q(x, \dots, z)$  then  $\delta_q = \sqrt{\left(\frac{\partial q}{\partial x} \delta x\right)^2 + \dots + \left(\frac{\partial q}{\partial z} \delta z\right)^2}$

$$r(R, G, B) = \frac{R}{R + G + B}$$

Normalized color value

$$\delta_r = \frac{\sqrt{R^2(\delta_B^2 + \delta_G^2) + (B + G)\delta_R^2}}{(B + G + R)^4}$$

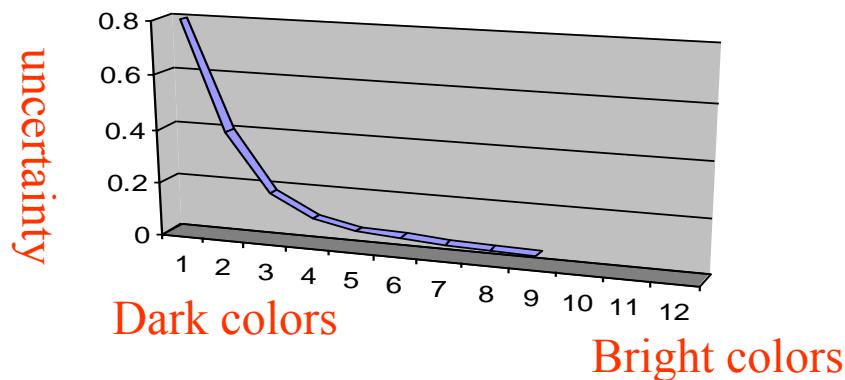
Predicted uncertainty of normalized color value

Measured values and predicted uncertainty

# Noise Propagation: Hue

$$H = \arctan\left(\frac{\sqrt{3}(G-B)}{(R-G)+(R-B)}\right)$$

$$\sigma_H = \frac{1}{2} \sqrt{3} \left( \frac{-2BR\sigma_G^2 + R^2\sigma_B^2 + \sigma_G^2 + G^2(\sigma_B^2 + \sigma_R^2) + B^2(\sigma_G^2 + \sigma_R^2) - 2G(R\sigma_B^2 + B\sigma_R^2)}{B^2 + G^2 - GR + R^2 - B(G + R)^2} \right)$$



# Histogram Construction

$$\tilde{f}(x) = \frac{1}{nh} (\text{number of } X_i \text{ in the same bin as } x)$$

$n$  is the number of pixels  $X_i$  in the image,  $h$  is the bin width and  $x$  the range of the data

Kernel density estimator:

$$\tilde{f}(x) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - X_i}{h}\right)$$

where  $K$  is a function satisfying  $\int K(x)dx = 1$

# Variable Kernel Density Estimation

The most frequently occurring noise is additive Gaussian noise.

It is widely used to model thermal noise and is the limiting behaviour of photon counting noise and film grain noise

$$K(x) = \frac{1}{\sqrt{2\pi}} \exp^{\frac{-x^2}{2}}$$

Variable kernel density for  $rg$  is defined by :

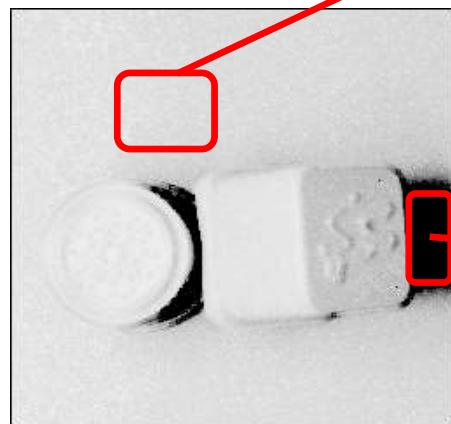
$$\tilde{f}(x) = \frac{1}{n} \sum_{i=1}^n \sigma_r^{-1} K\left(\frac{x-r}{\sigma_r}\right) \sigma_g^{-1} K\left(\frac{x-g}{\sigma_g}\right)$$

where  $\sigma_r, \sigma_g$  denote the uncertainties in normalized color

# Variable Kernel Density Estimation



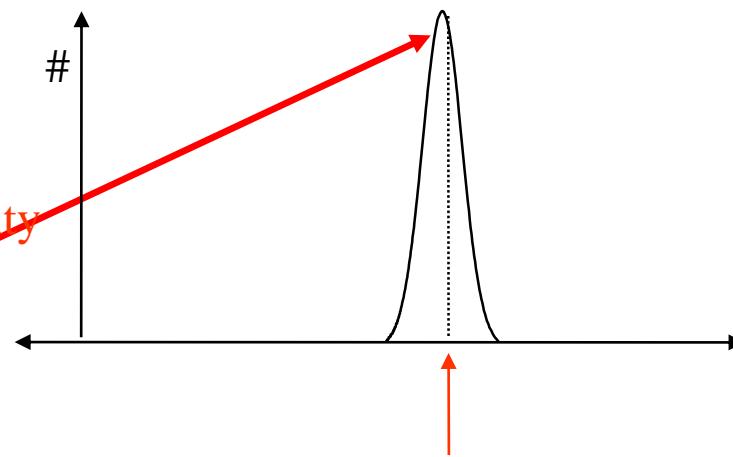
image



uncertainty

High certainty

Low certainty



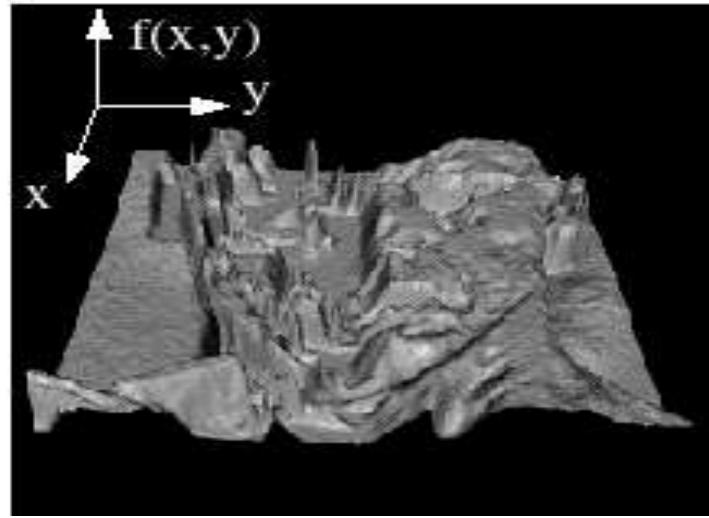
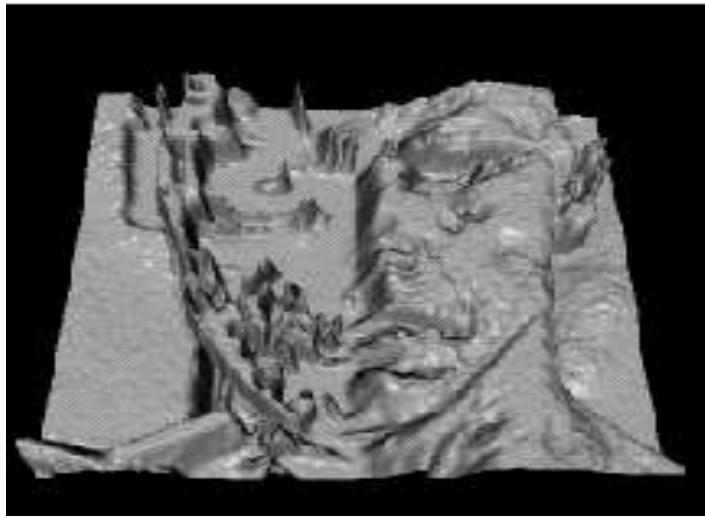
color invariant value

color invariant value

# Today's Class

- 1. Reflection Models**
- 2. Color Invariance**
- 3. Image Processing**

# Images as Functions



# Image Processing

An **image processing** operation typically defines a new image  $g$  in terms of an existing image  $f$ .

We can transform either the range of  $f$ .

$$g(x, y) = t(f(x, y))$$

Or the domain of  $f$ :

$$g(x, y) = f(t_x(x, y), t_y(x, y))$$

What kinds of operations can each perform?

# Point Processing

The simplest kind of range transformations are these independent of position x,y:

$$g = t(f)$$

This is called point processing.

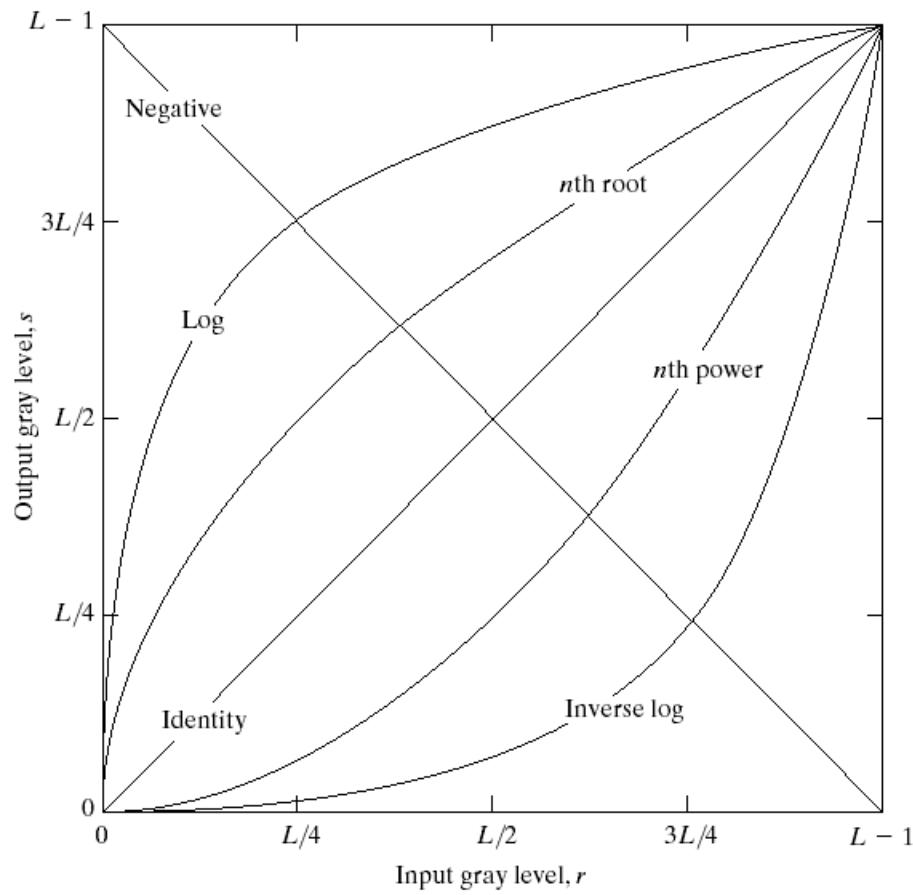
What can they do?

e.g.  $g = at(f) + b$  where  $a$  = gain (contrast) and  $b$  = bias (brightness)

**Important:** every pixel for himself – spatial information completely lost!

# Basic Point Processing

**FIGURE 3.3** Some basic gray-level transformation functions used for image enhancement.

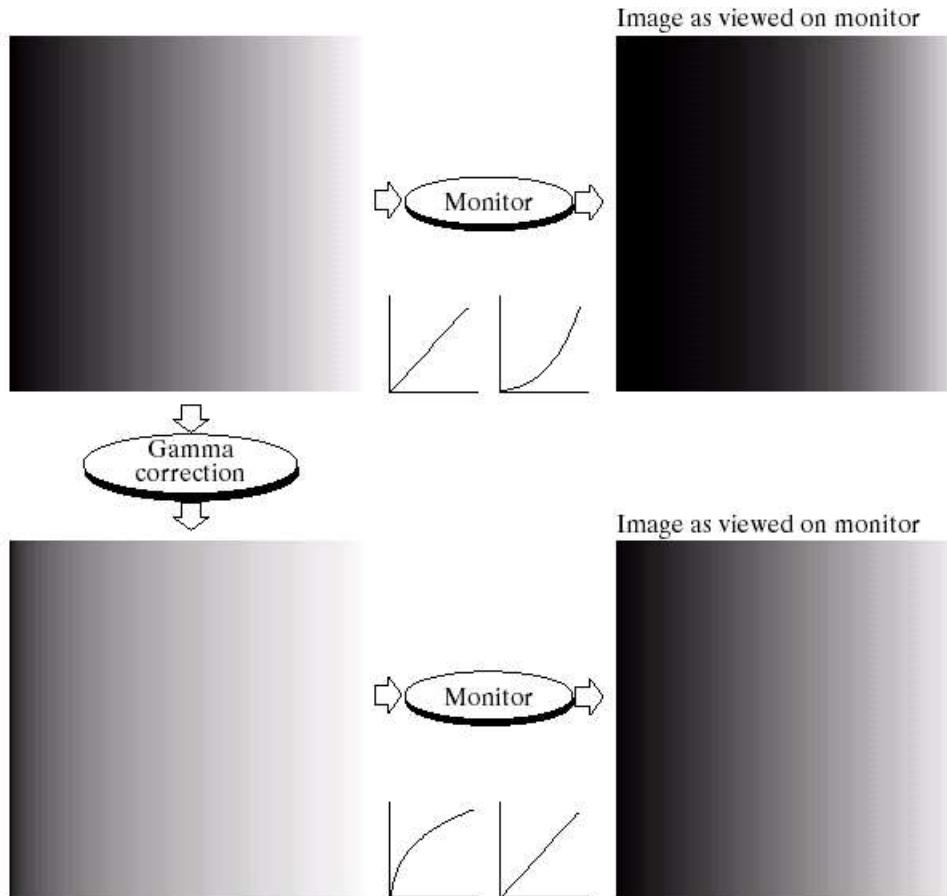


# Gamma Correction

a b  
c d

**FIGURE 3.7**

- (a) Linear-wedge gray-scale image.
- (b) Response of monitor to linear wedge.
- (c) Gamma-corrected wedge.
- (d) Output of monitor.



Gamma Measuring Applet:

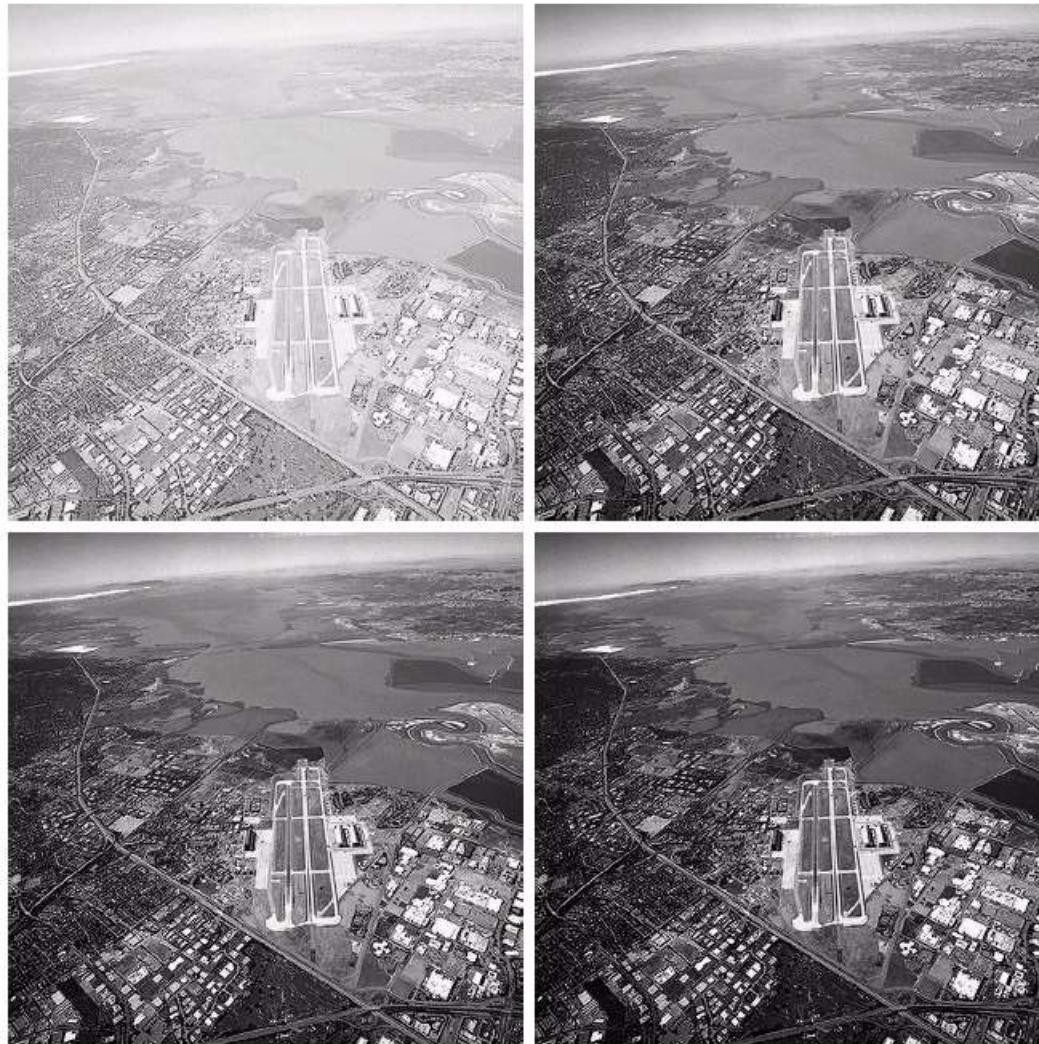
<http://www.cs.berkeley.edu/~efros/java/gamma/gamma.html>

# Image Enhancement

a b  
c d

**FIGURE 3.9**

(a) Aerial image.  
(b)–(d) Results of applying the transformation in Eq. (3.2-3) with  $c = 1$  and  $\gamma = 3.0, 4.0$ , and  $5.0$ , respectively.  
(Original image for this example courtesy of NASA.)



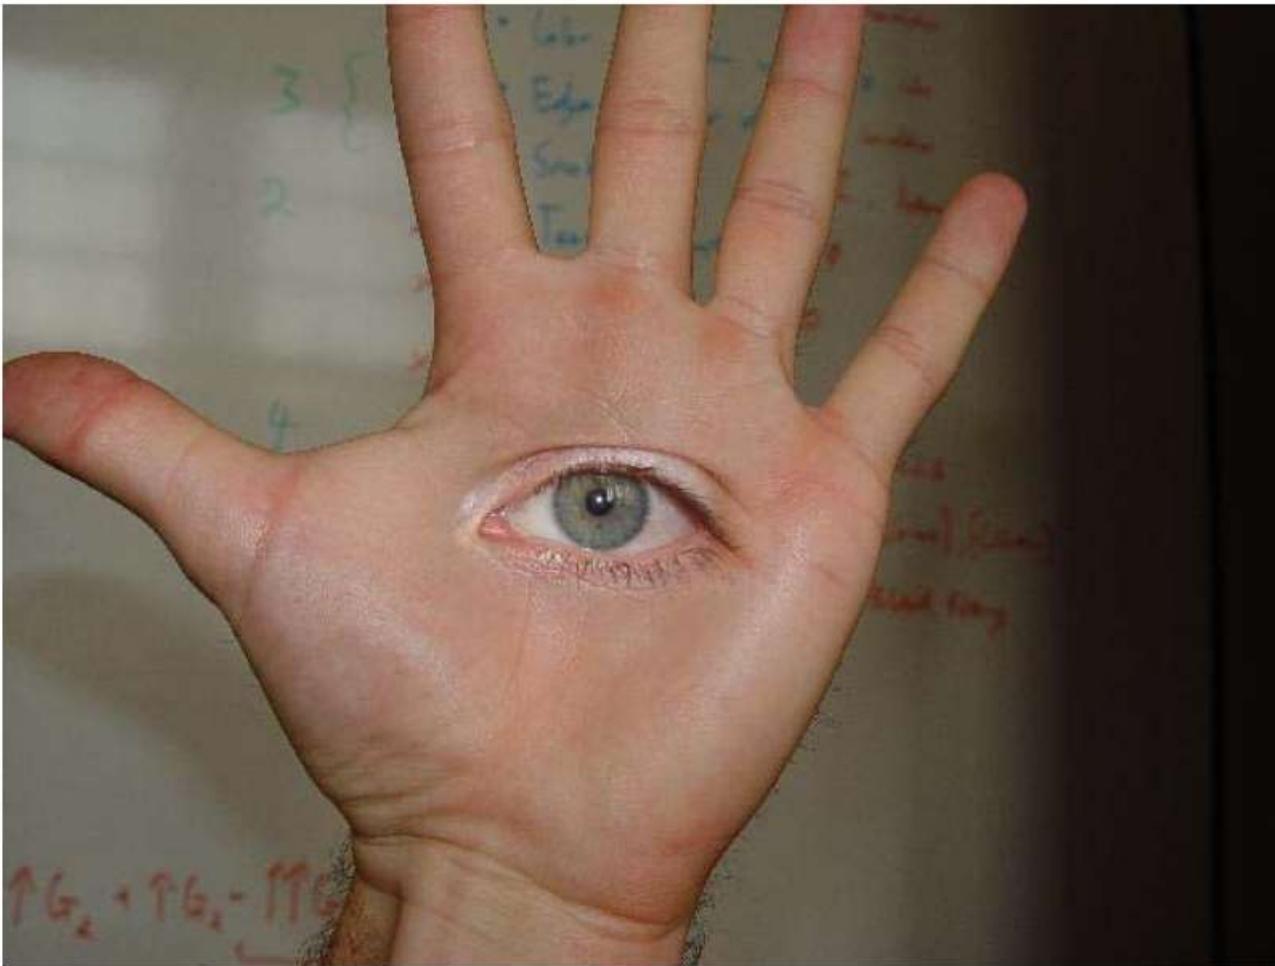
# Photo Editing: Matting



# Photo Editing: Matting



# Poisson Matting



© prof. d'martin

# Neighborhood Processing (filtering)

Q: What happens if I reshuffle all pixels within the image?



A: Its histogram won't change. No point processing will be affected...

Need spatial information to capture this.

# Linear Filtering

$g[\cdot, \cdot]$

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Slide credit: David Lowe (UBC)

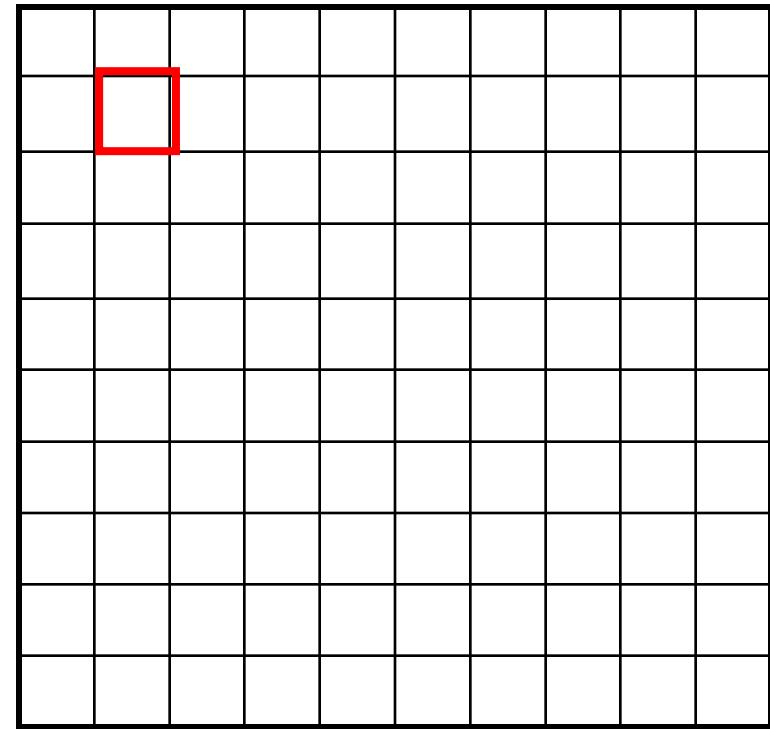
# Image Filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

 $f[.,.]$  $h[.,.]$ 

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0



$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$

# Image Filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$

0	10									

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

# Image Filtering

$$g[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	0	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

			0	10	20				

$$h[m, n] = \sum_{k, l} g[k, l] f[m+k, n+l]$$

# Image Filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$


0    10    20    30

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

# Image Filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

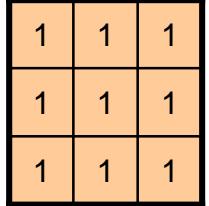
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$


0    10    20    30    30

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

# Image Filtering

$$g[\cdot, \cdot] \frac{1}{9}$$


$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

	0	10	20	30	30				

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

# Image Filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

 $f[.,.]$  $h[.,.]$ 

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30					

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

# Image Filtering

$$g[\cdot, \cdot] \begin{matrix} 1 \\ 1 \\ 9 \end{matrix} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \\ \hline \end{array}$$

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

# Box Filter

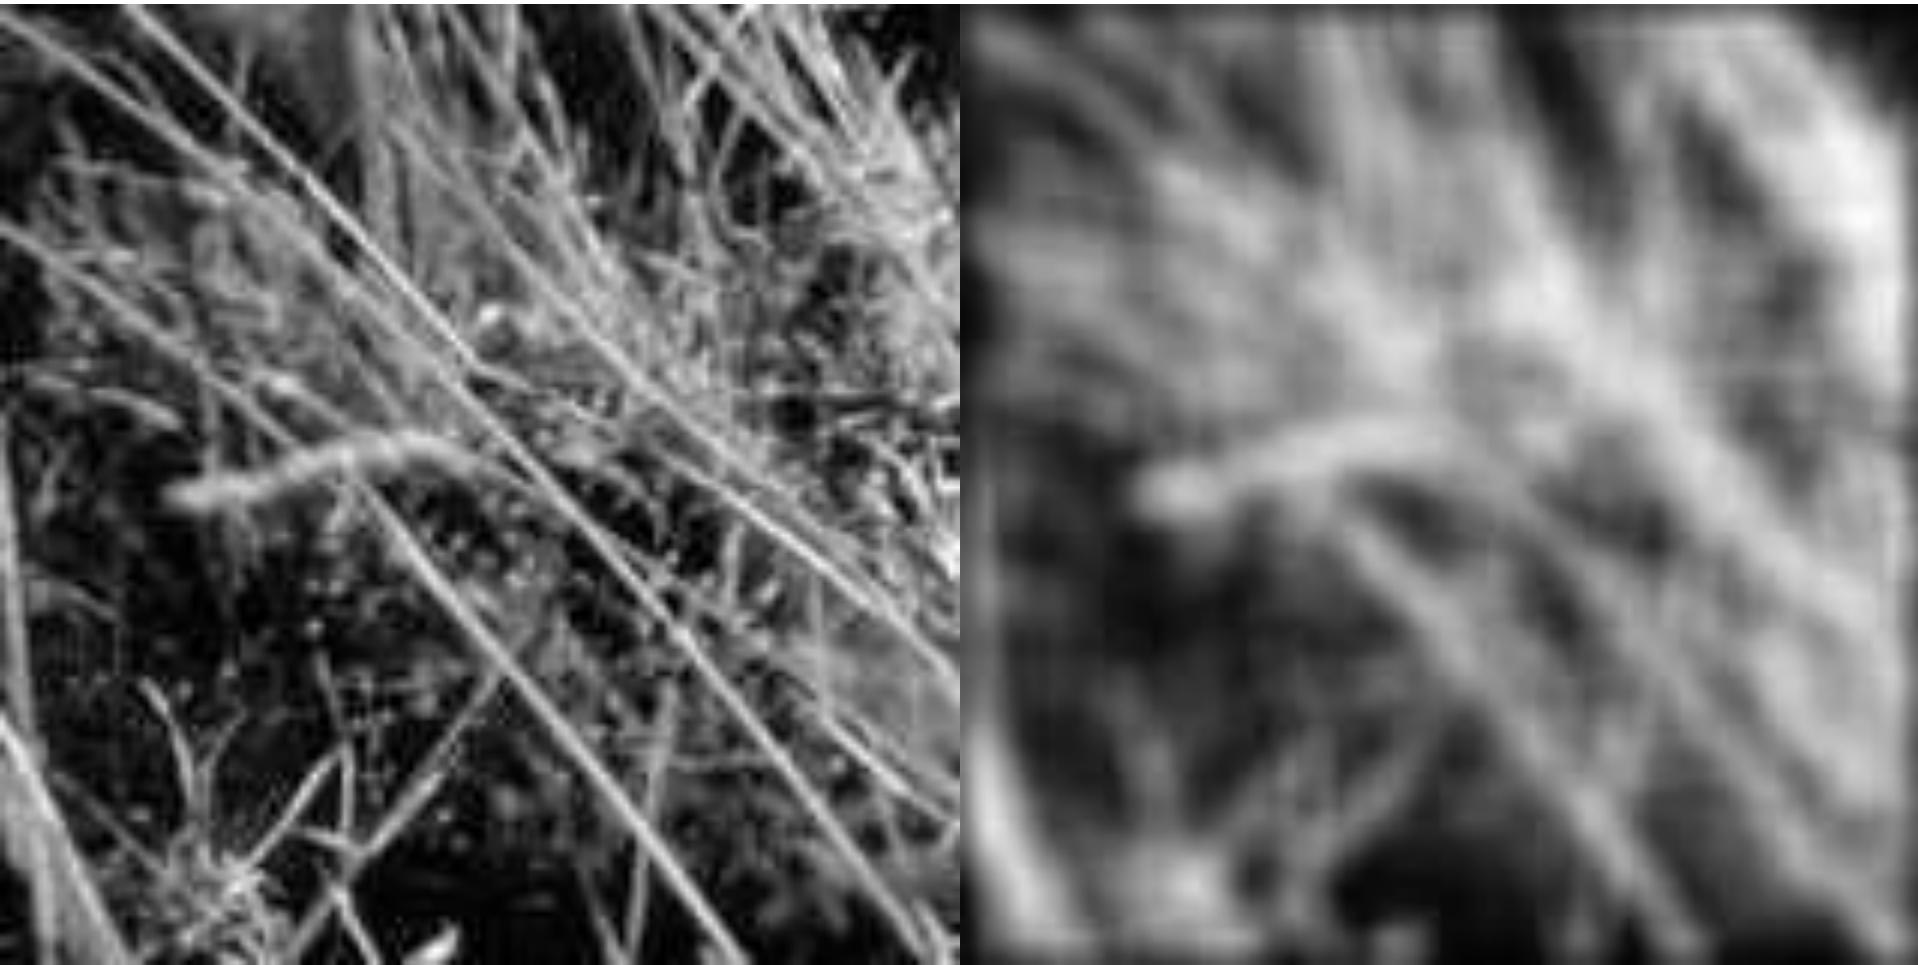
What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

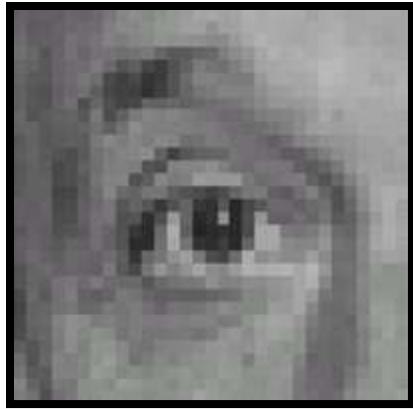
$$g[\cdot, \cdot] = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Slide credit: David Lowe (UBC)

# Smoothing with Box Filter



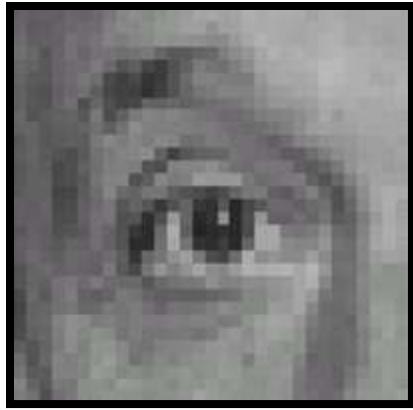
# Practice with Linear Filters



0	0	0
0	1	0
0	0	0

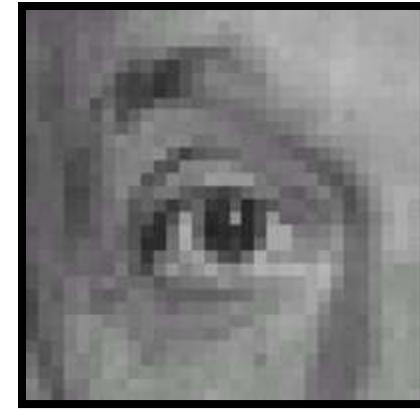
?

# Practice with Linear Filters



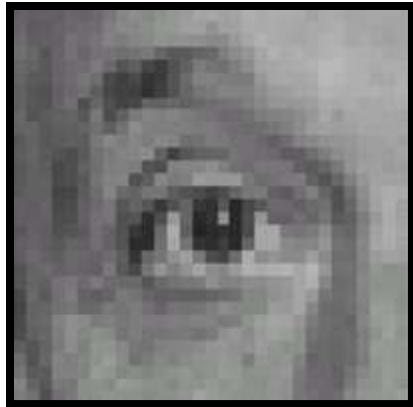
Original

0	0	0
0	1	0
0	0	0



Filtered  
(no change)

# Practice with Linear Filters

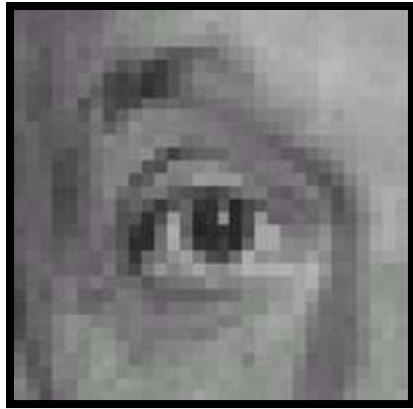


Original

0	0	0
0	0	1
0	0	0

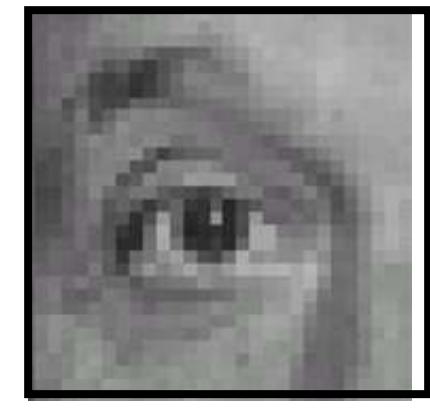
?

# Practice with Linear Filters



Original

0	0	0
0	0	1
0	0	0



Shifted left  
By 1 pixel

# Practice with Linear Filters



0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1

?

# Practice with Linear Filters

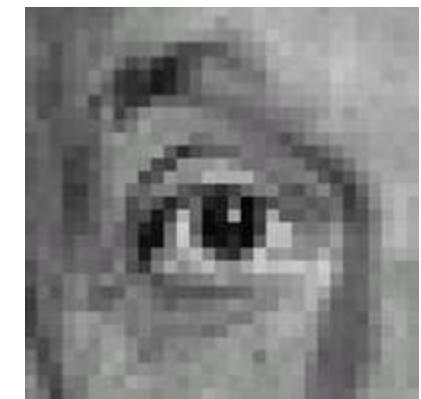


Original

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array}$$

-

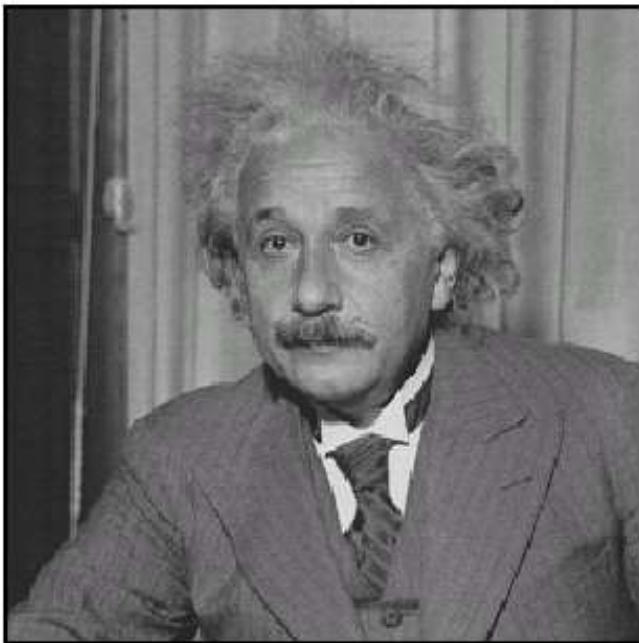
$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



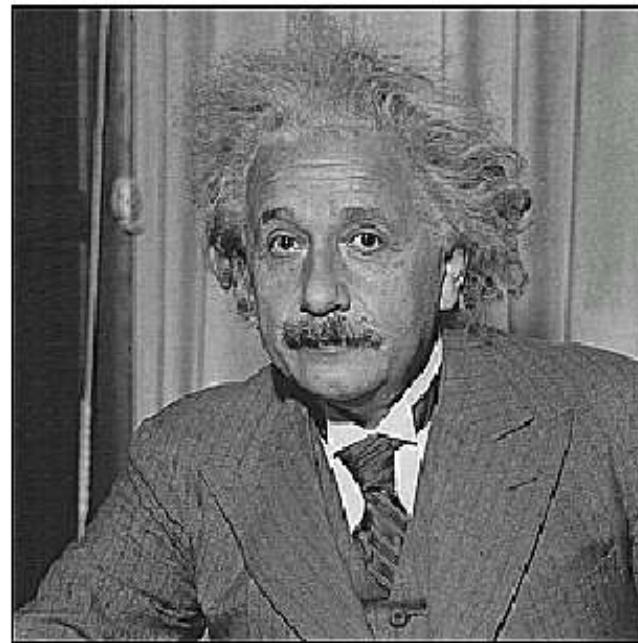
## Sharpening filter

- Accentuates differences with local average

# Sharpening

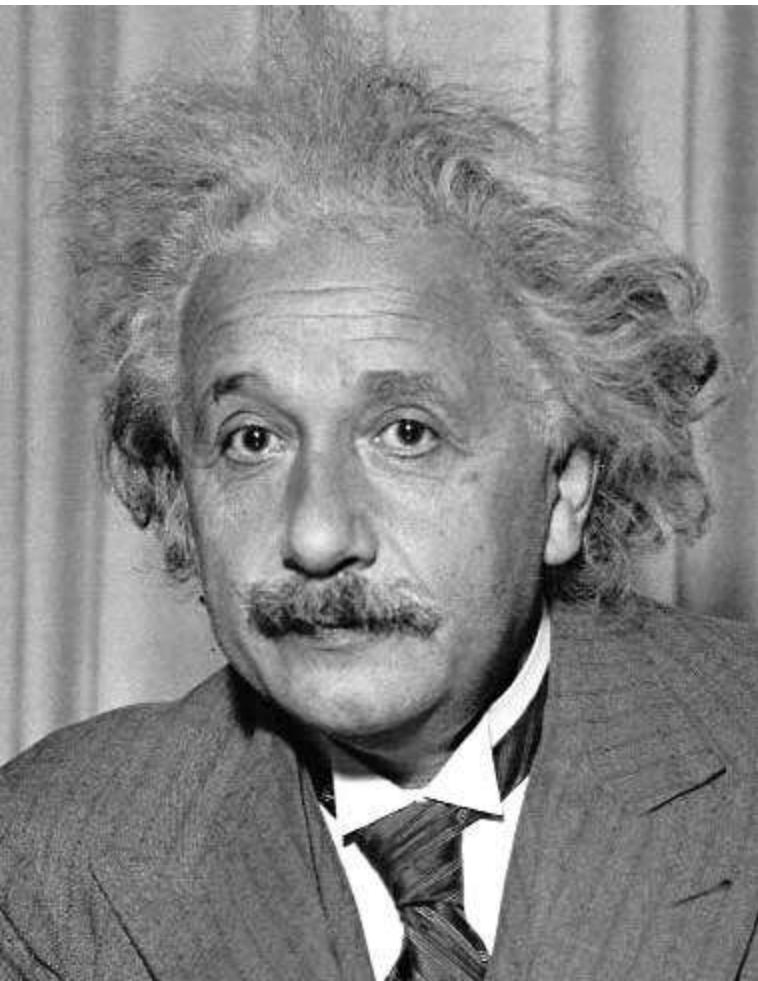


**before**



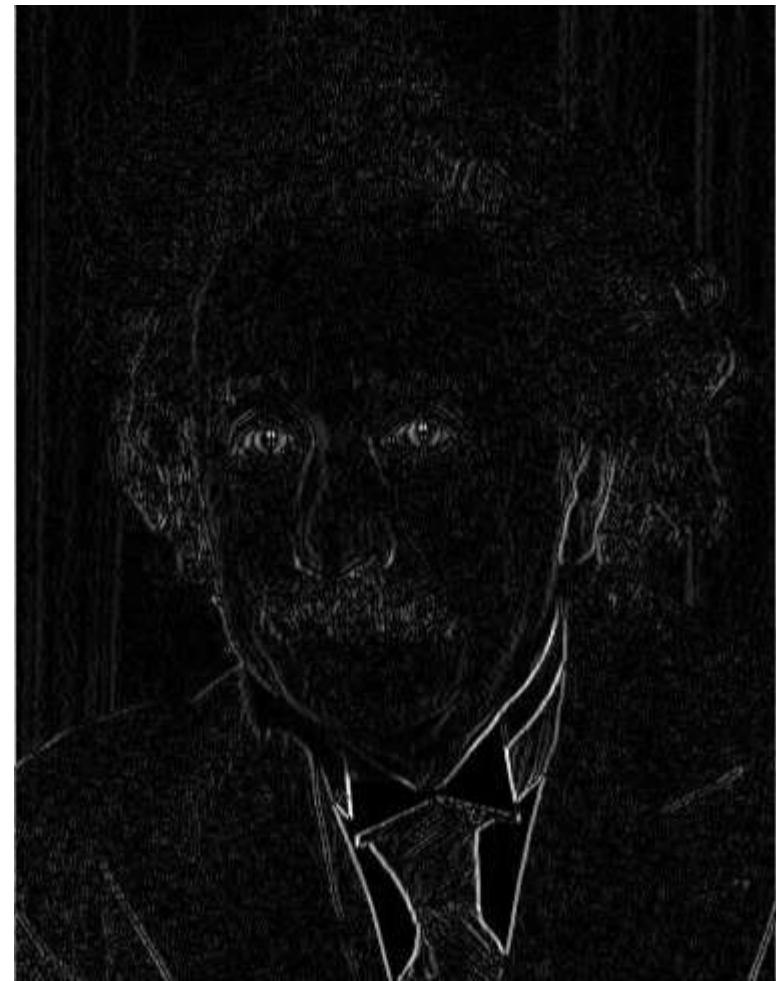
**after**

# Other Filters



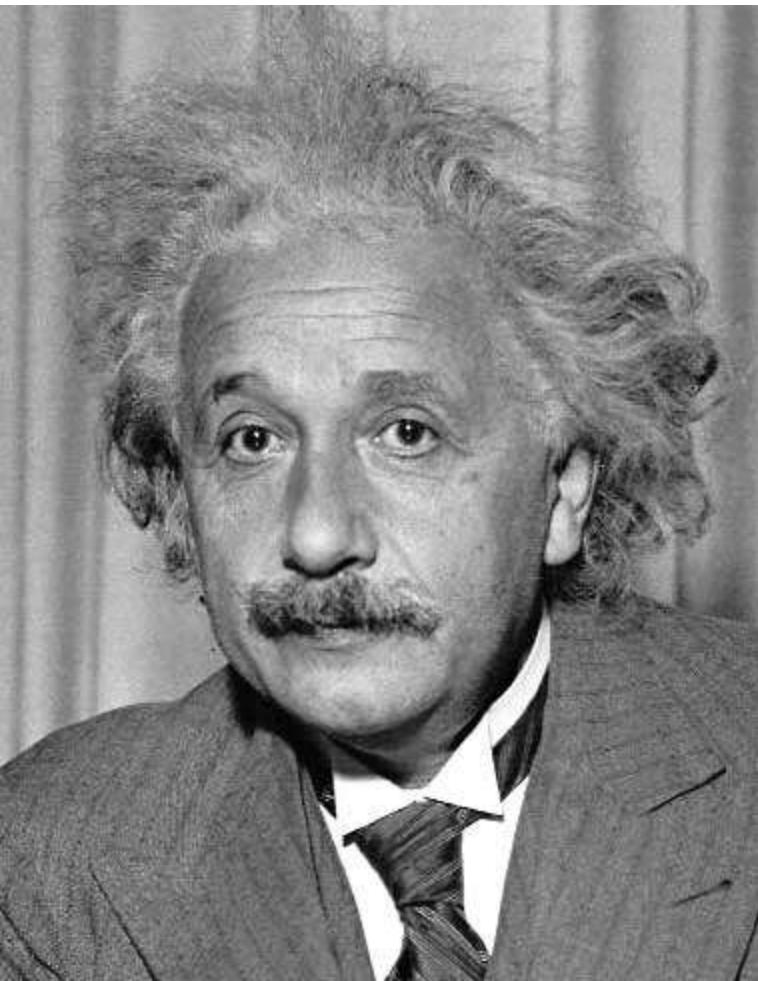
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge  
(absolute value)<sup>100</sup>

# Other Filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge  
(absolute value)<sup>101</sup>

# Filtering vs. Convolution

- 2d filtering
  - $h = \text{filter2}(g, f);$  or  $h = \text{imfilter}(f, g);$

$g = \text{filter}$      $f = \text{image}$

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

- 2d convolution
  - $h = \text{conv2}(g, f);$

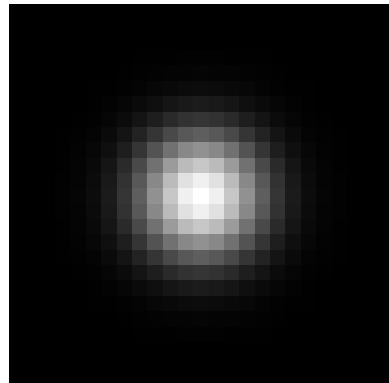
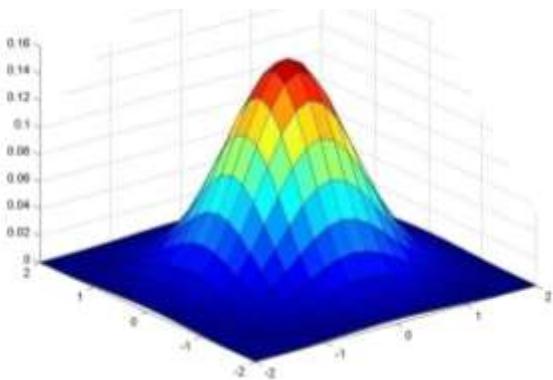
$$h[m, n] = \sum_{k, l} g[k, l] f[m - k, n - l]$$

# Properties

- Commutative:  $a * b = b * a$ 
  - Conceptually no difference between filter and signal
- Associative:  $a * (b * c) = (a * b) * c$ 
  - Often apply several filters one after another:  $((a * b_1) * b_2) * b_3$
  - This is equivalent to applying one filter:  $a * (b_1 * b_2 * b_3)$
- Distributes over addition:  $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out:  $ka * b = a * kb = k(a * b)$
- Identity: unit impulse  $e = [0, 0, 1, 0, 0]$ ,  
 $a * e = a$

# Important Filter: Gaussian

Weight contributions of neighboring pixels by nearness



0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

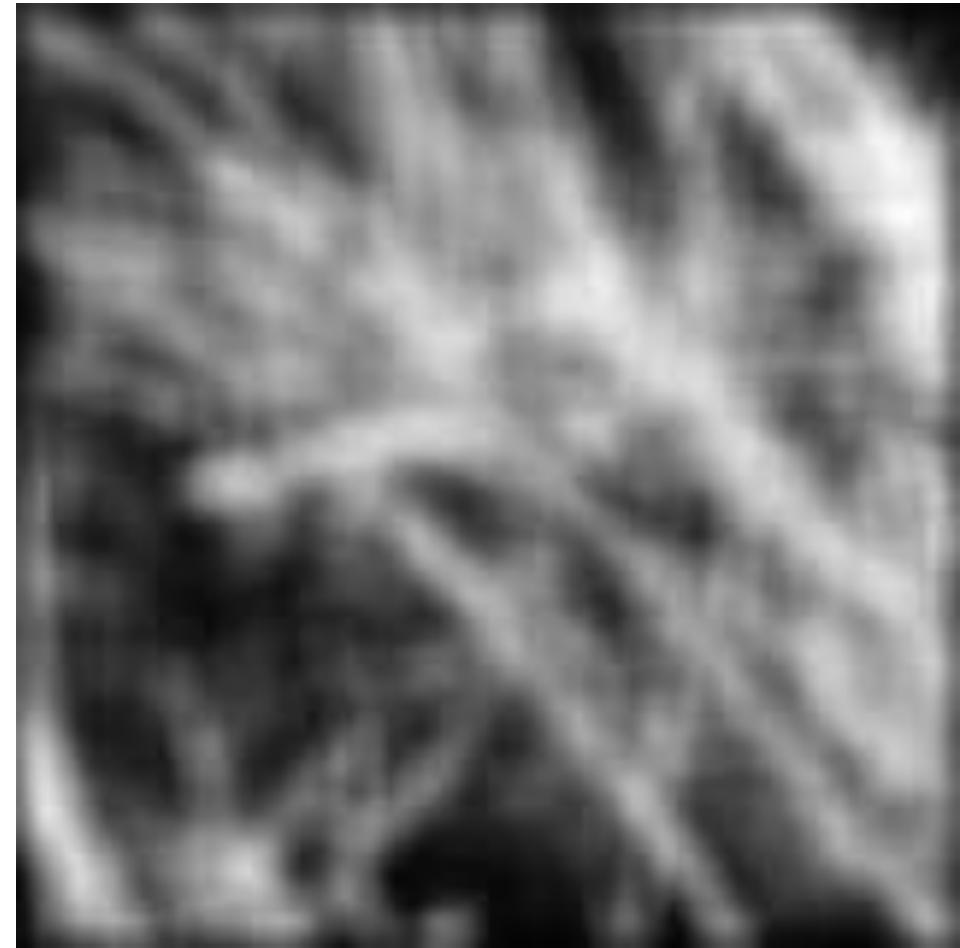
$5 \times 5, \sigma = 1$

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

# Smoothing with Gaussian filter



# Smoothing with Box Filter



# Gaussian Filters

- Remove “high-frequency” components from the image (low-pass filter)
  - Images become more smooth
- Convolution with self is another Gaussian
  - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
  - Convolving two times with Gaussian kernel of width  $\sigma$  is same as convolving once with kernel of width  $\sigma\sqrt{2}$
- *Separable* kernel
  - Factors into product of two 1D Gaussians

# Separability of the Gaussian Filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

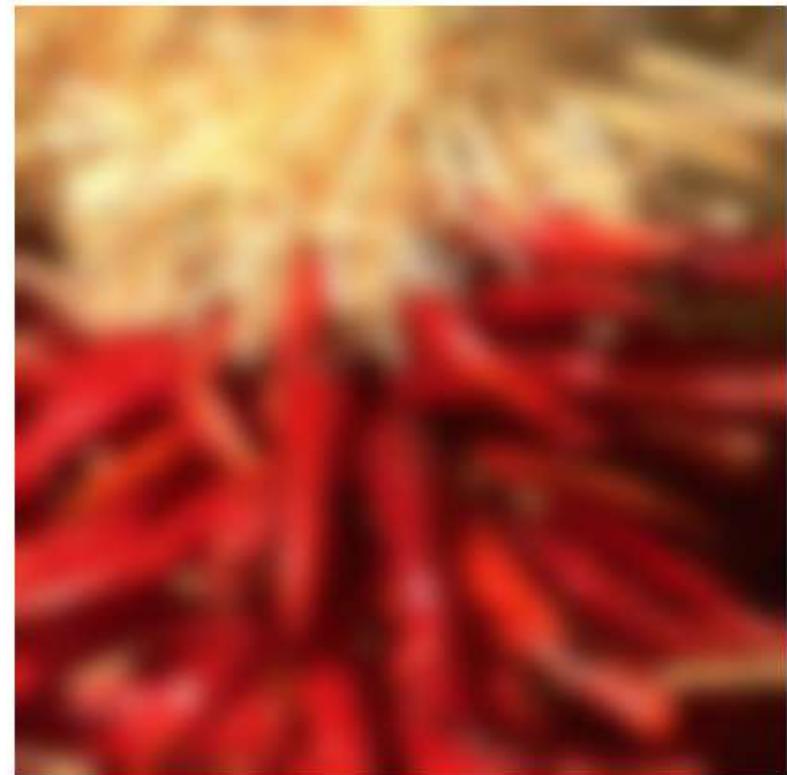
# Separability Example

The filter factors  
into a product of 1D  
filters:

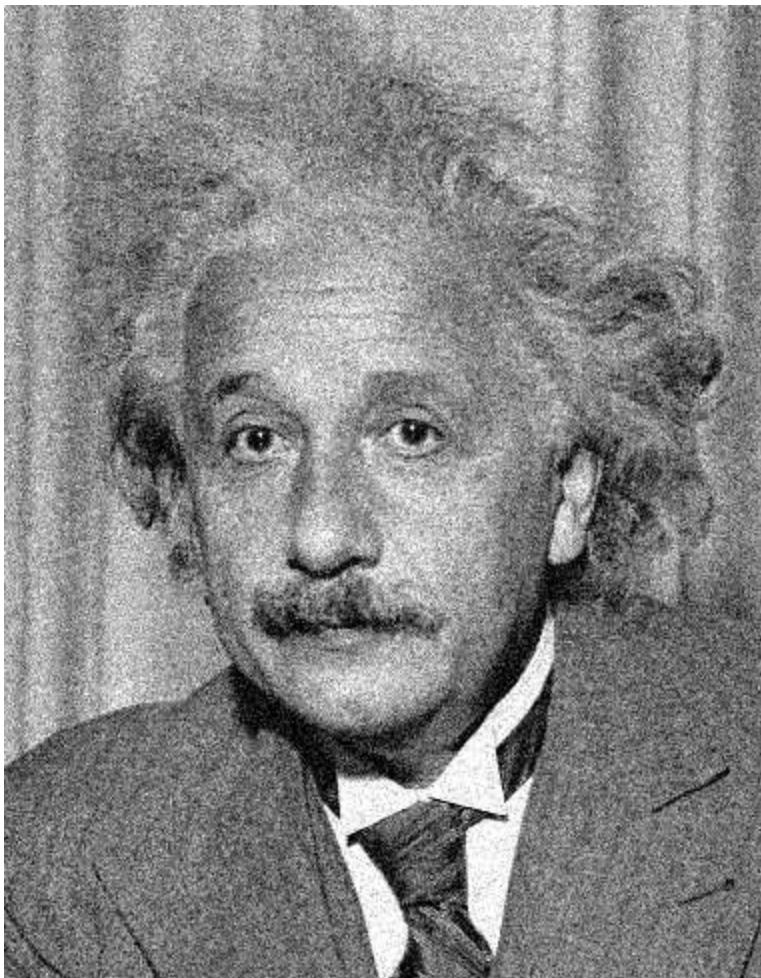
$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} = \begin{matrix} 1 \\ 2 \\ 1 \end{matrix} \times \begin{matrix} 1 & 2 & 1 \end{matrix}$$

# Practical Matters

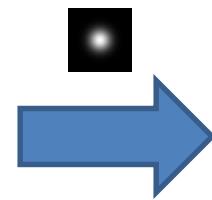
- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - wrap around
    - copy edge
    - reflect across edge



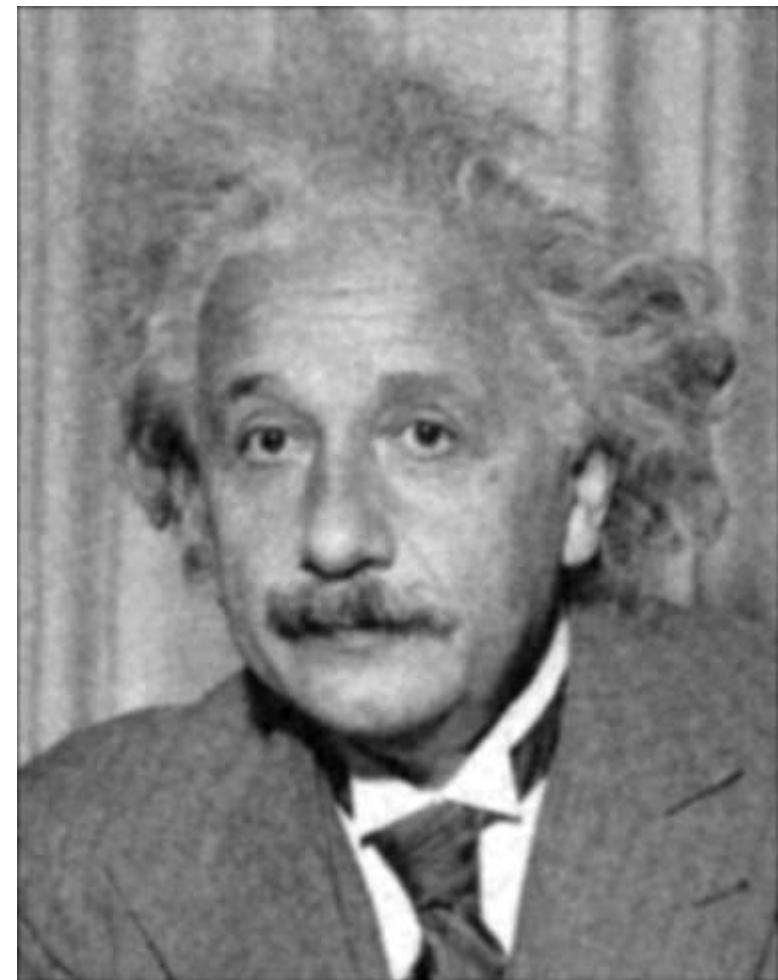
# Denoising



Additive Gaussian Noise



Gaussian  
Filter



# Salt-and-Pepper Noise: Gaussian

3x3



5x5

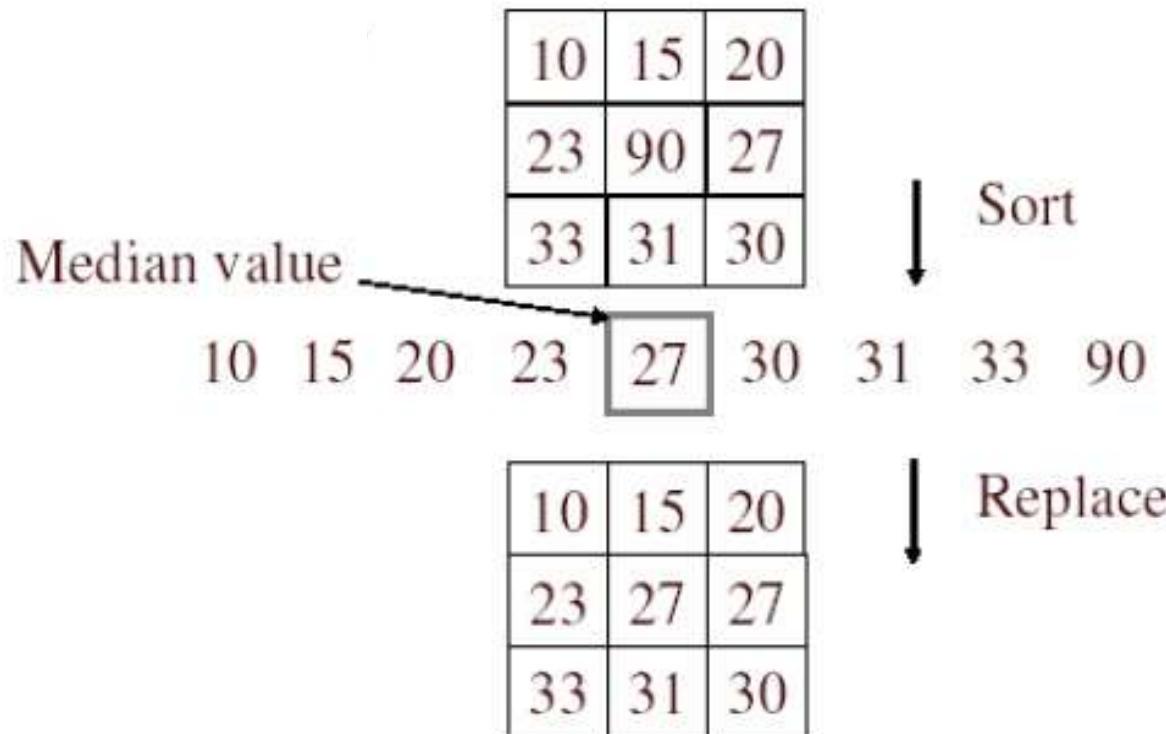


7x7



# Alternative Idea: Median Filtering

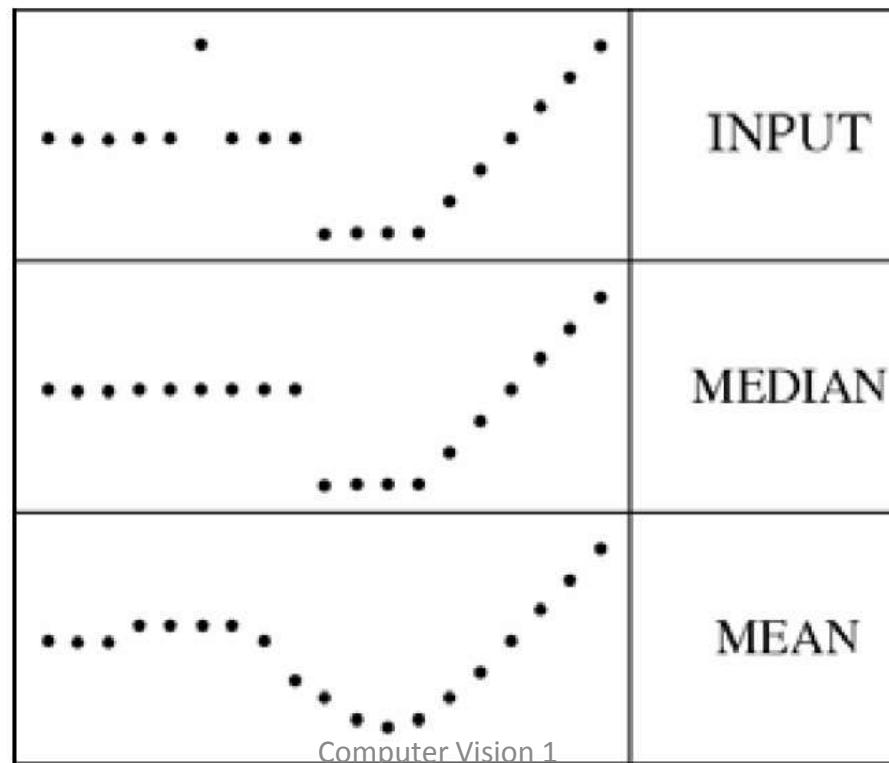
- A **median filter** operates over a window by selecting the median intensity in the window



# Median Filter

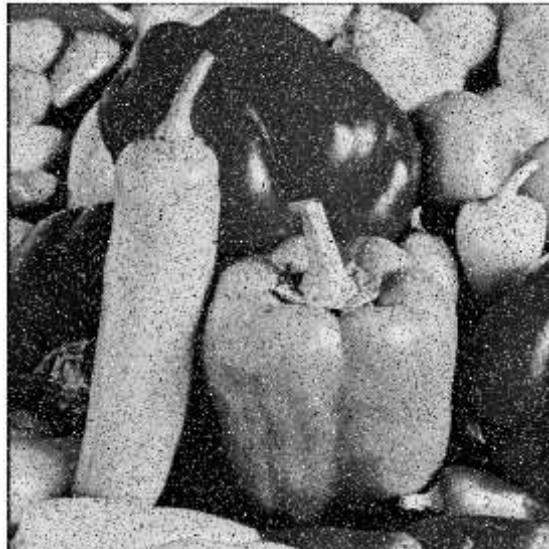
- What advantage does median filtering have over Gaussian filtering?
  - Robustness to outliers

filters have width 5 :

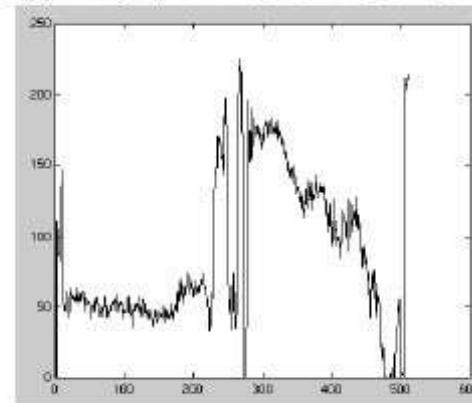
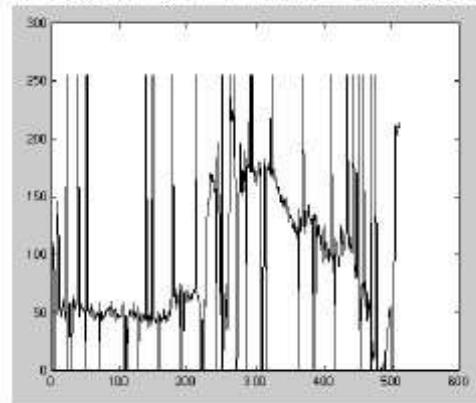


# Median Filter

Salt-and-pepper noise



Median filtered



# Median vs. Gaussian Filtering

3x3



5x5



7x7



Gaussian

Median



# Other Non-linear Filters

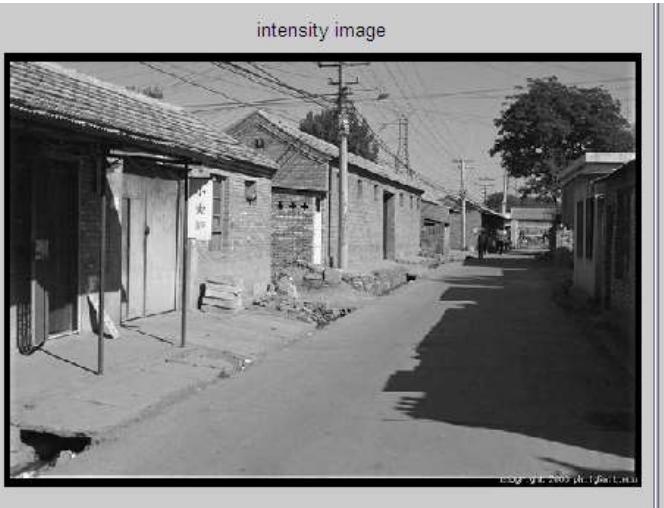
- Weighted median (pixels further from center count less)
- Clipped mean (average, ignoring few brightest and darkest pixels)
- Bilateral filtering (weight by spatial distance *and* intensity difference)



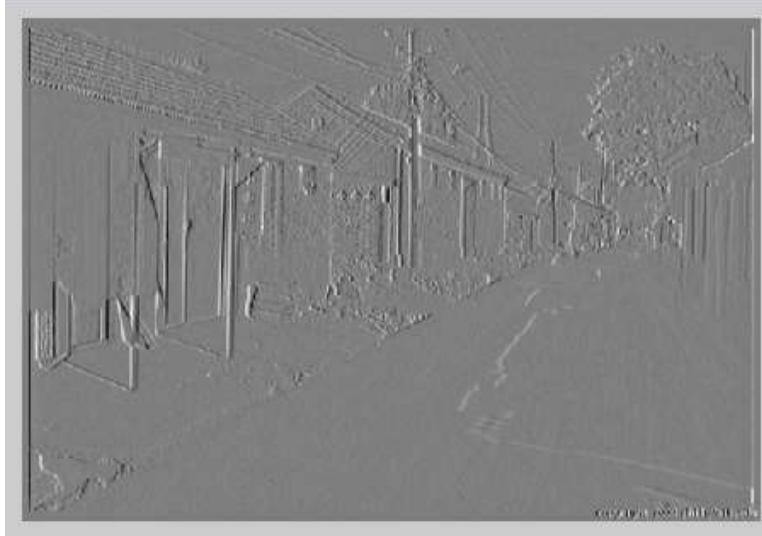
Bilateral filtering

# Review: Edge Detection

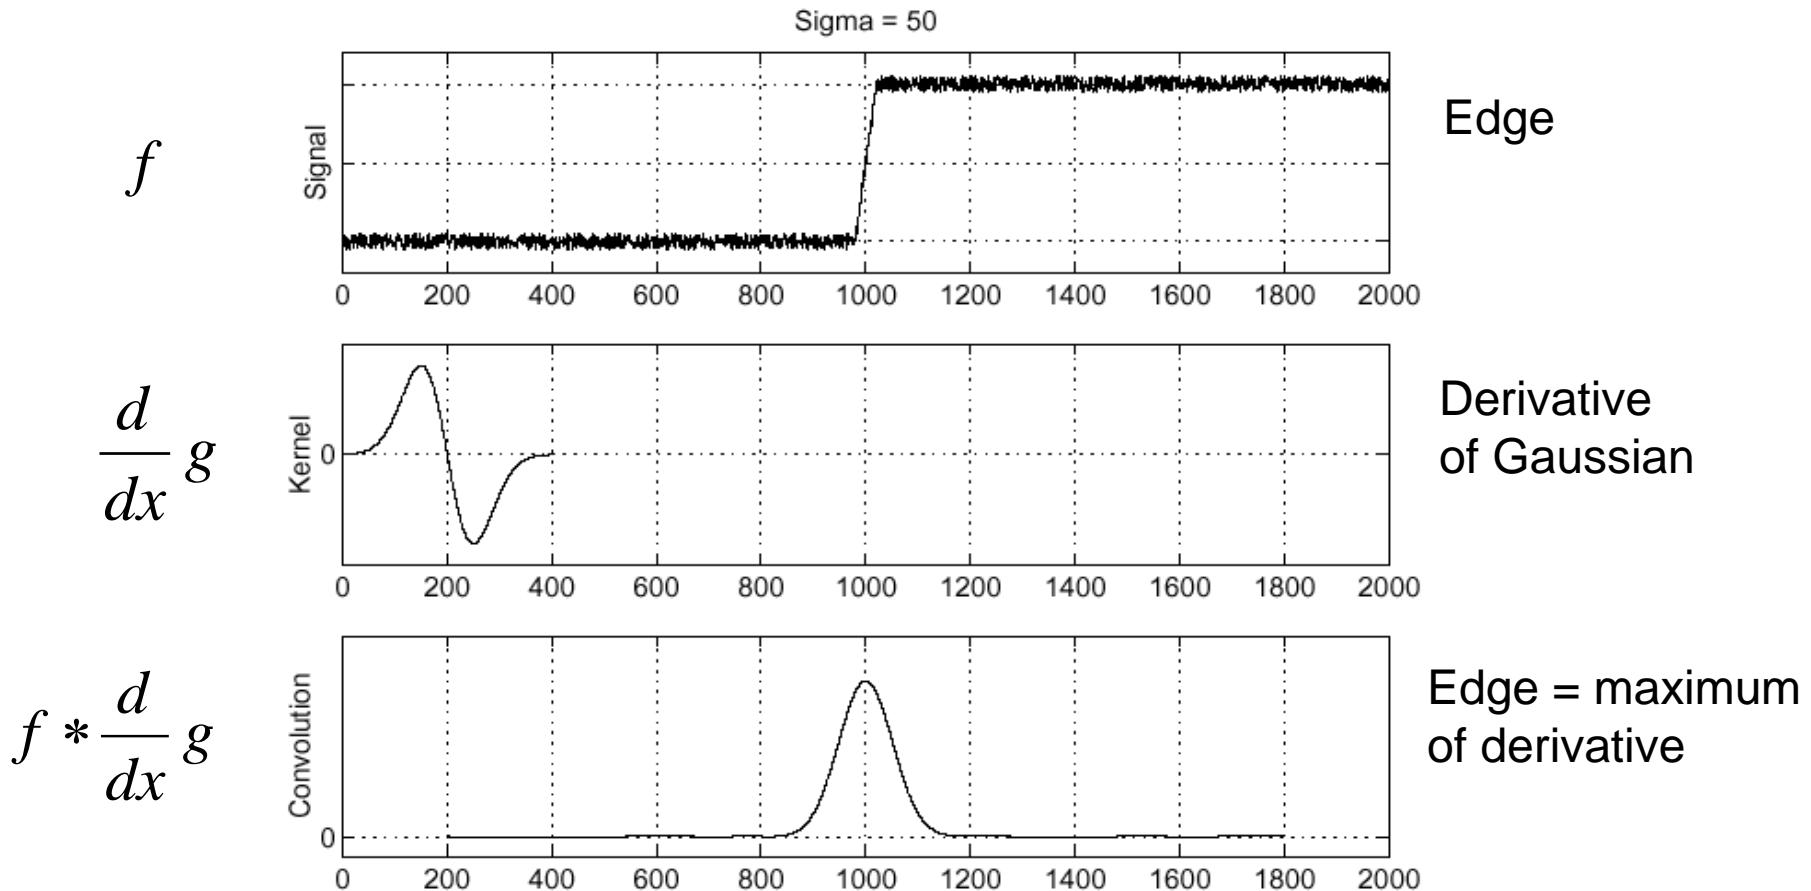
1	0	-1
2	0	-2
1	0	-1



$$\begin{matrix} * & \end{matrix} = \begin{matrix} \text{edge map} \end{matrix}$$

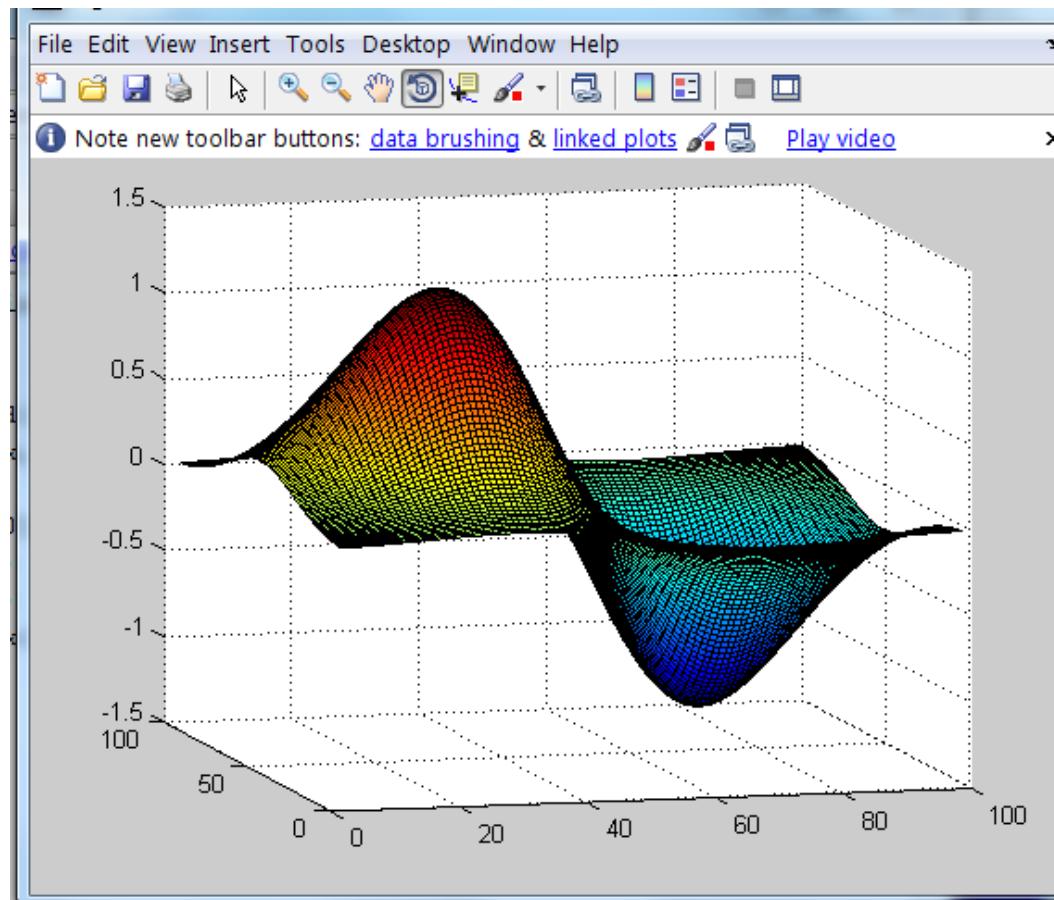


# Review: Edge Detection

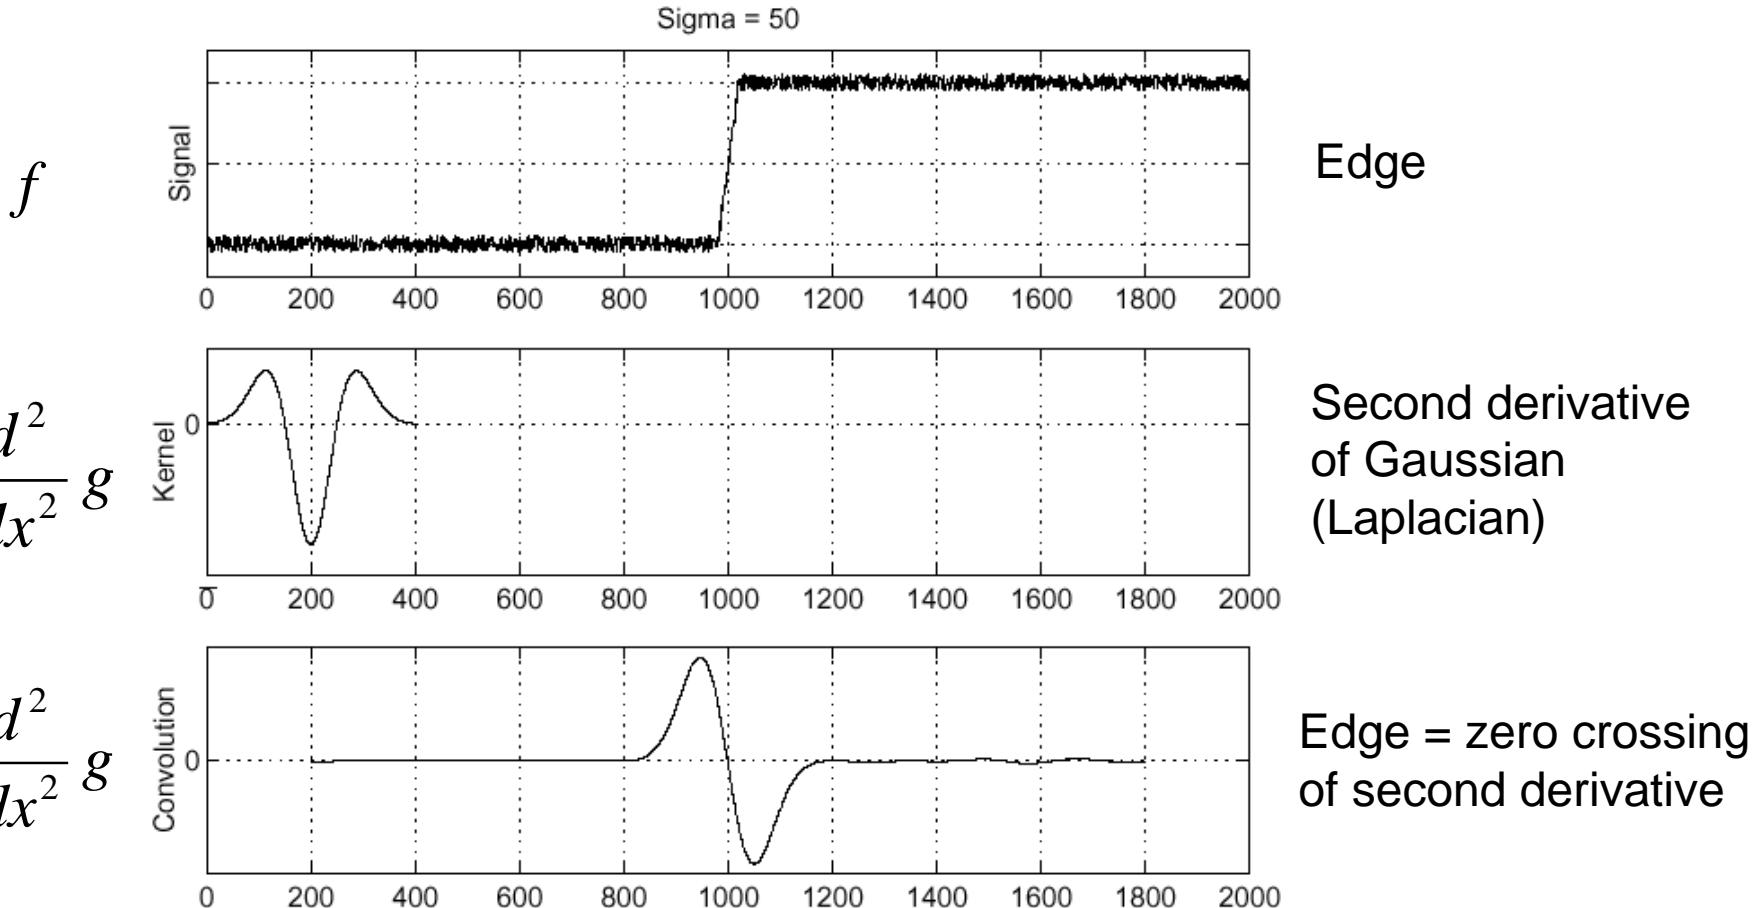


# Review

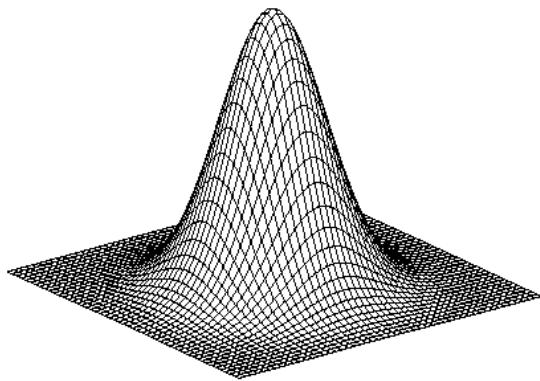
## Derivative of Gaussian



# Review: Edge Detection

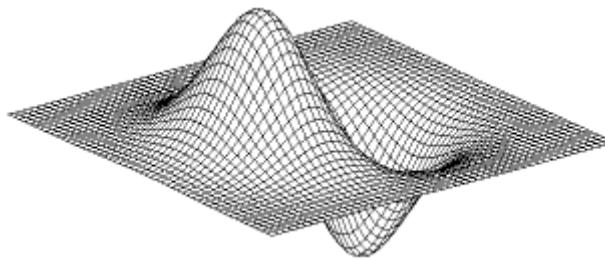


# 2D Edge Detection Filters



Gaussian

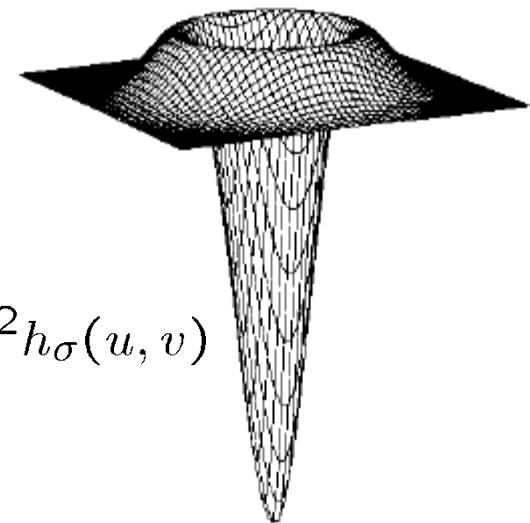
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

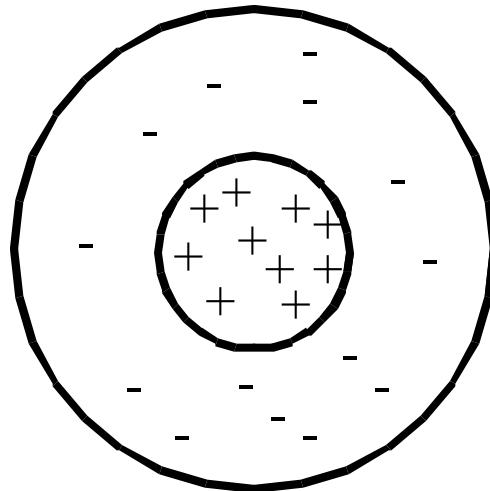
Laplacian of Gaussian



$\nabla^2$  is the **Laplacian** operator:

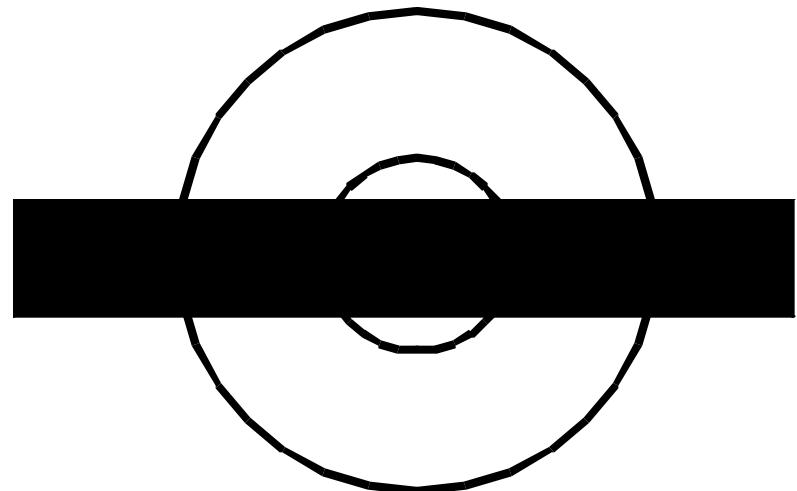
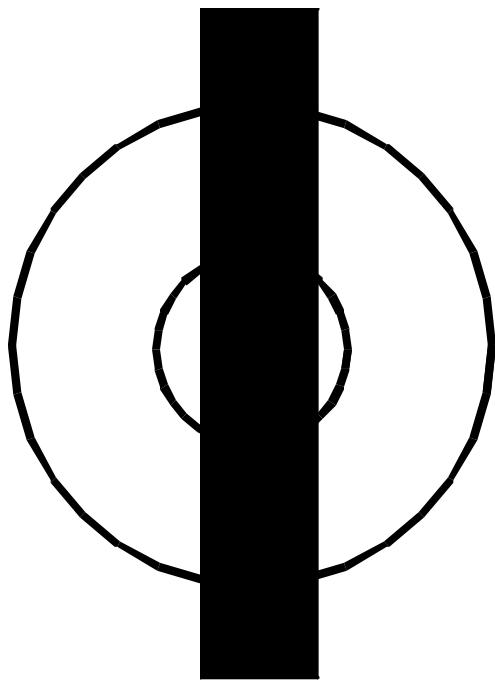
$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

# “On-Center” Ganglion Cell

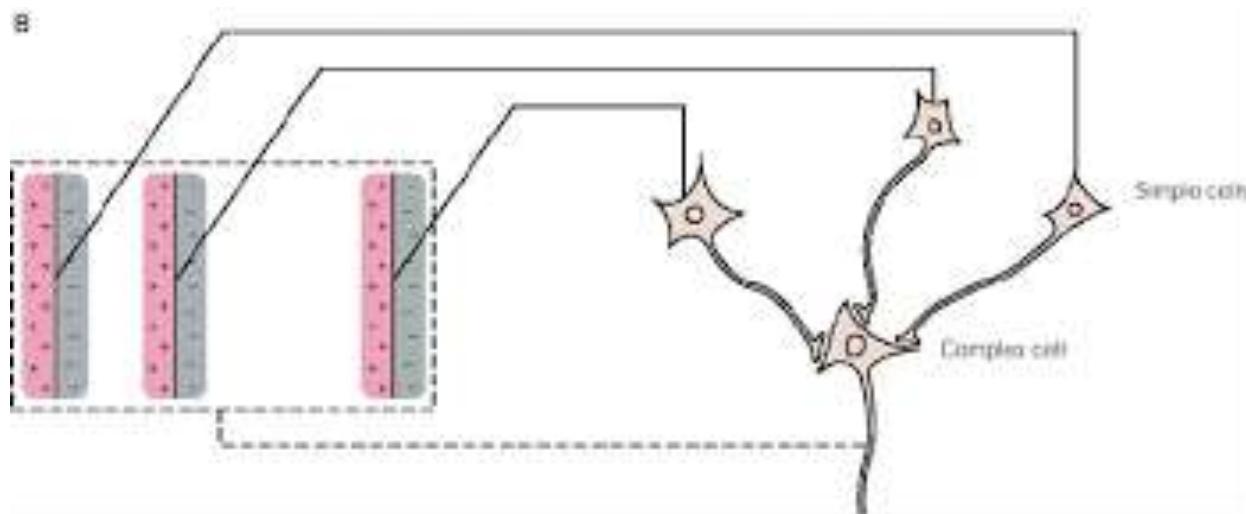
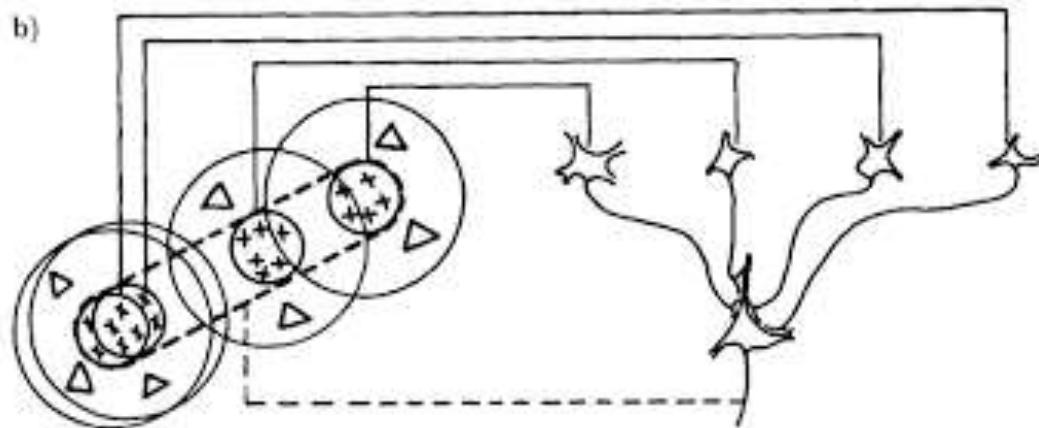


Responds maximally to light increments in the center,  
and light decrements in the surround.

Ganglion cells have no orientation preference.

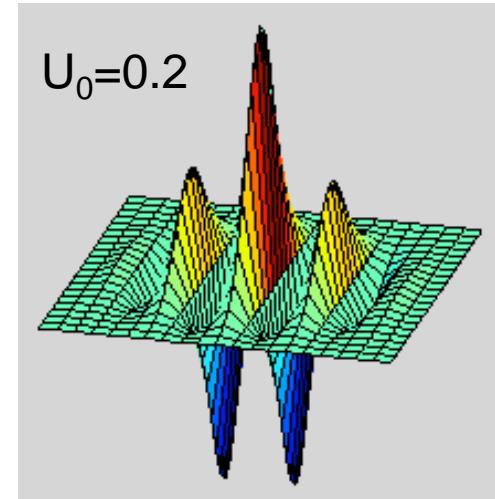
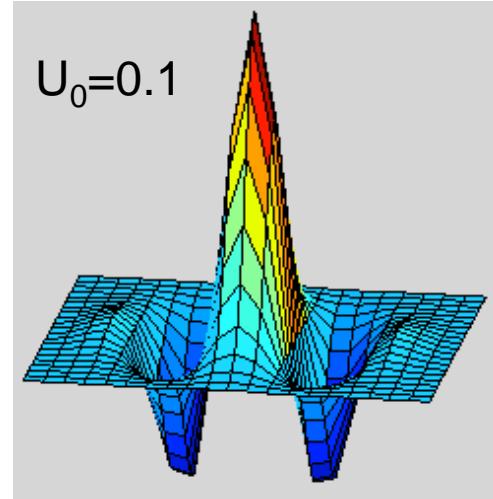
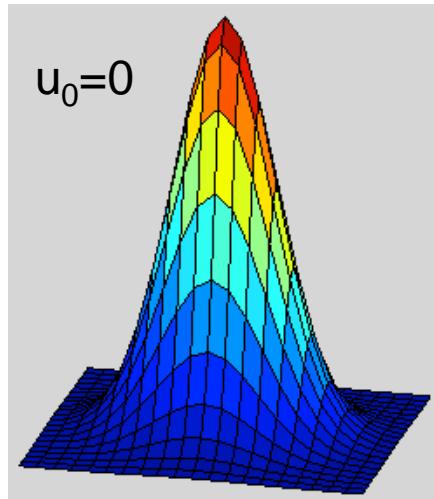


# Model of how center-surround cells can be building blocks for simple cells.

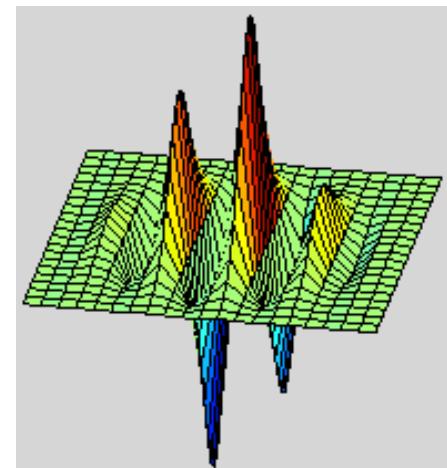
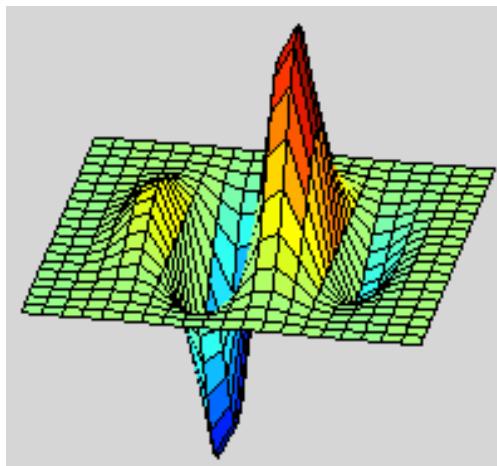


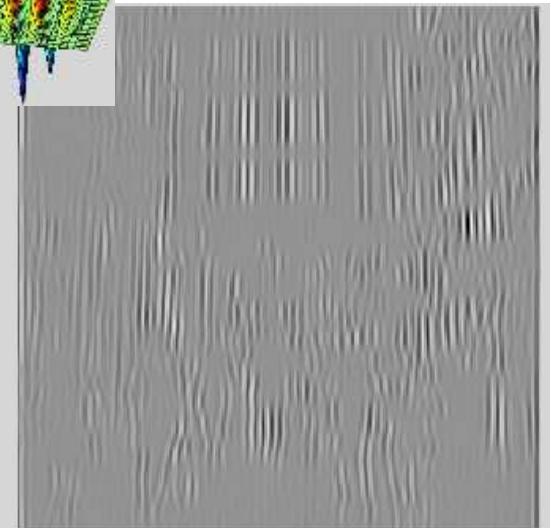
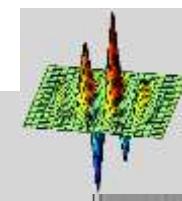
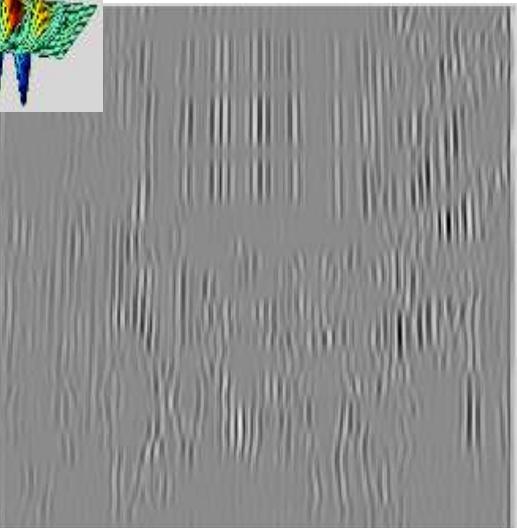
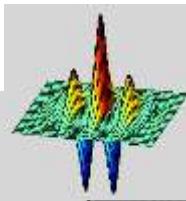
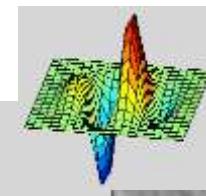
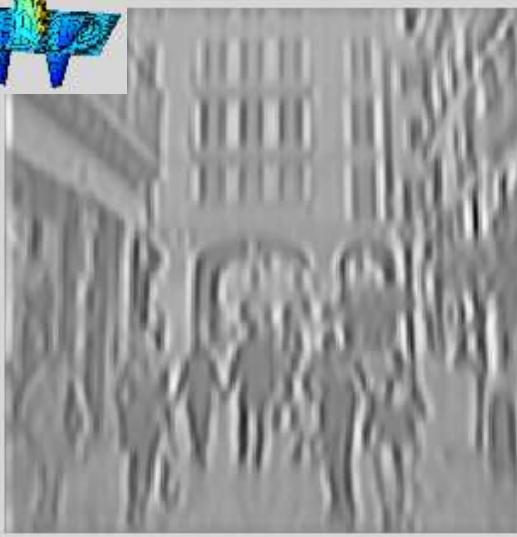
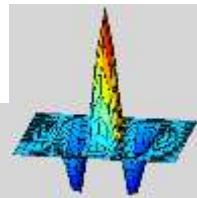
# Gabor wavelets

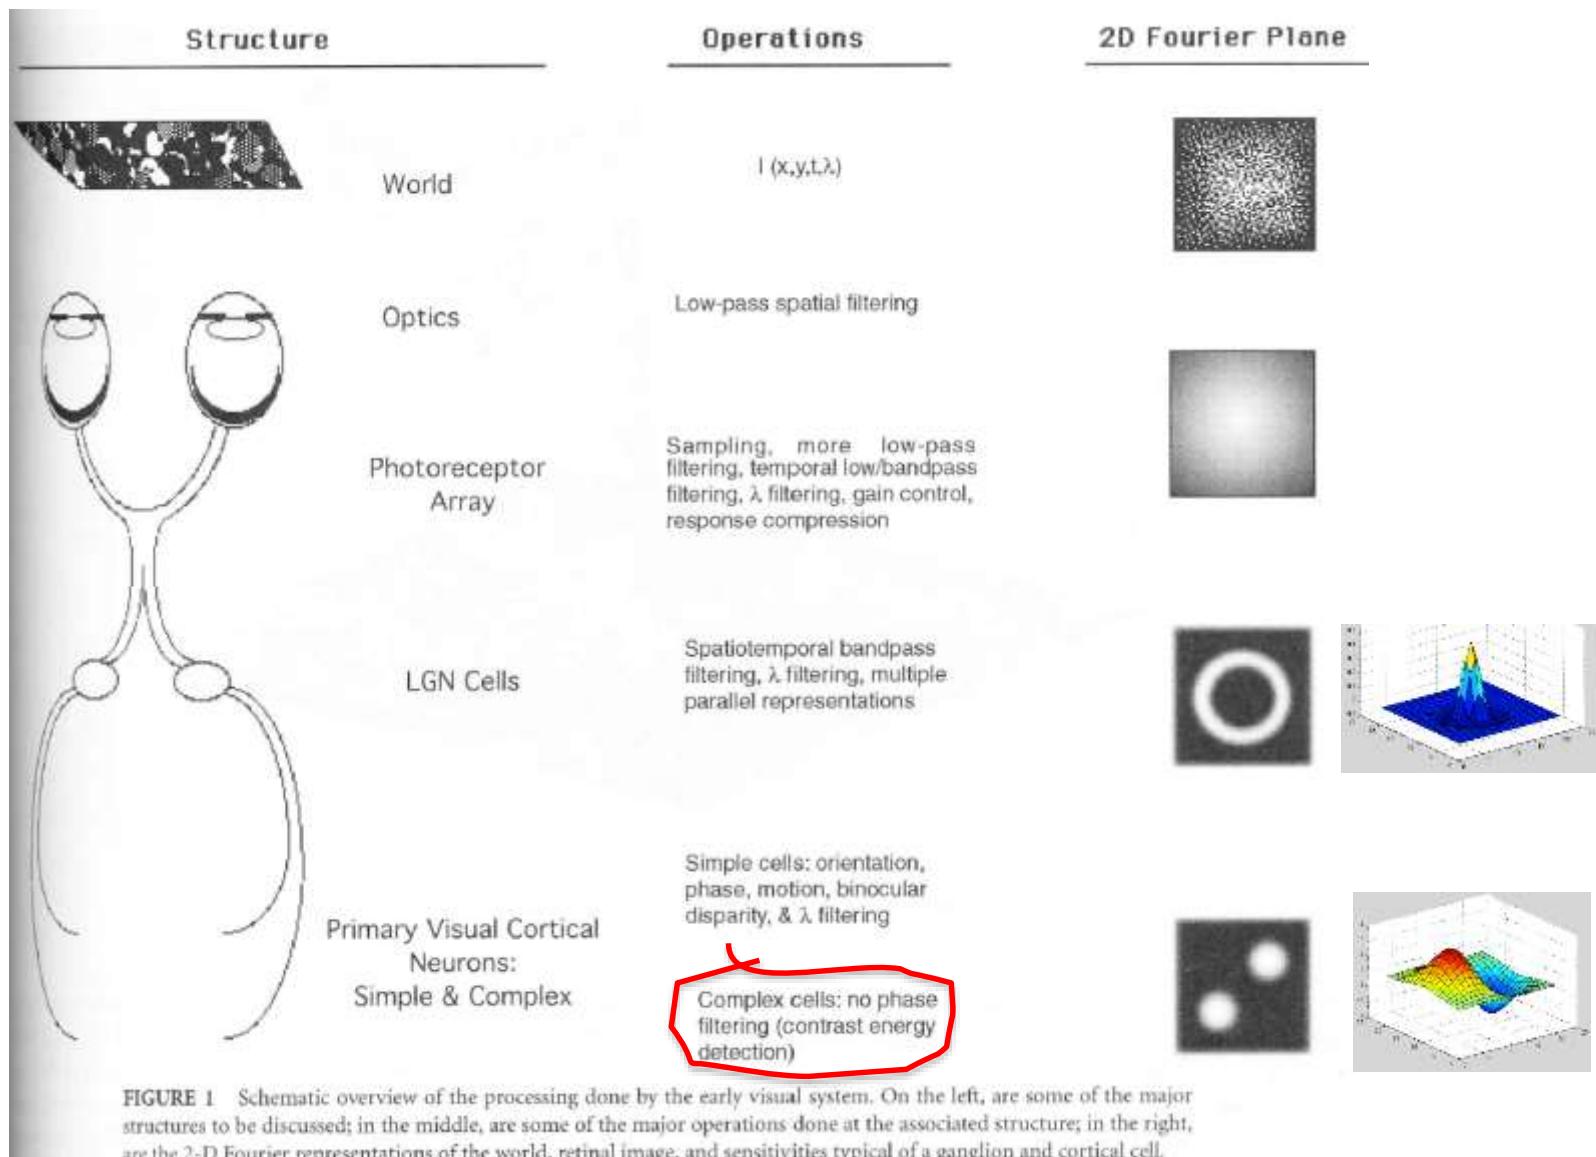
$$\psi_c(x, y) = e^{-\frac{x^2+y^2}{2\sigma^2}} \cos(2\pi u_0 x)$$



$$\psi_s(x, y) = e^{-\frac{x^2+y^2}{2\sigma^2}} \sin(2\pi u_0 x)$$



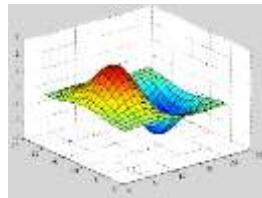




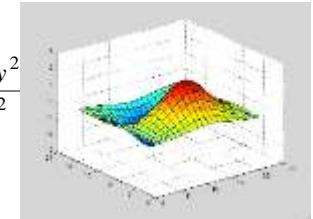
# Steerable filters

Derivatives of a Gaussian:

$$h_x(x, y) = \frac{\partial h(x, y)}{\partial x} = \frac{-x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



$$h_y(x, y) = \frac{\partial h(x, y)}{\partial y} = \frac{-y}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

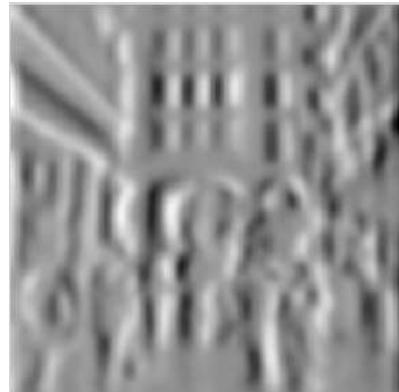


An arbitrary orientation can be computed as a linear combination of those two basis functions:

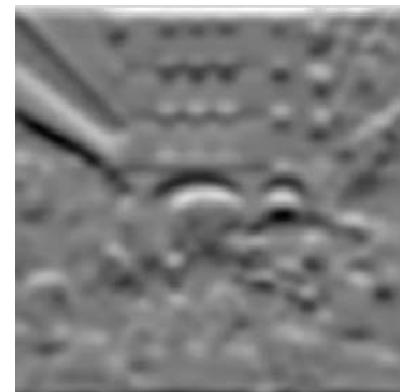
$$h_\alpha(x, y) = \cos(\alpha)h_x(x, y) + \sin(\alpha)h_y(x, y)$$

The representation is “shiftable” on orientation: We can interpolate any other orientation from a finite set of basis functions.

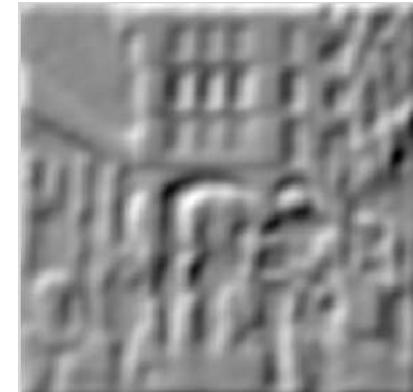
$\cos(\langle \rangle)$



$+\sin(\langle \rangle)$

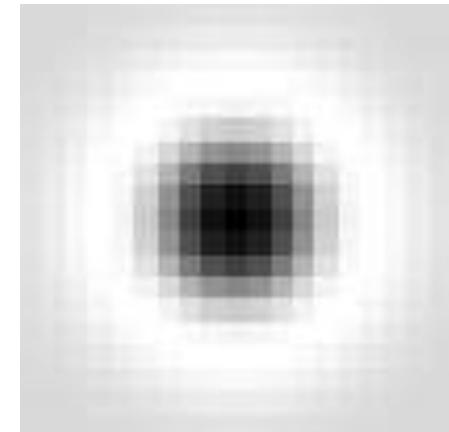
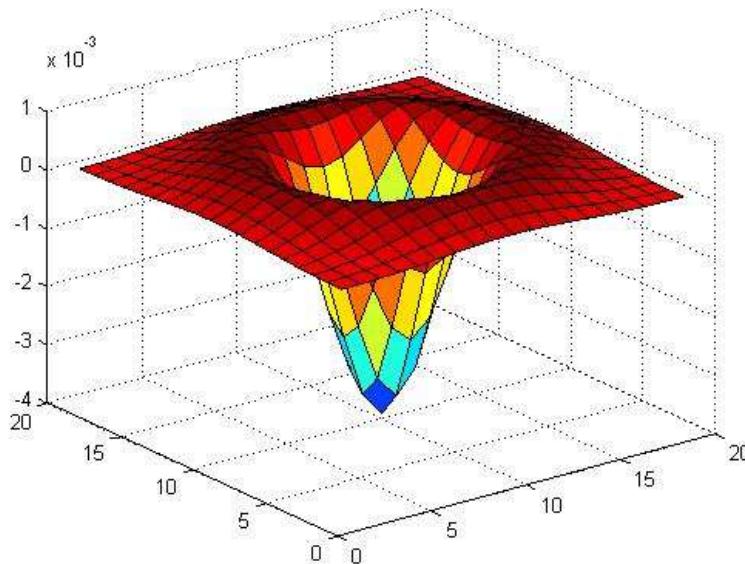


=



# Laplacian of Gaussian

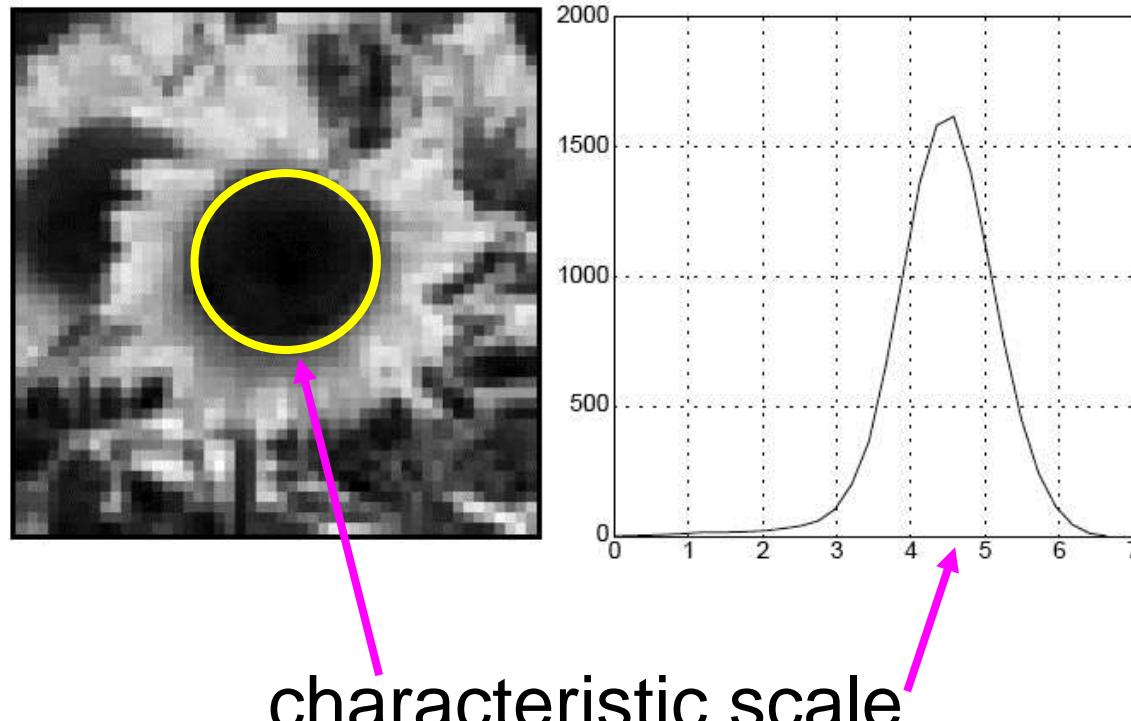
- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D



$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

# Blob Detection in 2D

- We define the *characteristic scale* as the scale that produces peak of Laplacian response

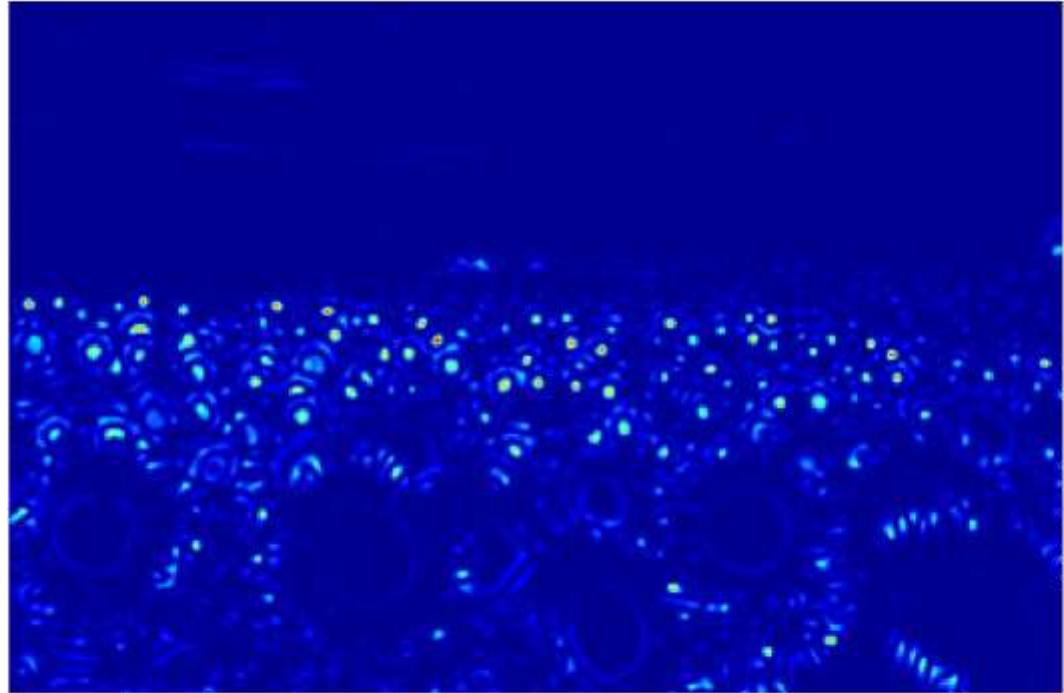


# Example

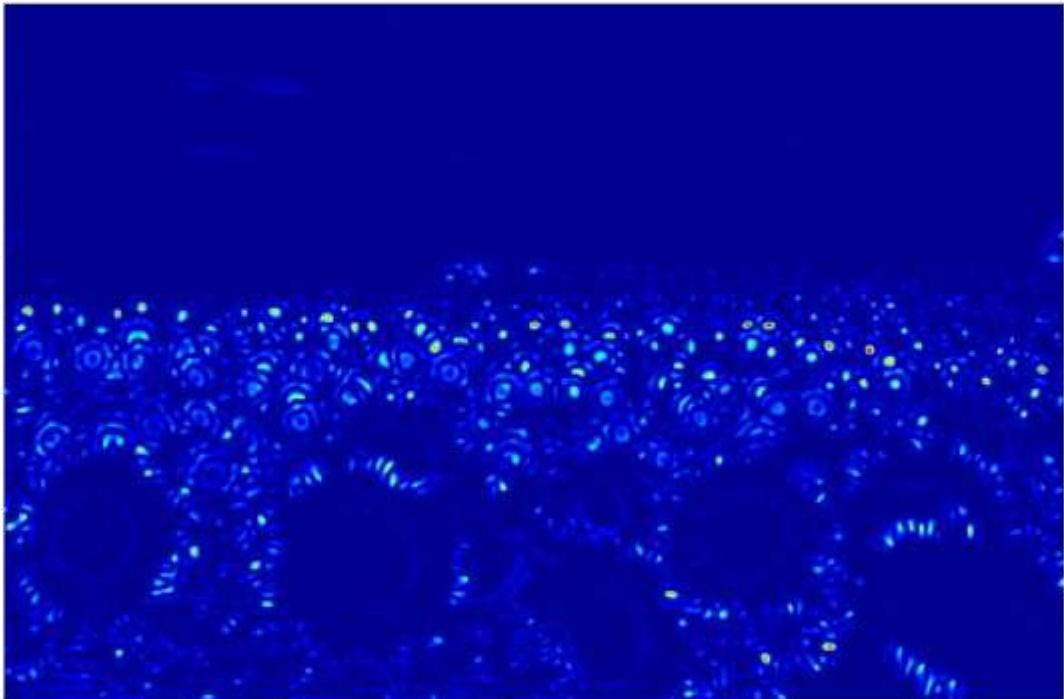
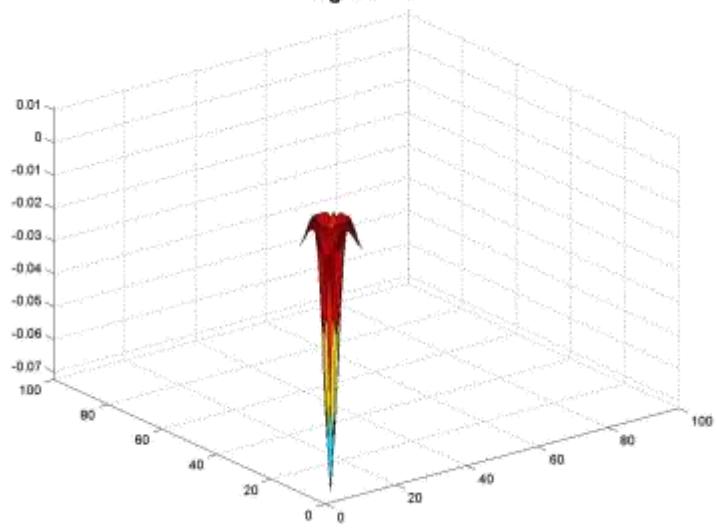
Original image  
at  $\frac{3}{4}$  the size

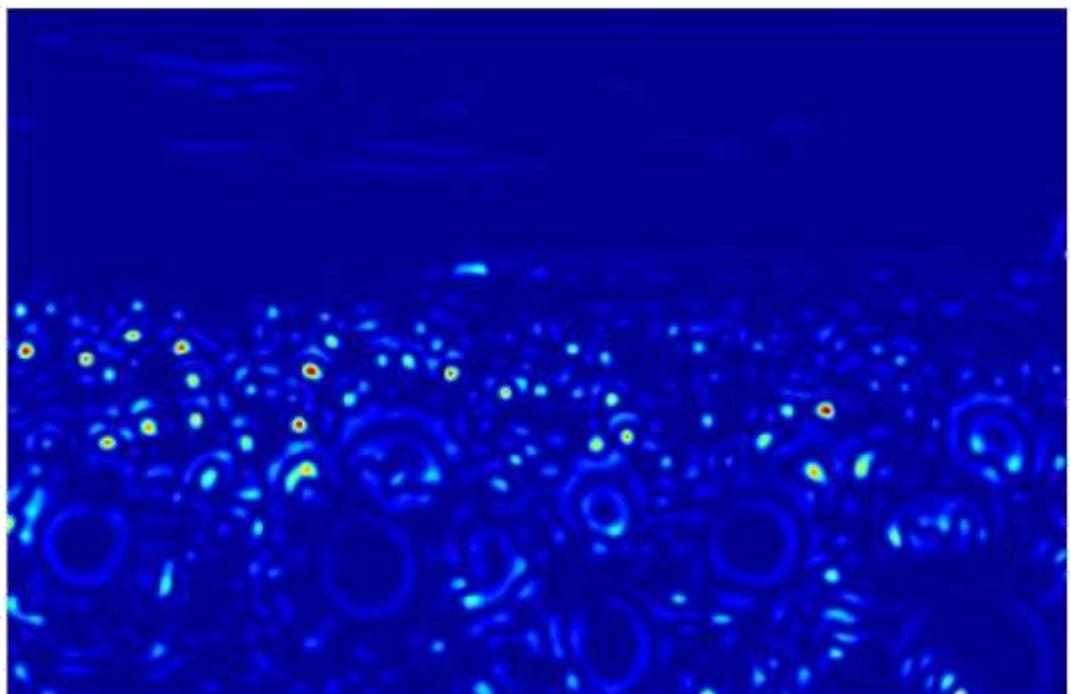
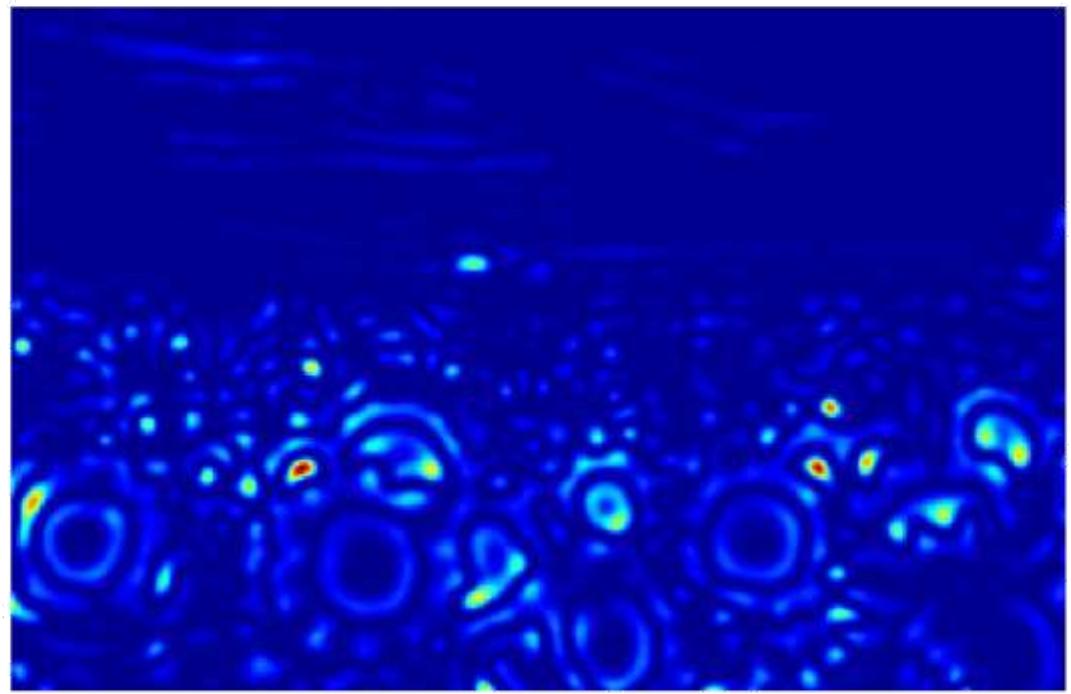


Original image  
at  $\frac{3}{4}$  the size

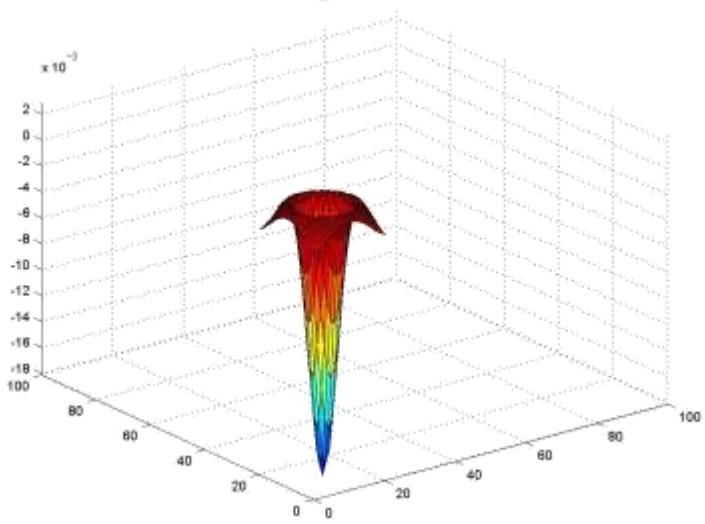


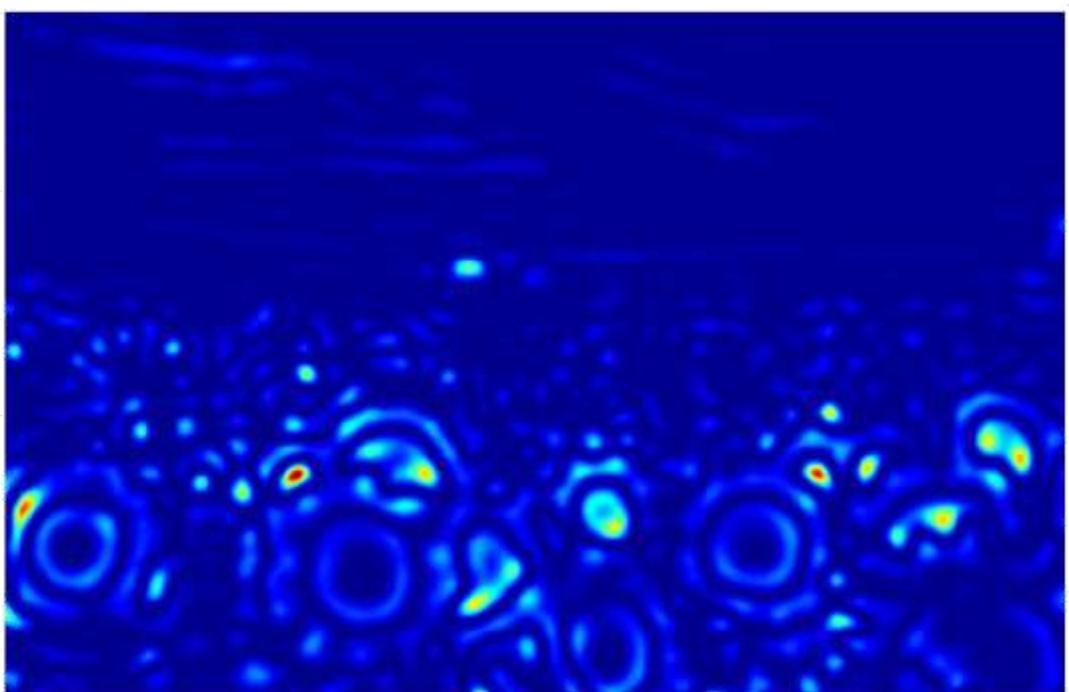
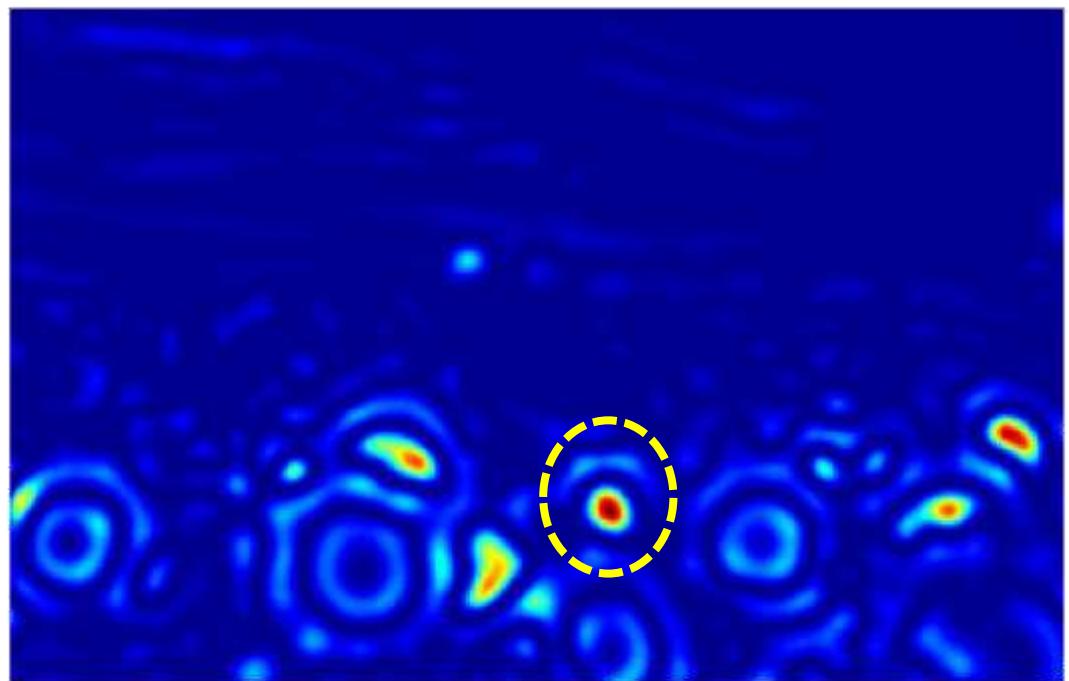
**sigma=2.1**



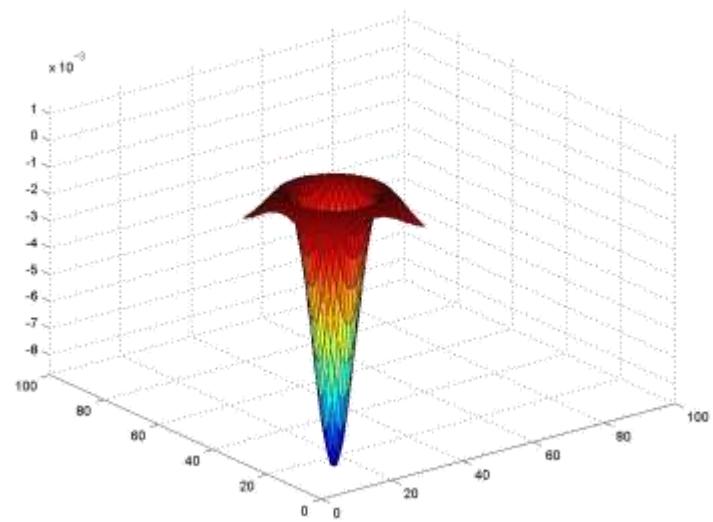


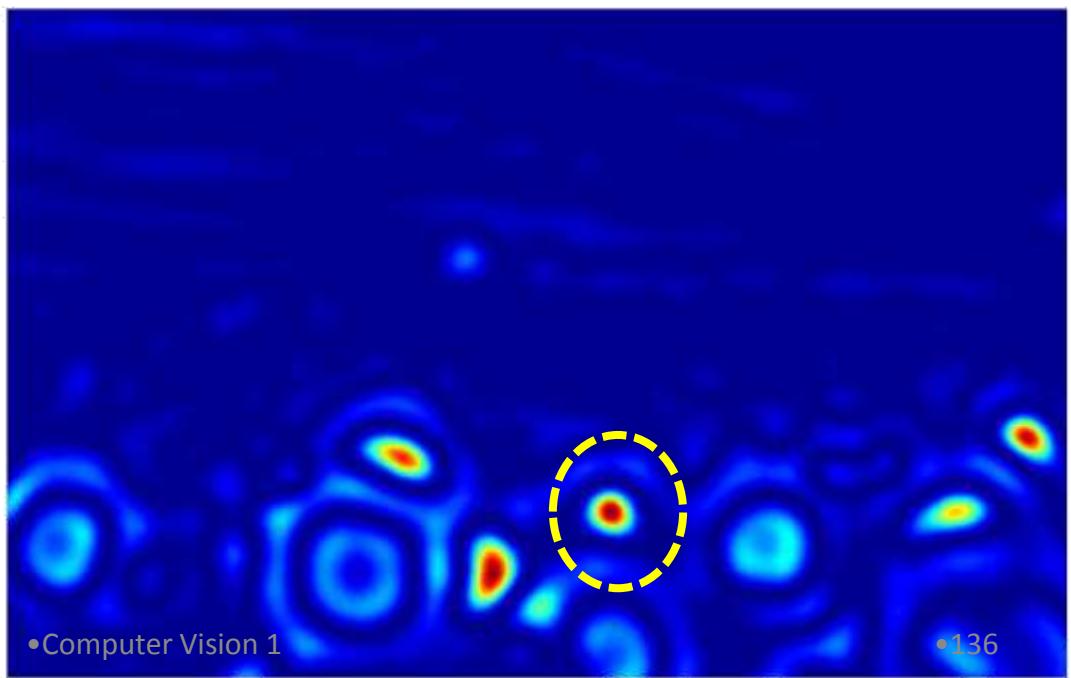
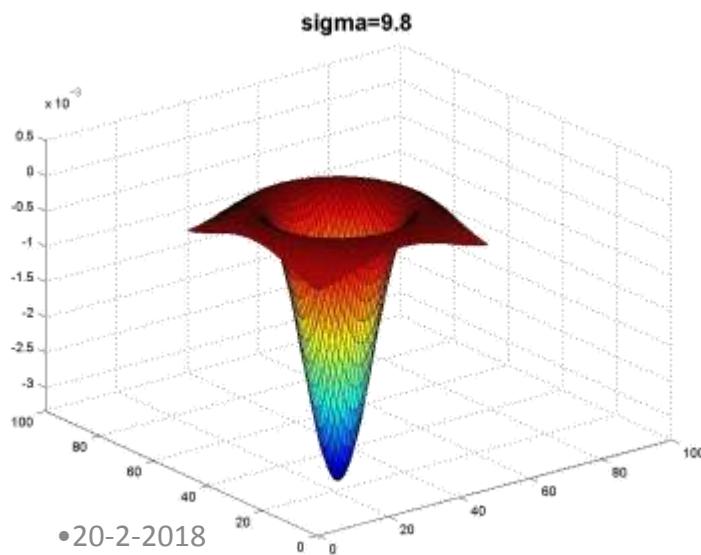
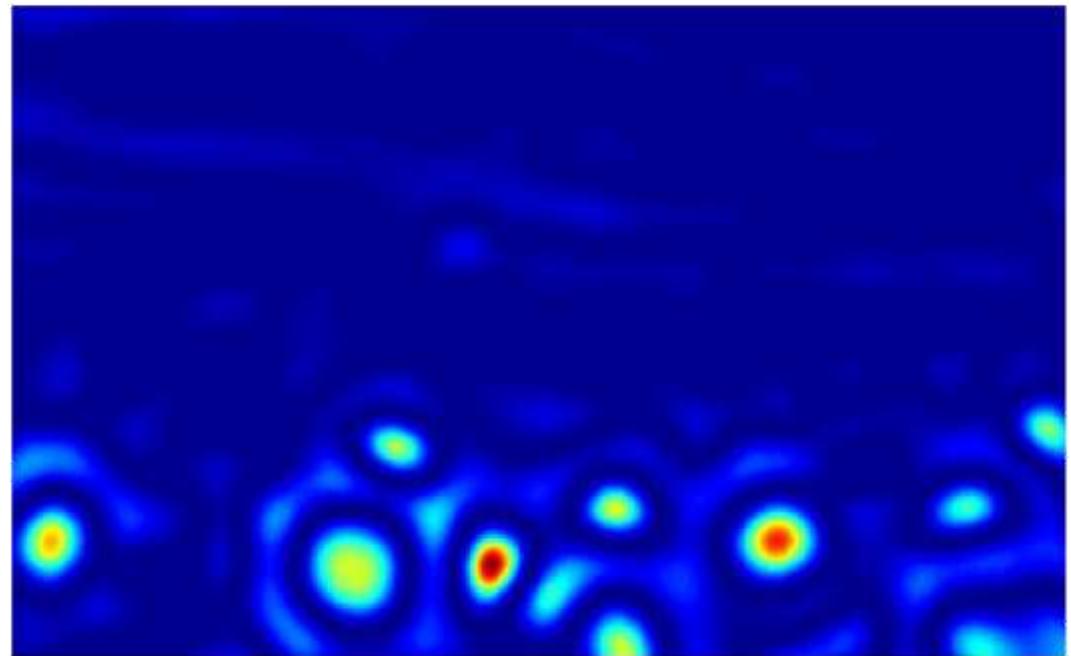
**sigma=4.2**

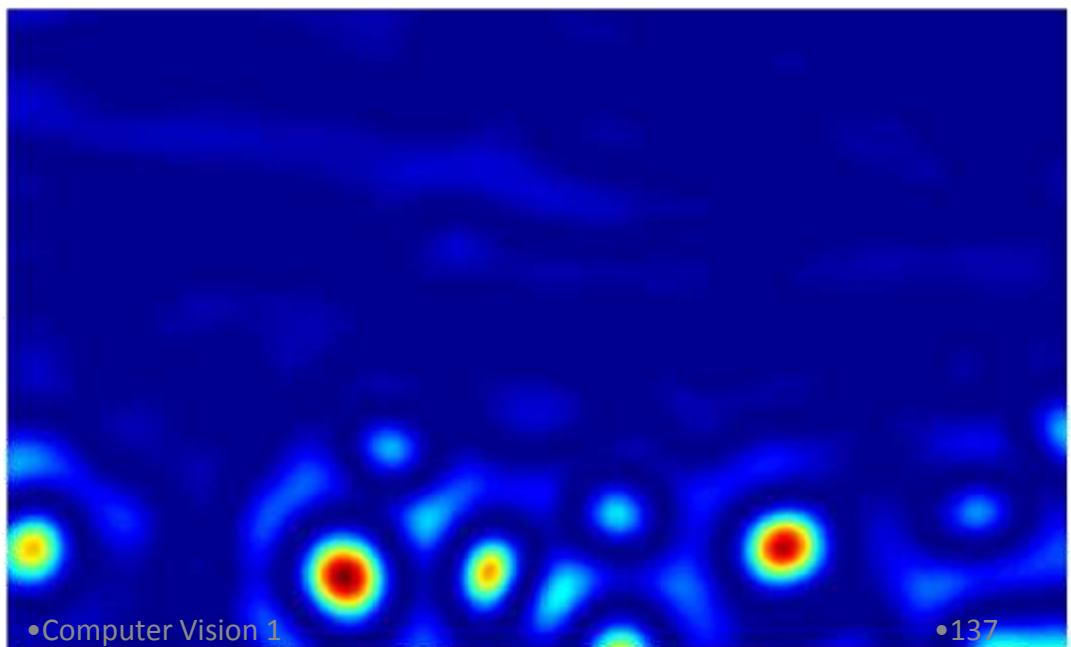
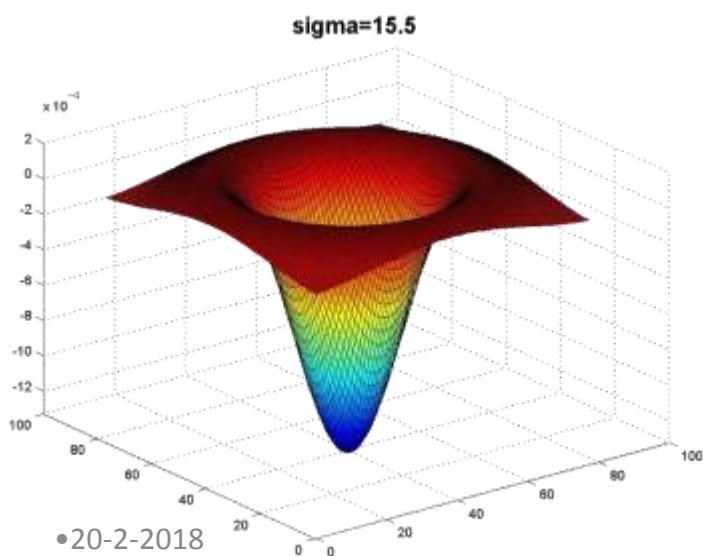
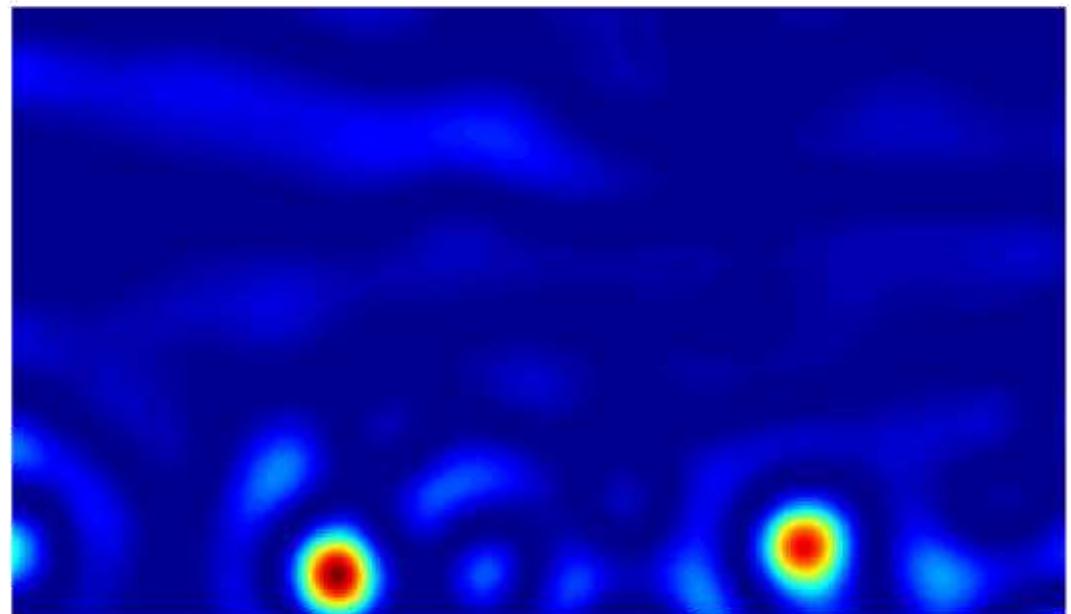




**sigma=6**







•Computer Vision 1

•137

# Scale Invariant Interest Points

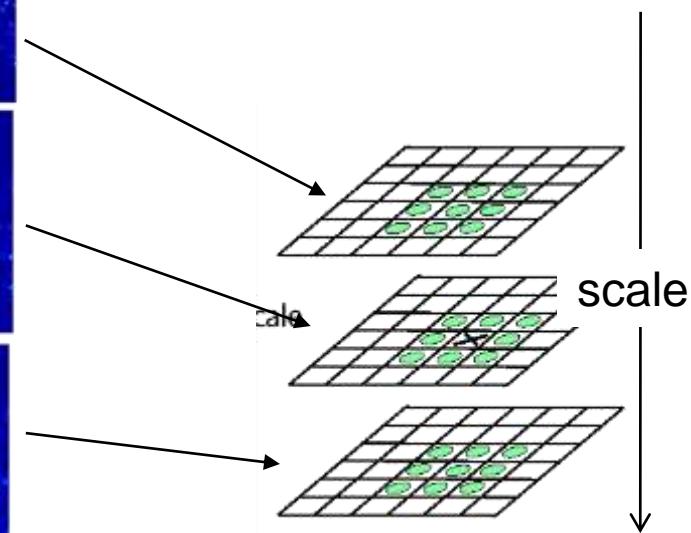
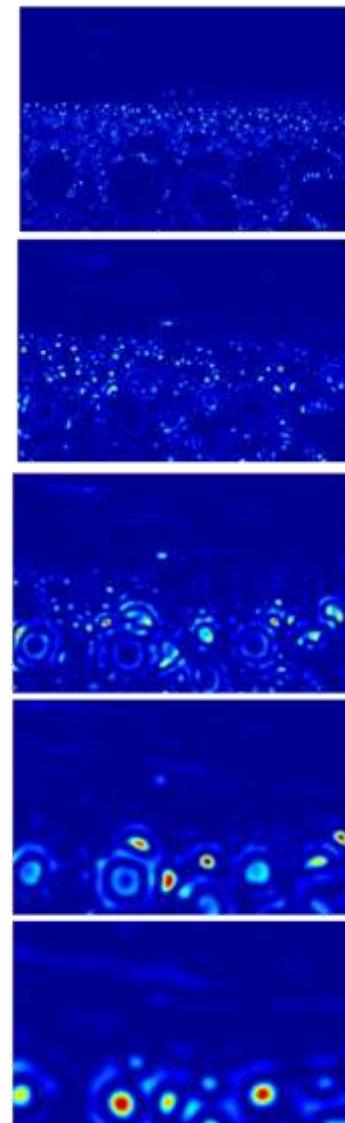
Interest points are local maxima in both position and scale.



$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow \sigma_3$$

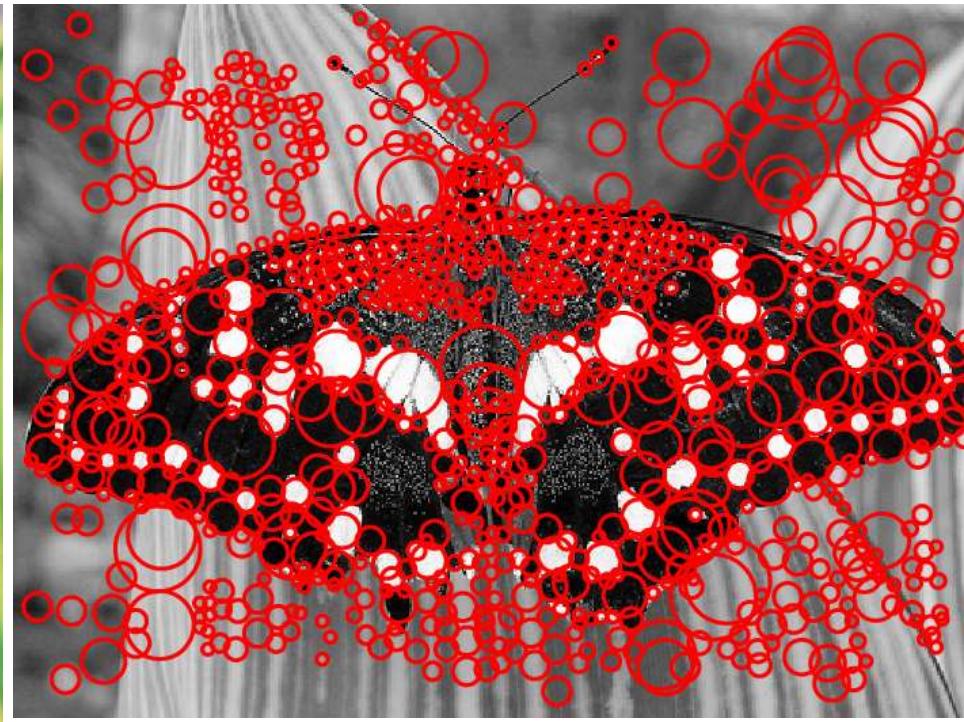
Squared filter response maps

$\sigma_5$   
 $\sigma_4$   
 $\sigma_3$   
 $\sigma_2$   
 $\sigma_1$



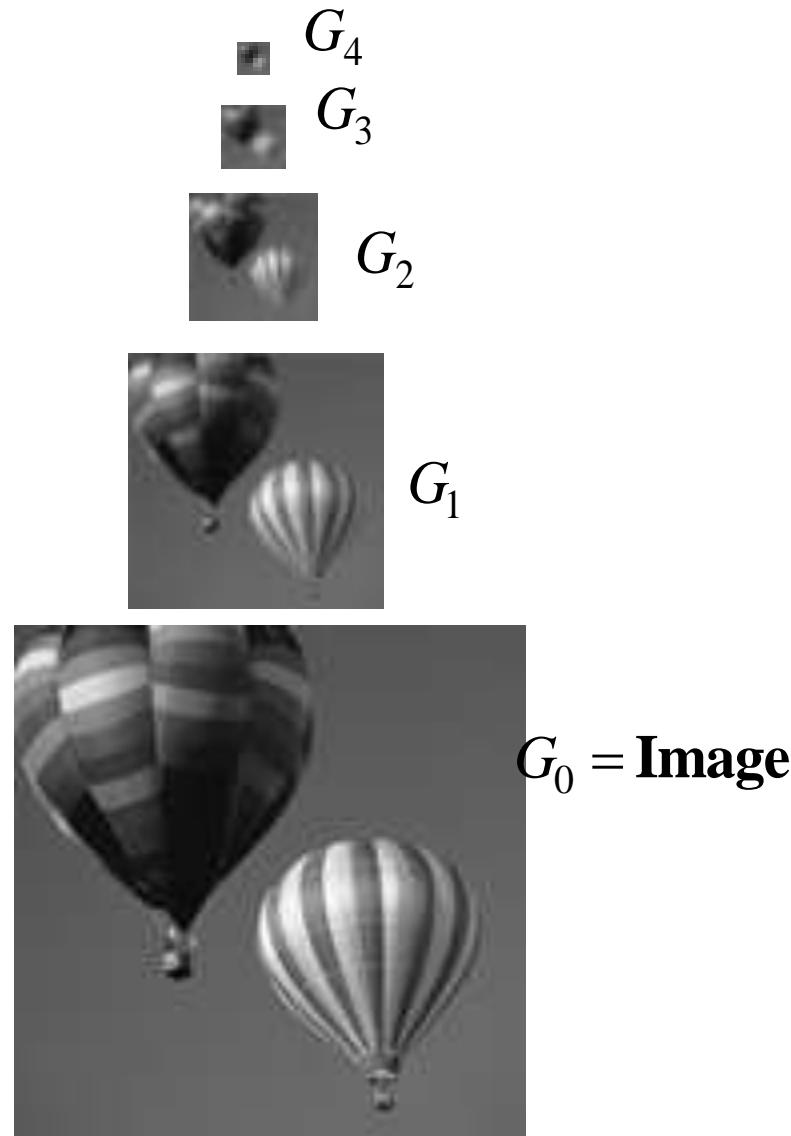
⇒ List of  
( $x$ ,  $y$ ,  $\sigma$ )

# Scale-space Blob Detector: Example



# The Gaussian Pyramid

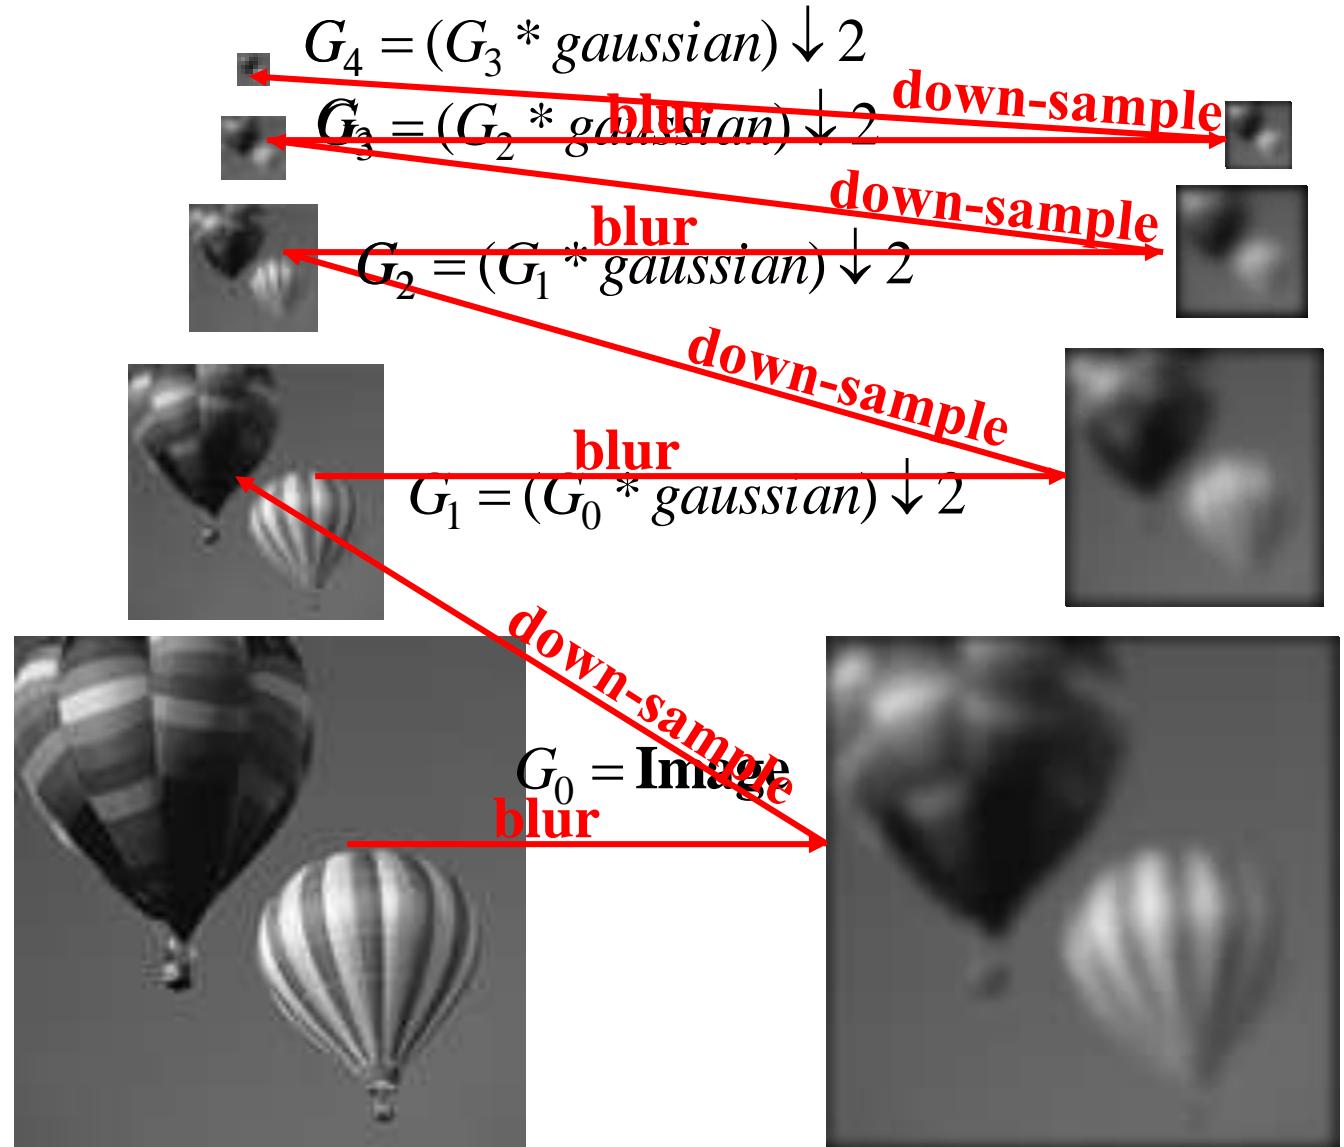
Low resolution



High resolution

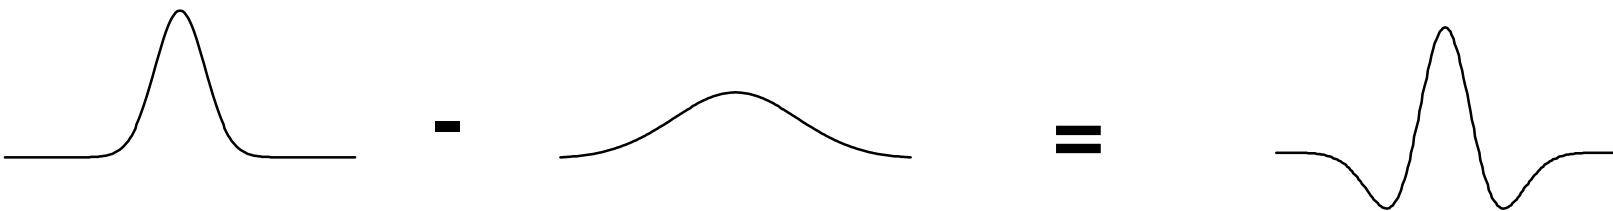
# The Gaussian Pyramid

Low resolution



High resolution

# Laplacian ~ Difference of Gaussians



**DOG = Difference Of Gaussians**

More details on Gaussian and Laplacian pyramids can be found in the paper by Burt and Adelson

# Template Matching with Image Pyramids

Input: Image, Template

1. Match template at current scale
2. Downsample image
3. Repeat 1-2 until image is very small
4. Take responses above some threshold, perhaps with non-maxima suppression

# Summary

- Edge Detection
  - Non-maximum suppression on gradient.
  - Zero-crossings for Laplacian.
  - Scale invariance
    - Gaussian pyramids. Coarse-to-fine search, multi-scale detection.
    - Laplacian pyramid. More compact image representation

# *Computer Vision 1*

*(total #slides 145 | Lecture 3)*

## Summary

1. Color invariants
2. Image processing
3. Filtering and convolutions